

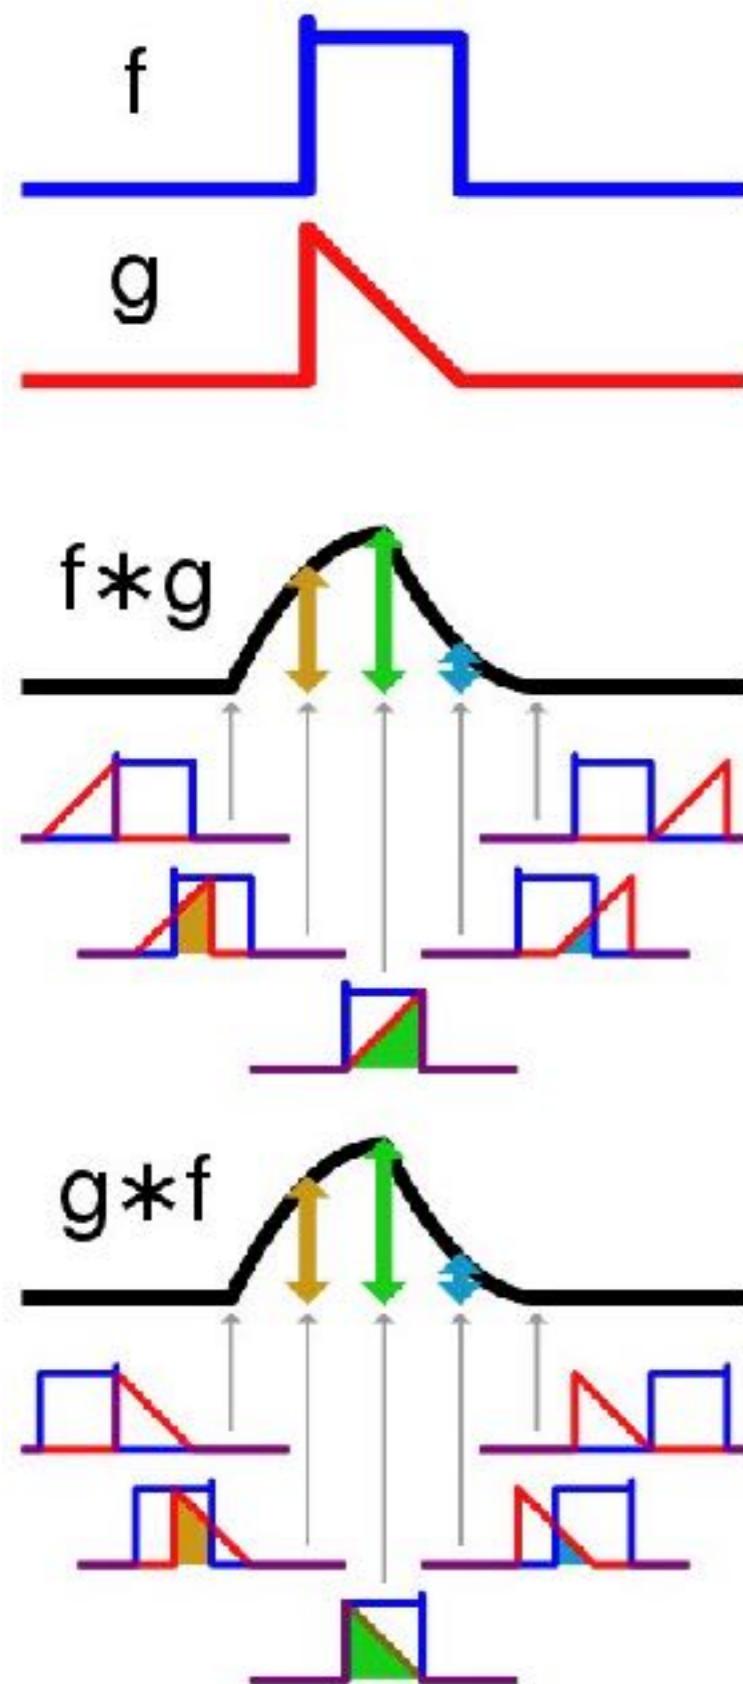
Image and Signal Processing

TechX CV – Lecture 2

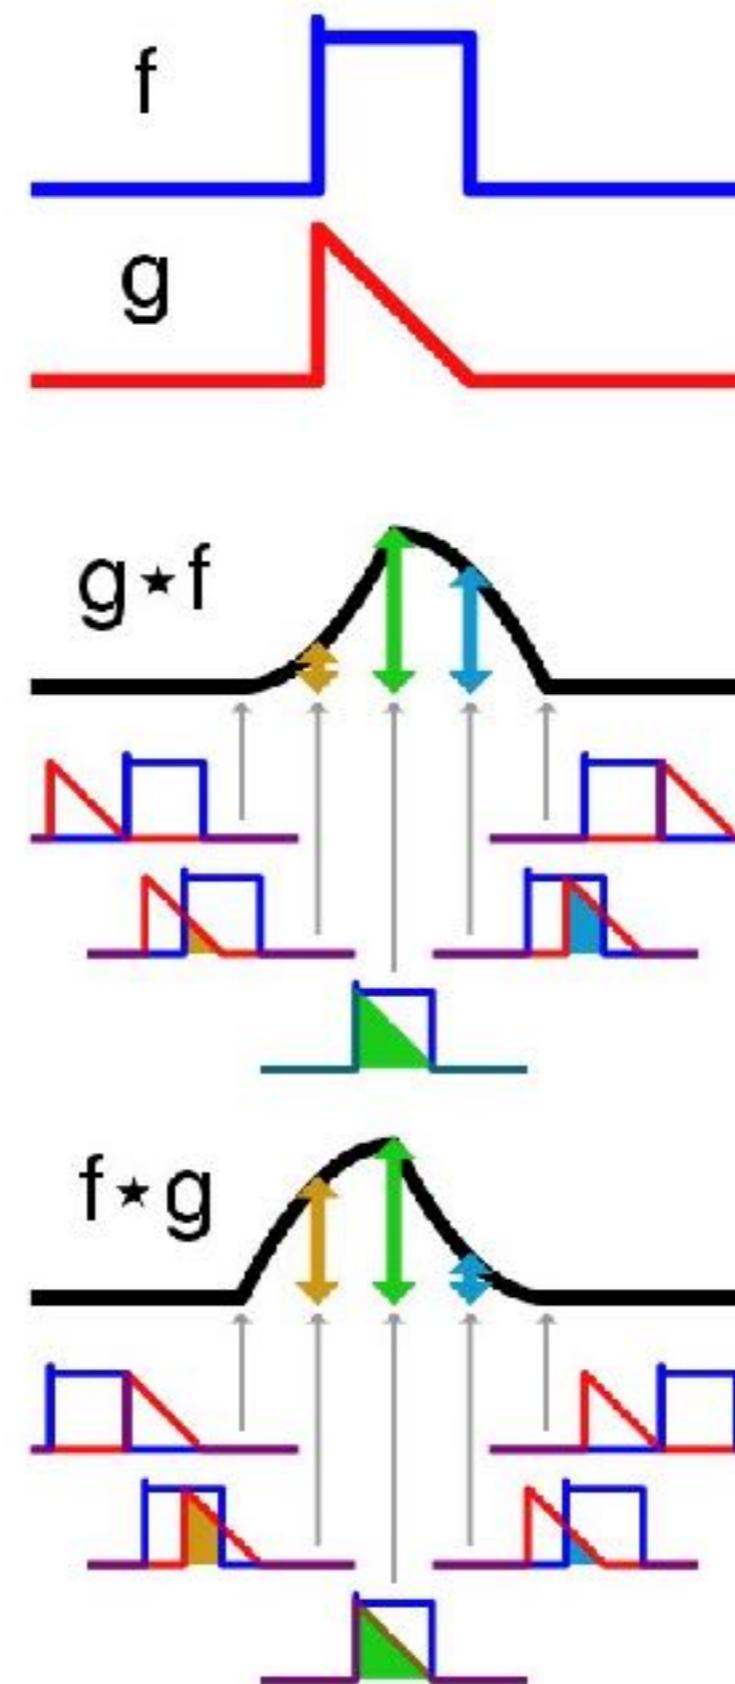
Cecilia Zhang

Recap

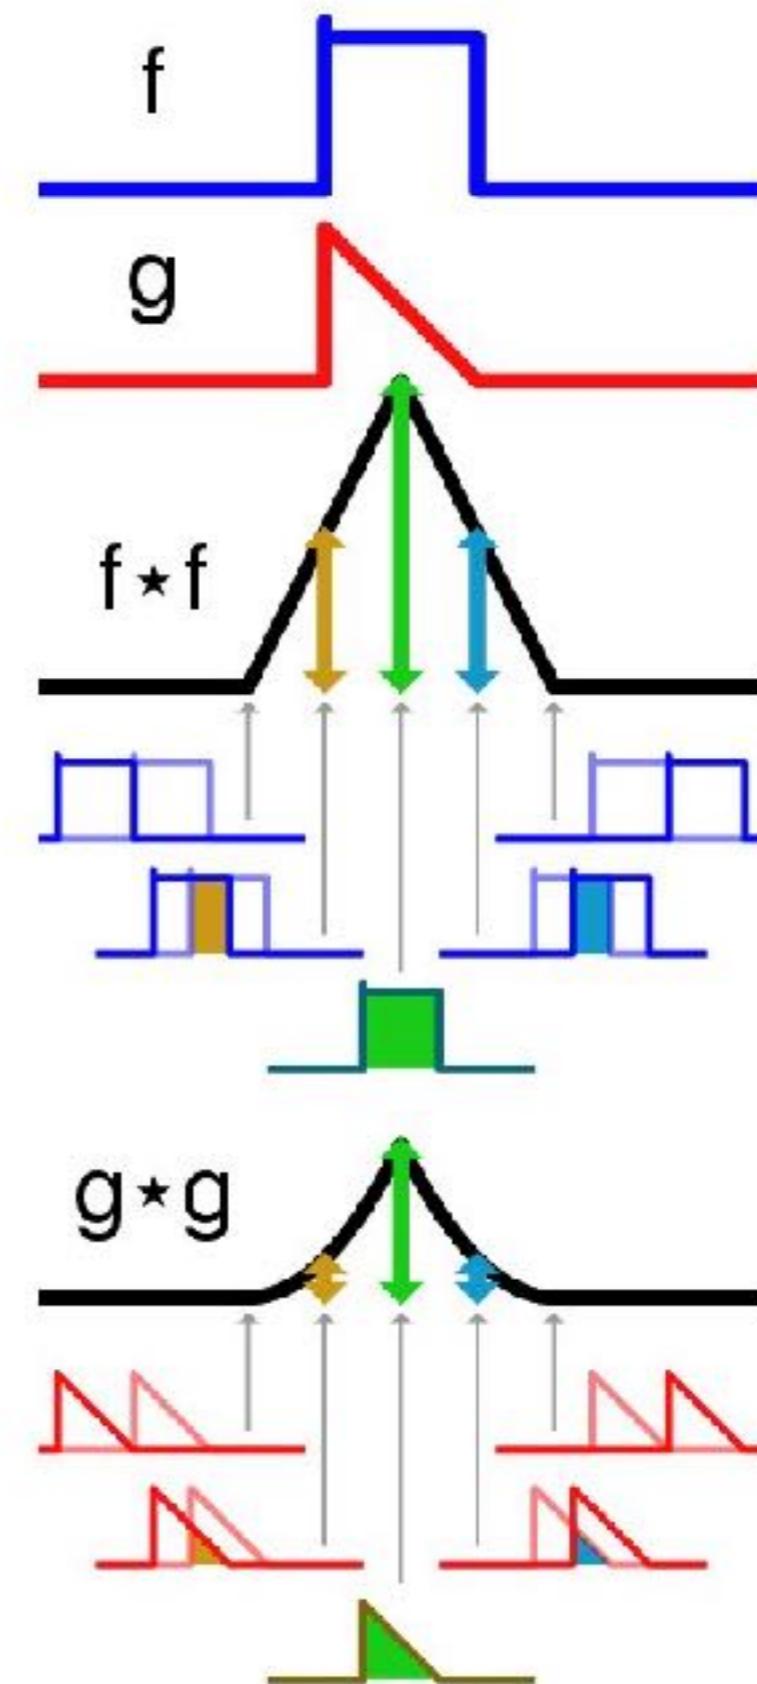
Convolution



Cross-correlation



Autocorrelation



Review Questions

1. Write down a 3×3 filter that returns a positive value if the average value of the 4-adjacent neighbors is less than the center and a negative value otherwise
2. Write down a filter that will compute the gradient in the x-direction:

$$\text{grad}_x(y, x) = \text{im}(y, x+1) - \text{im}(y, x) \text{ for each } x, y$$

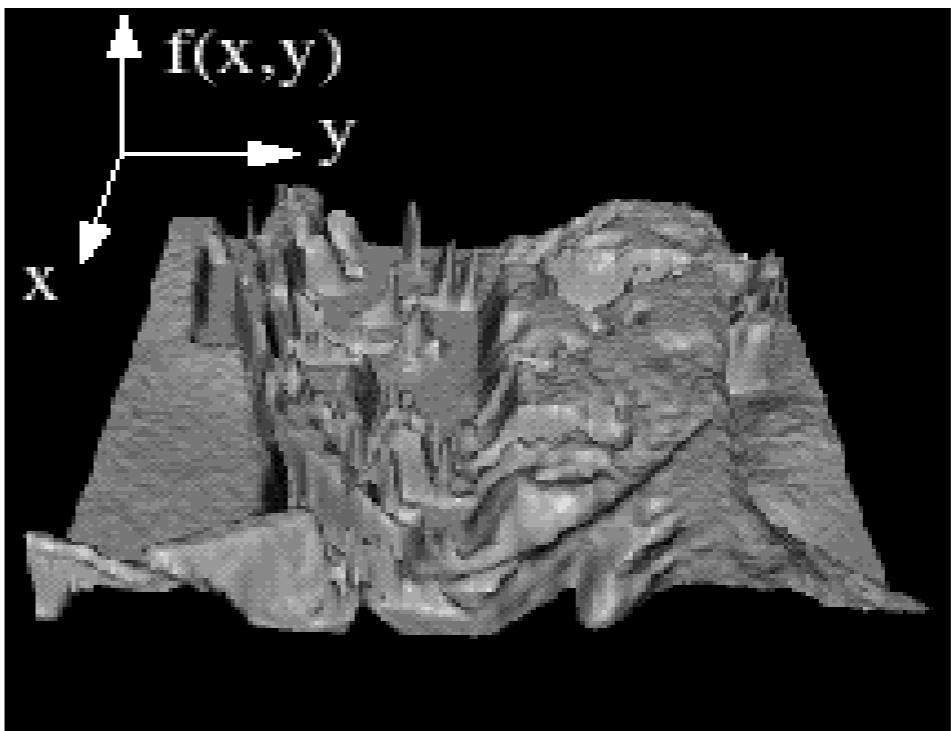
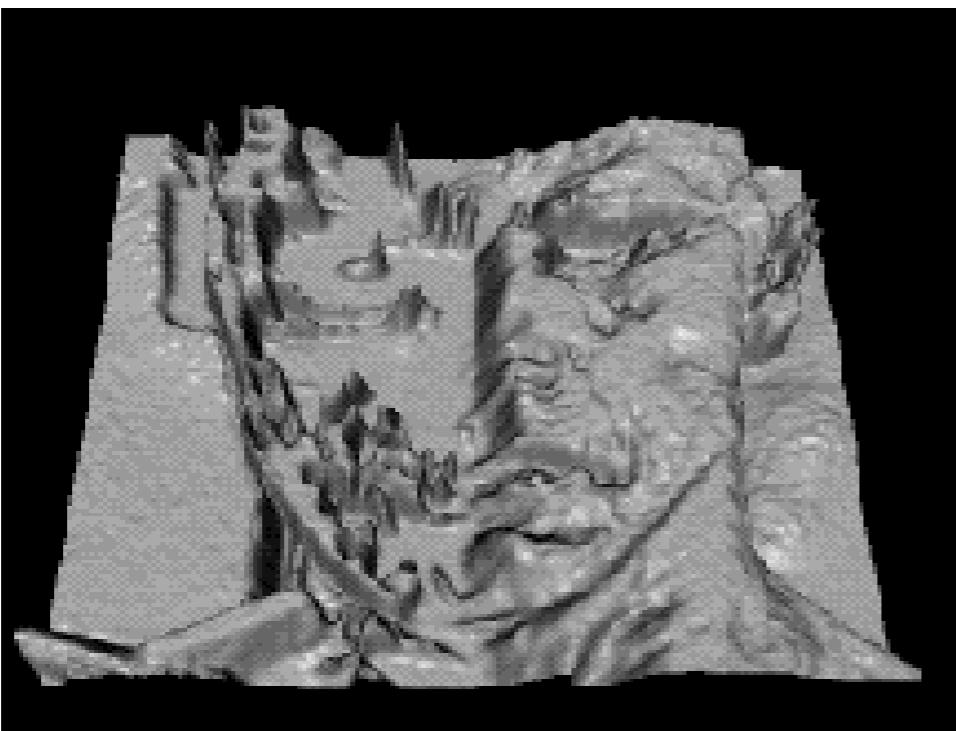
Recap

- Filtering — Convolution
- Box filter
- Gaussian filter
- Impact of scale / width of smoothing filter
- Separable filters more efficient
- Template matching
- Normalized cross-correlation
- Pyramid — Gaussian and Laplacian
- Image blending

Image Representations

- **Pixels**: great for spatial resolution, poor access to frequency
- **Pyramids/filter banks**: balance between spatial and frequency information
- Today — **Fourier transform**: great for frequency, not for spatial info

High frequency



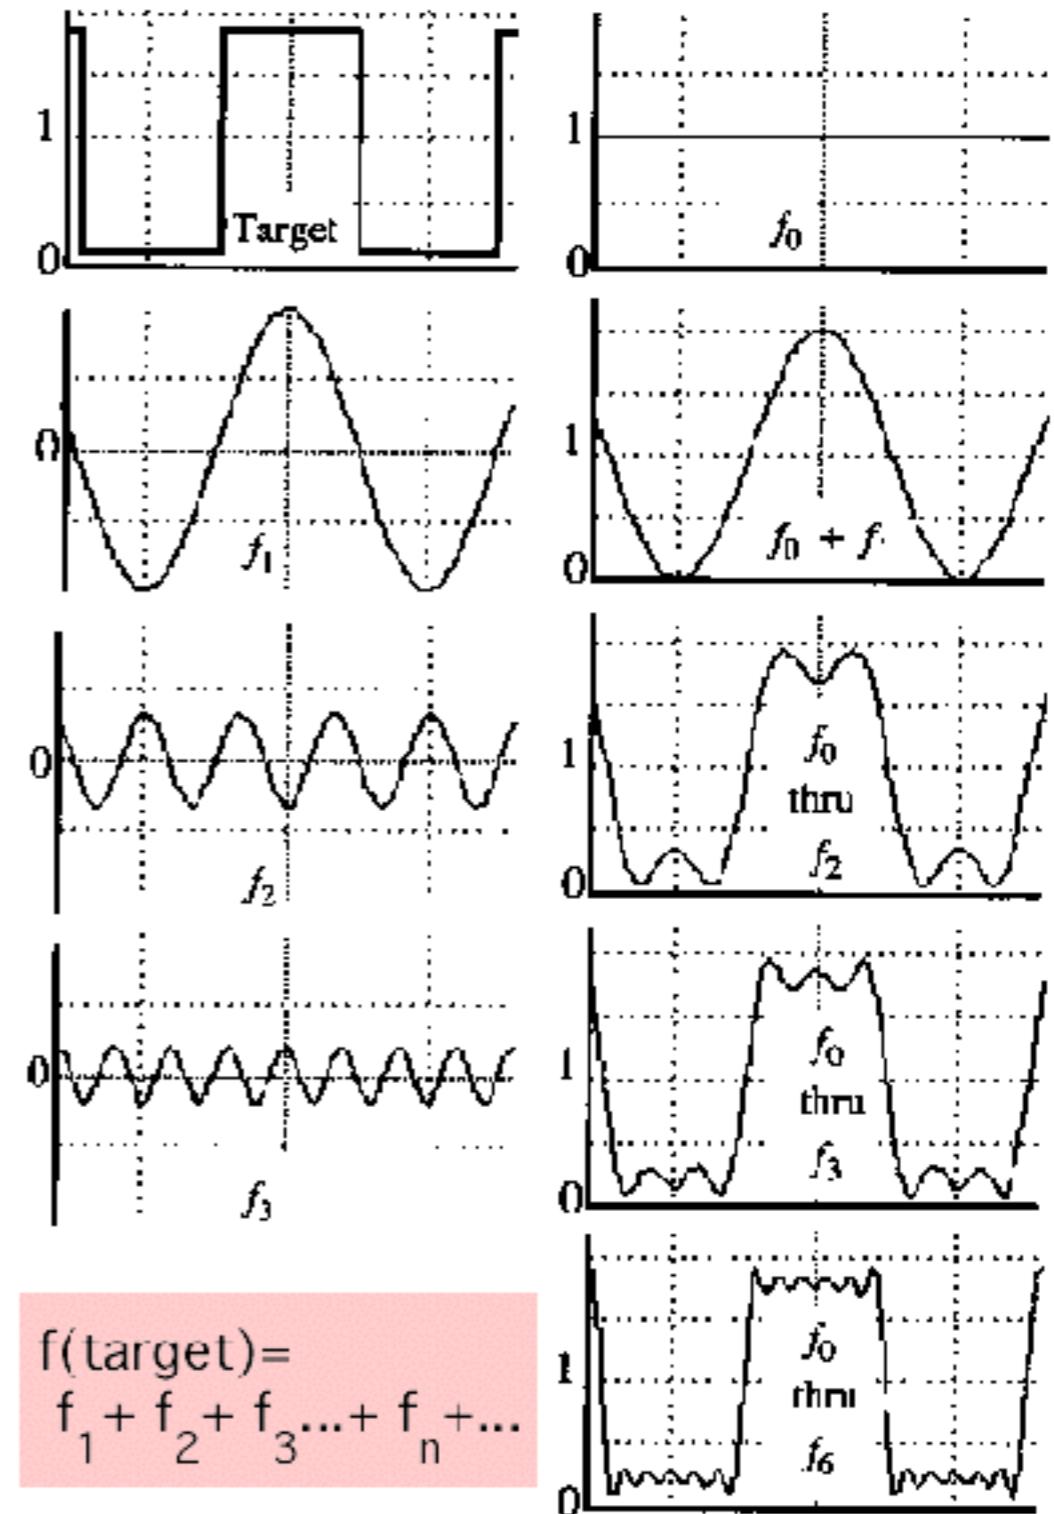
Think in terms of Frequency

A sum of sines

Our building block:

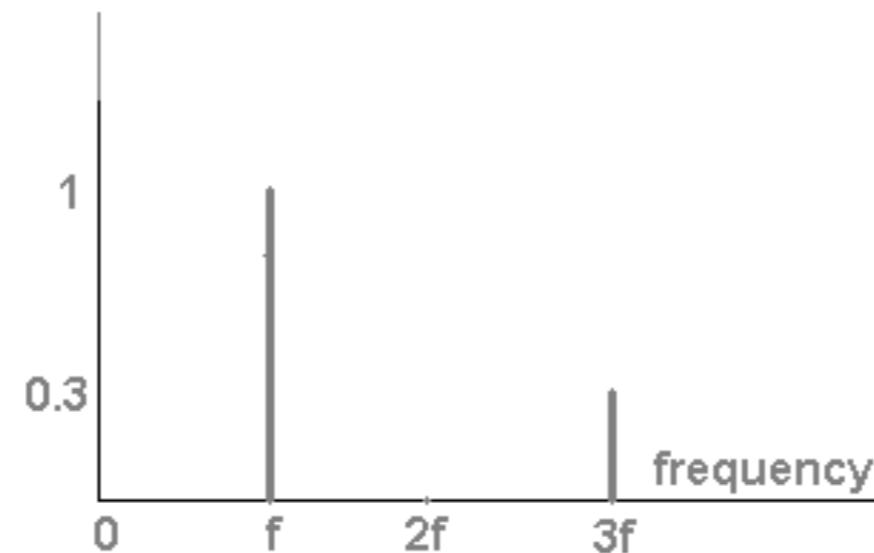
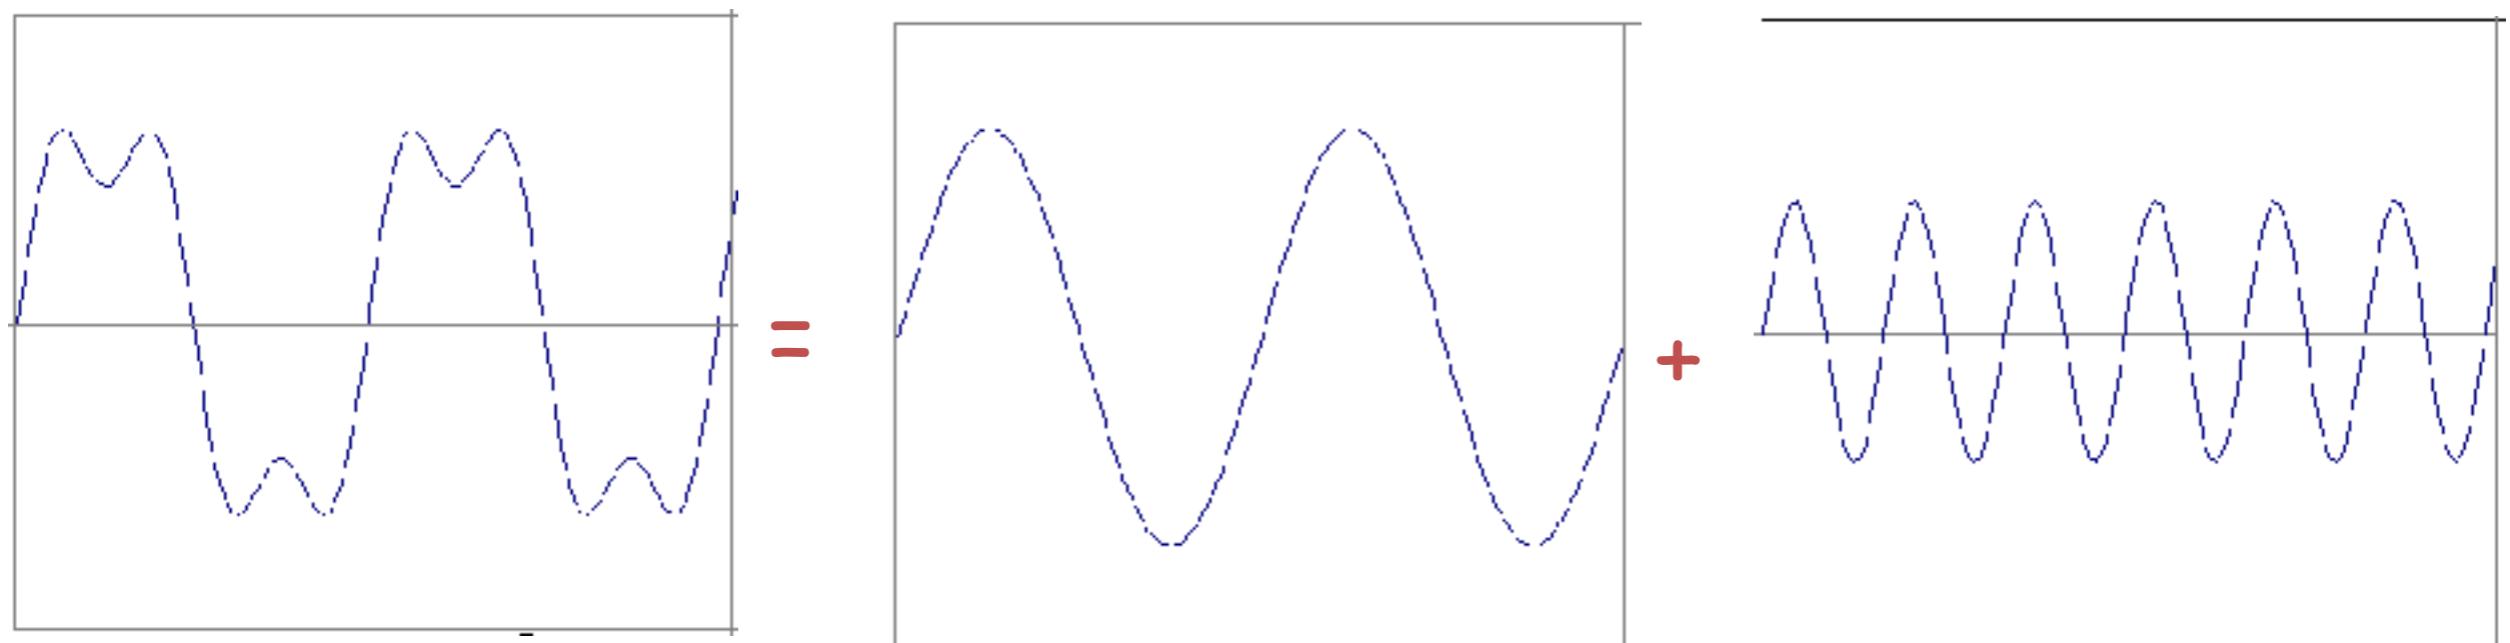
$$A \sin(\omega x + \phi)$$

Add enough of them to get any signal $f(x)$ you want!

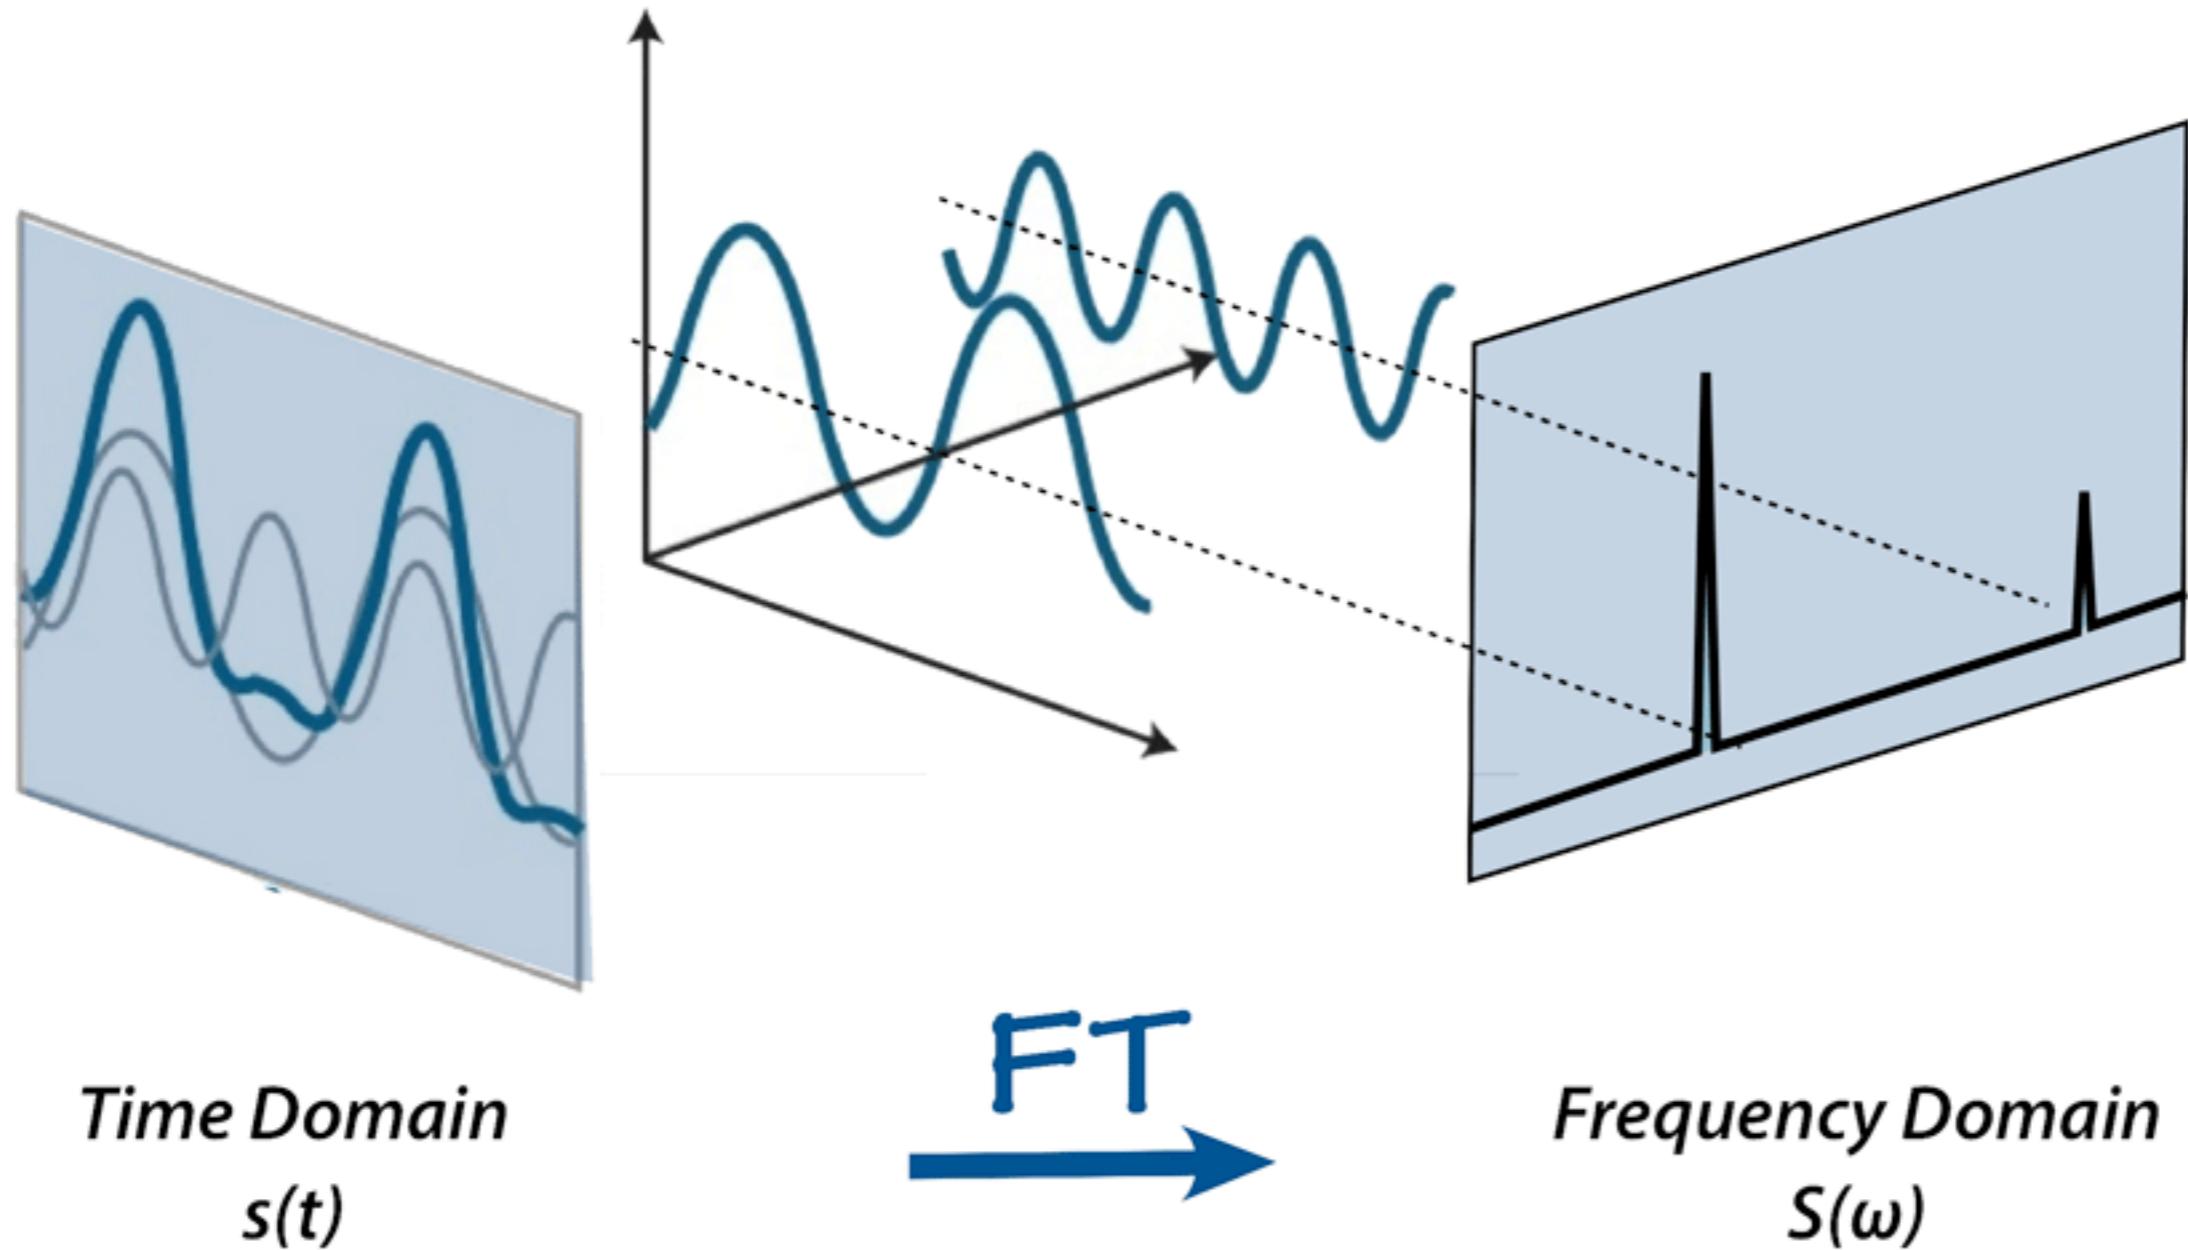


1D Frequency Spectra

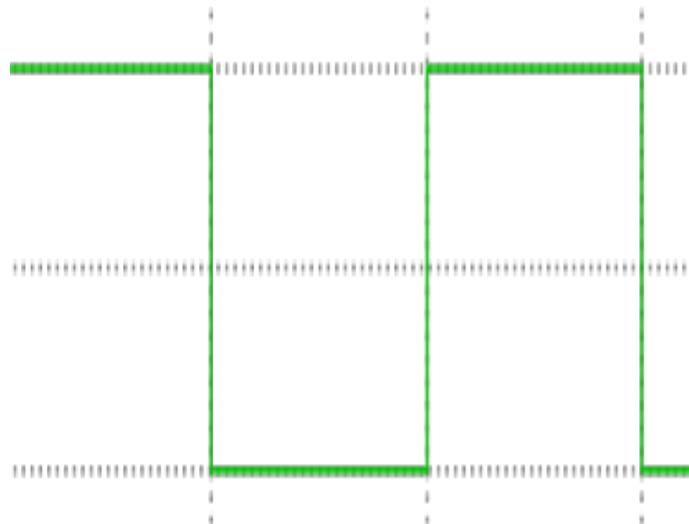
- example : $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f) t)$



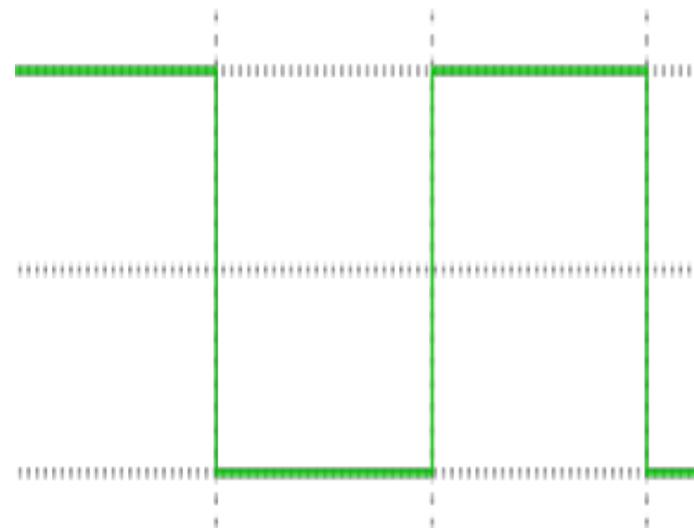
1D Frequency Spectra



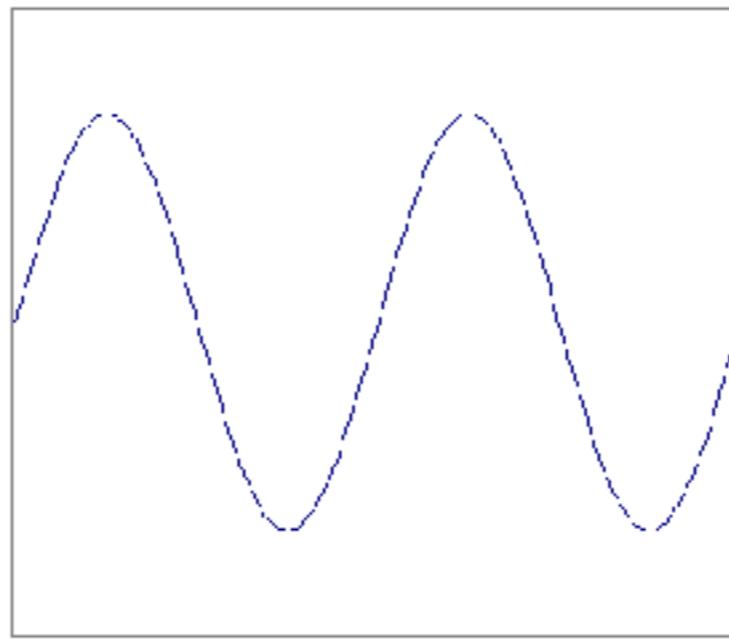
Frequency Spectra



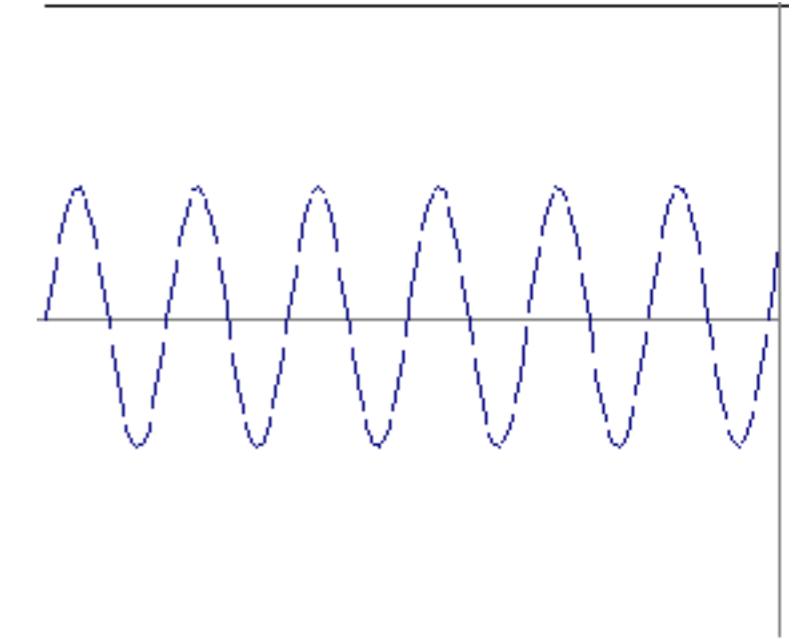
Frequency Spectra



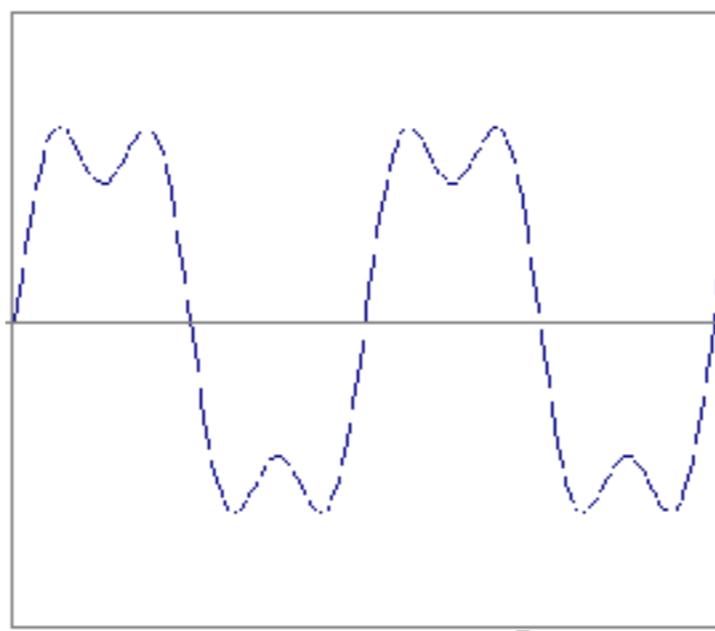
=



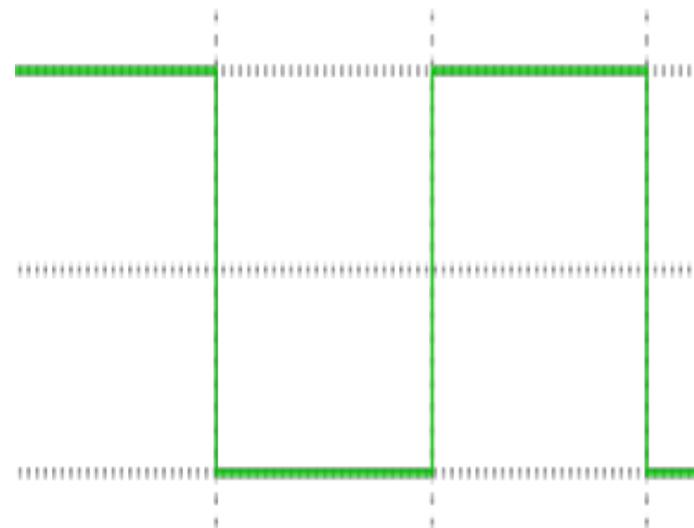
+



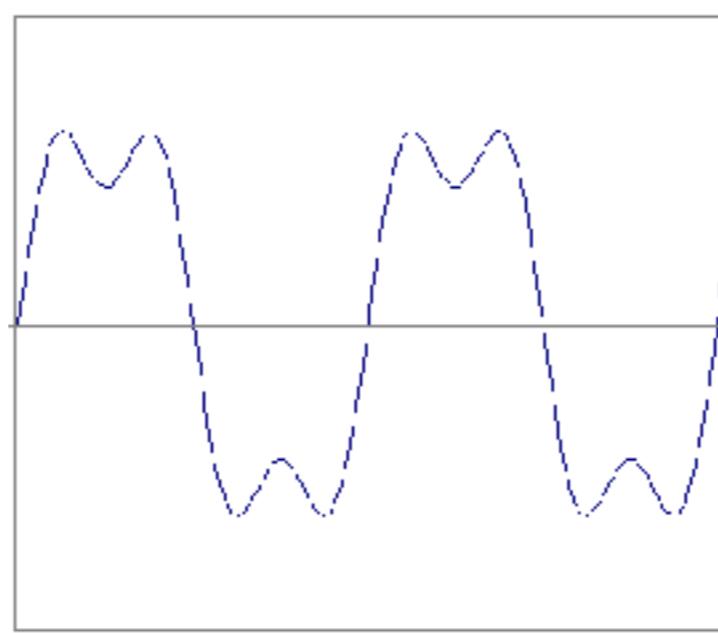
=



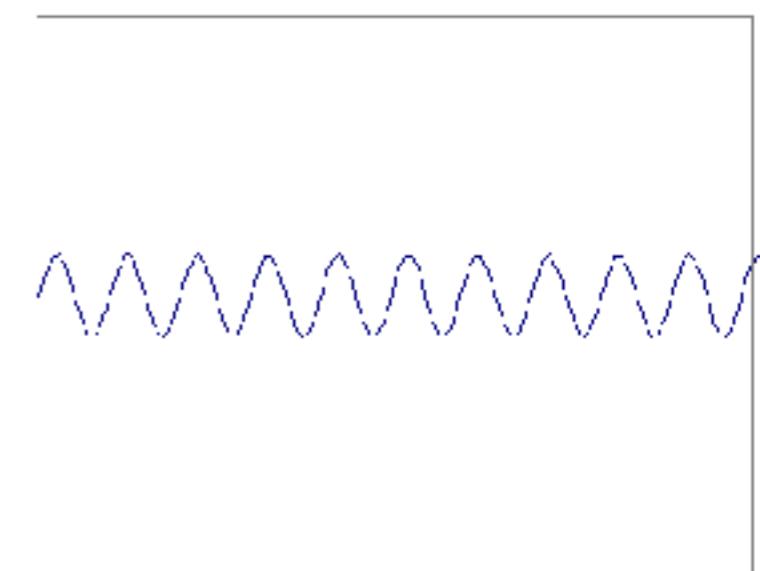
Frequency Spectra



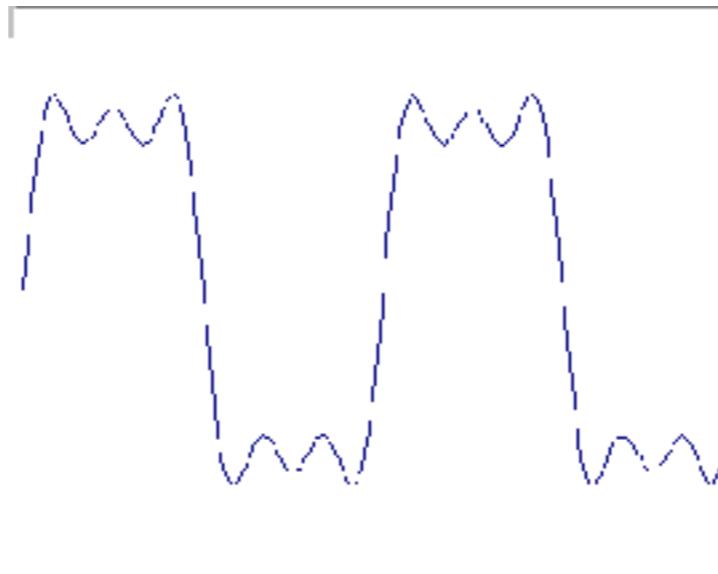
=



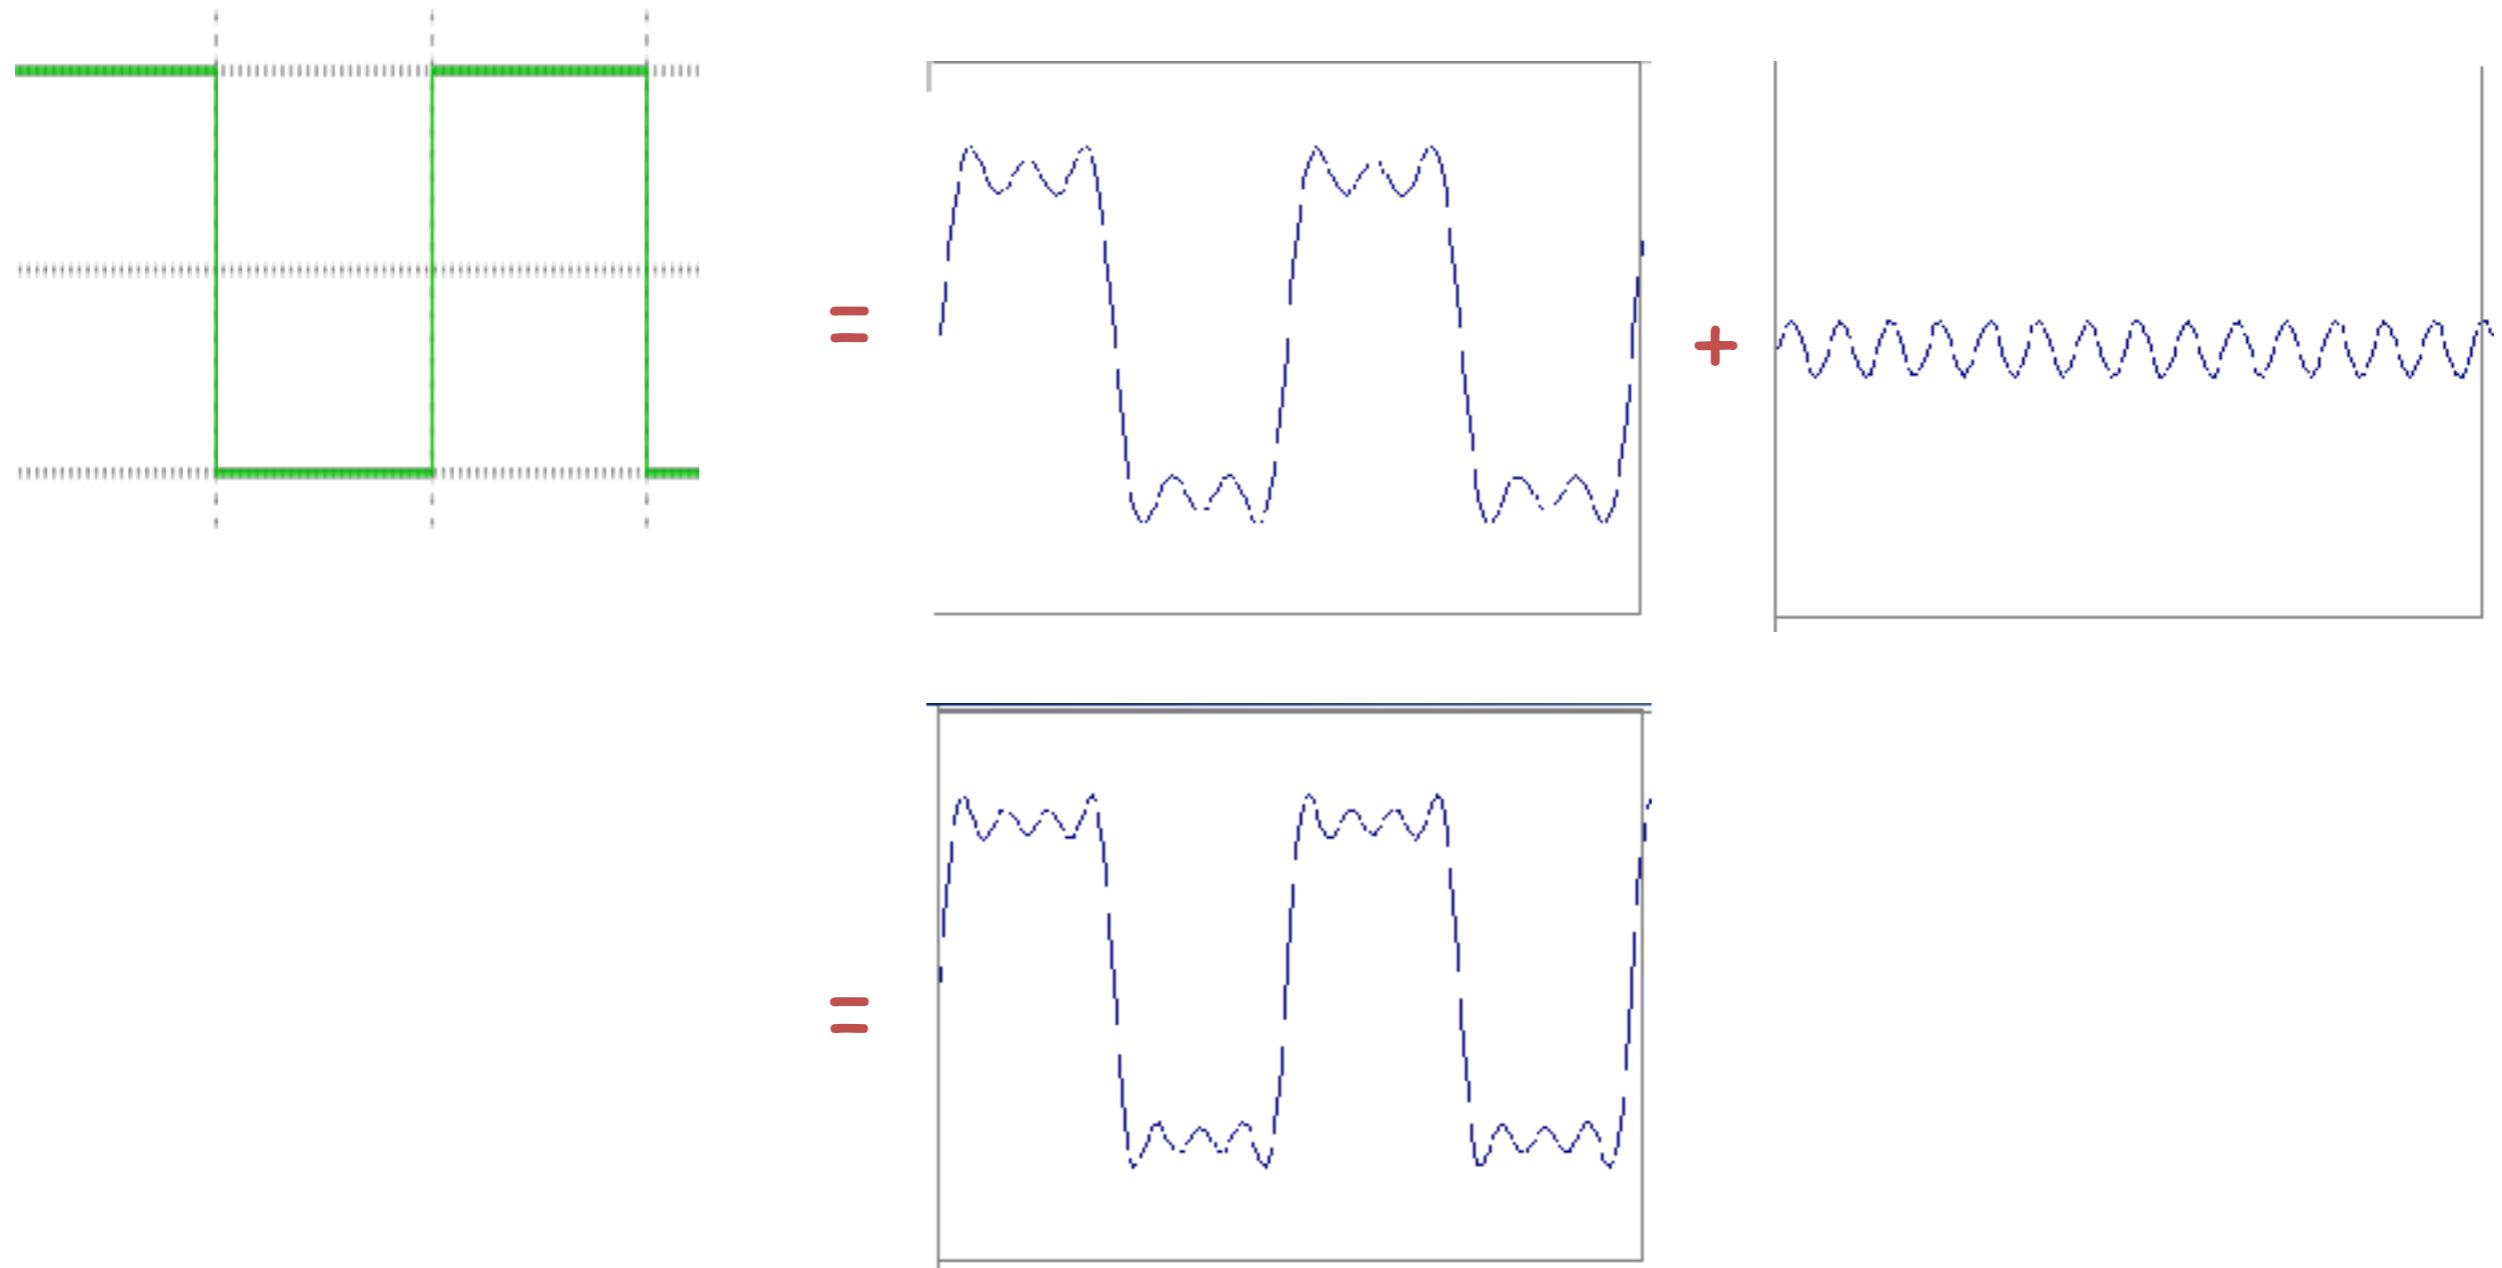
+



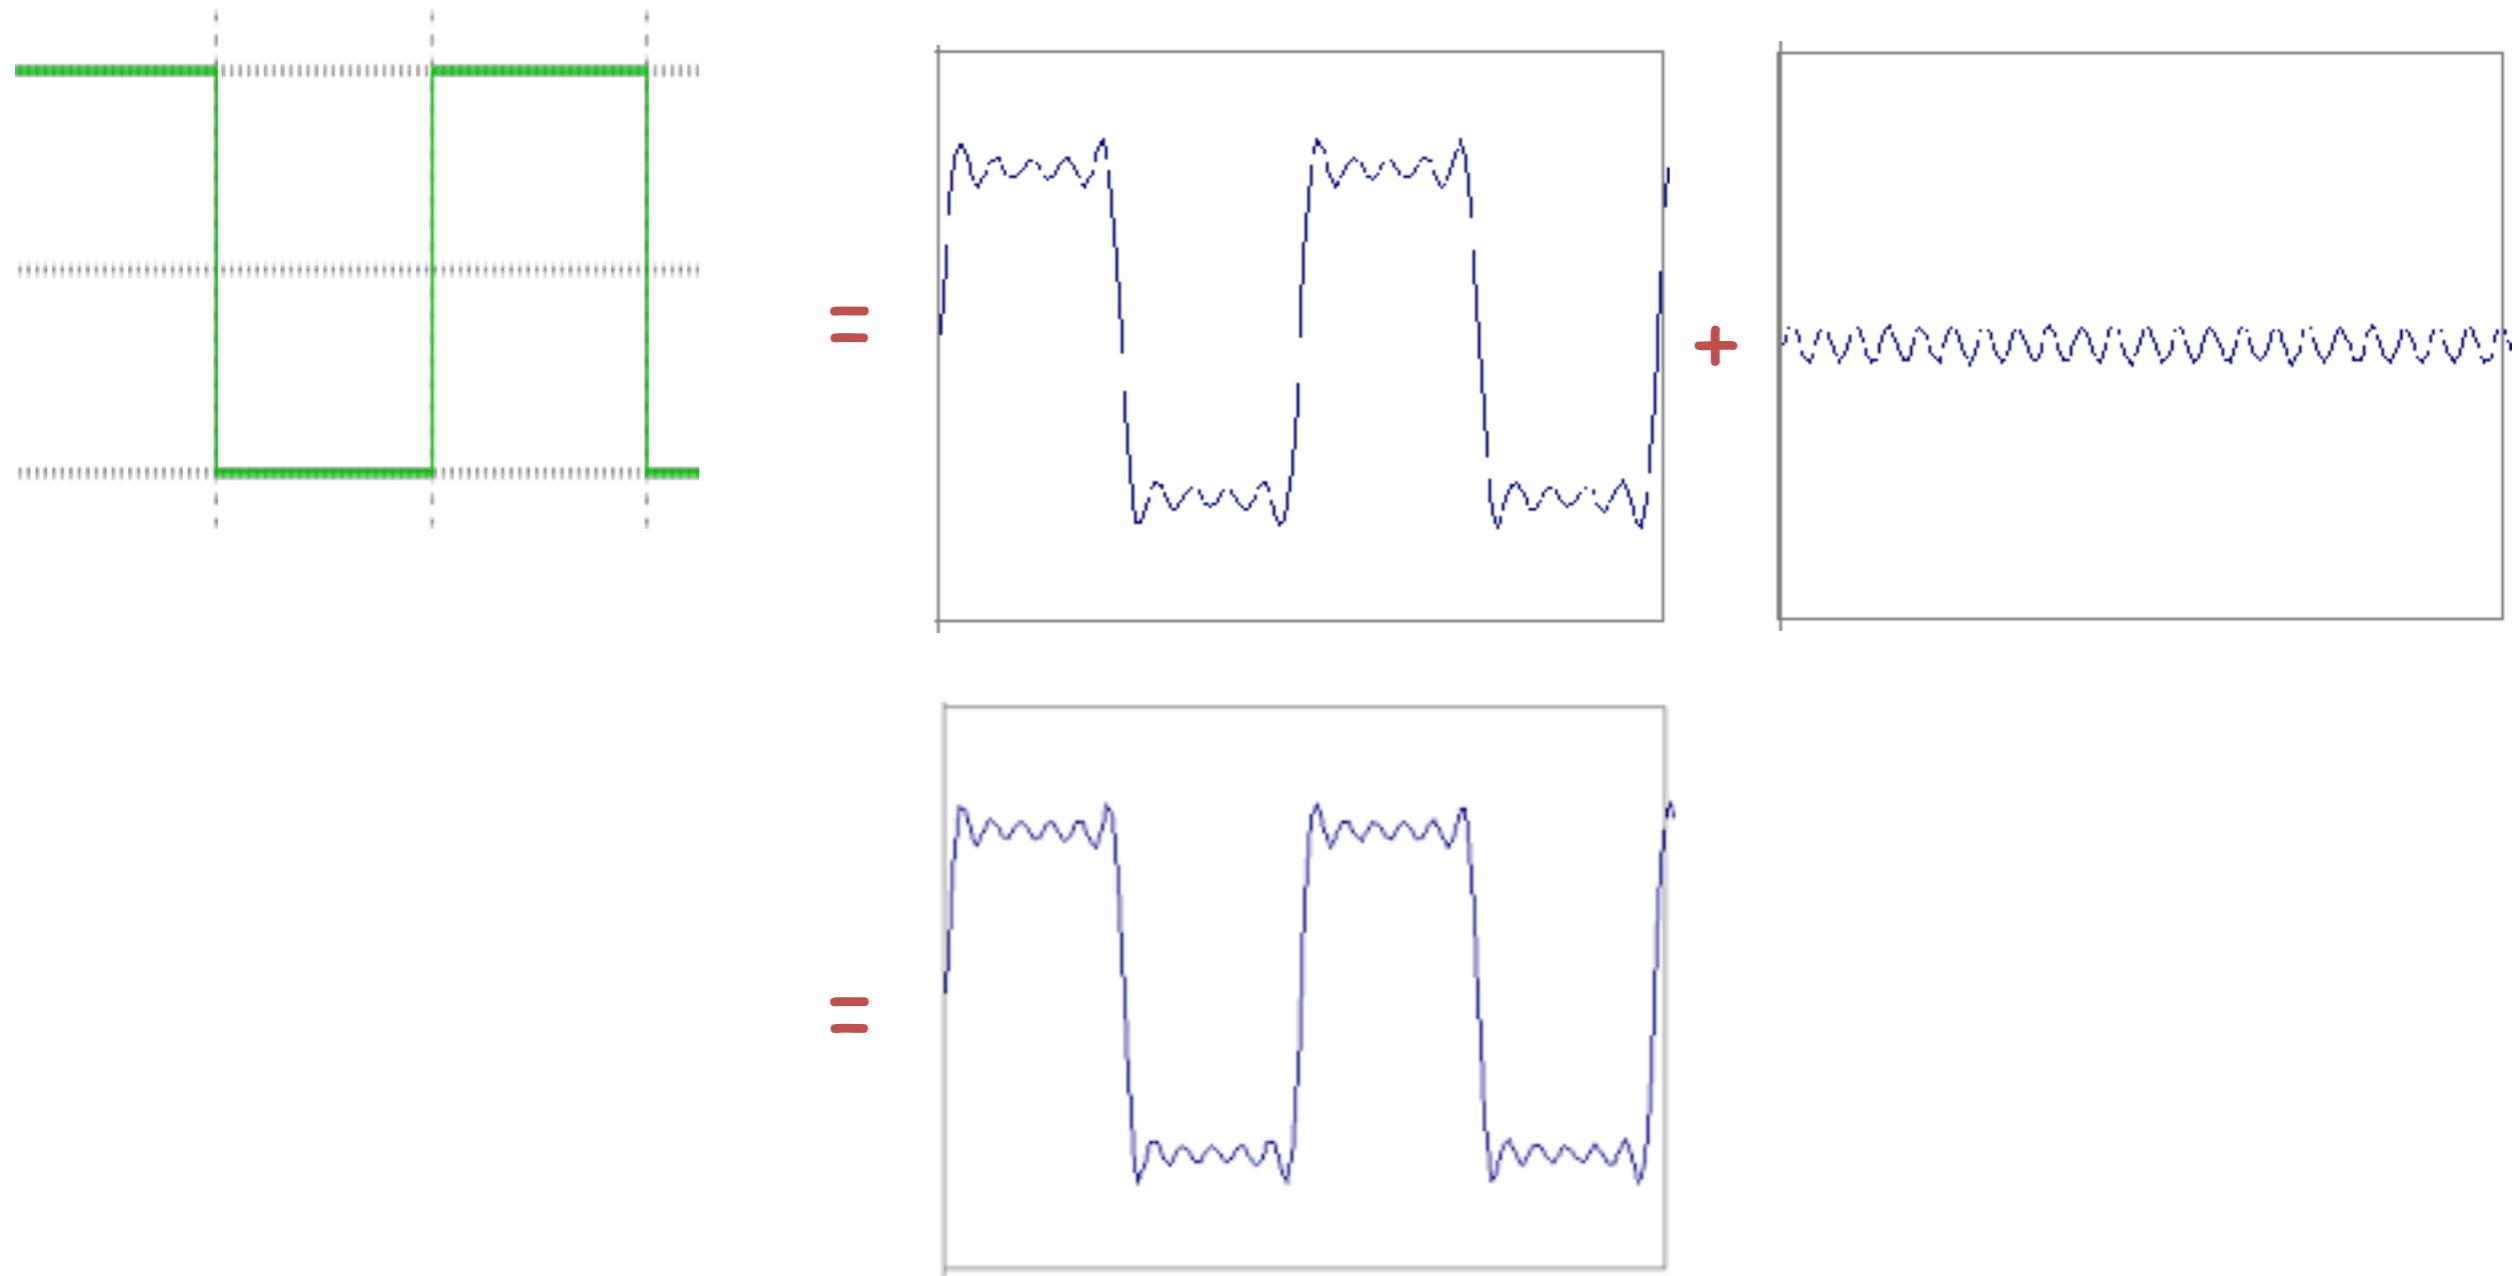
=



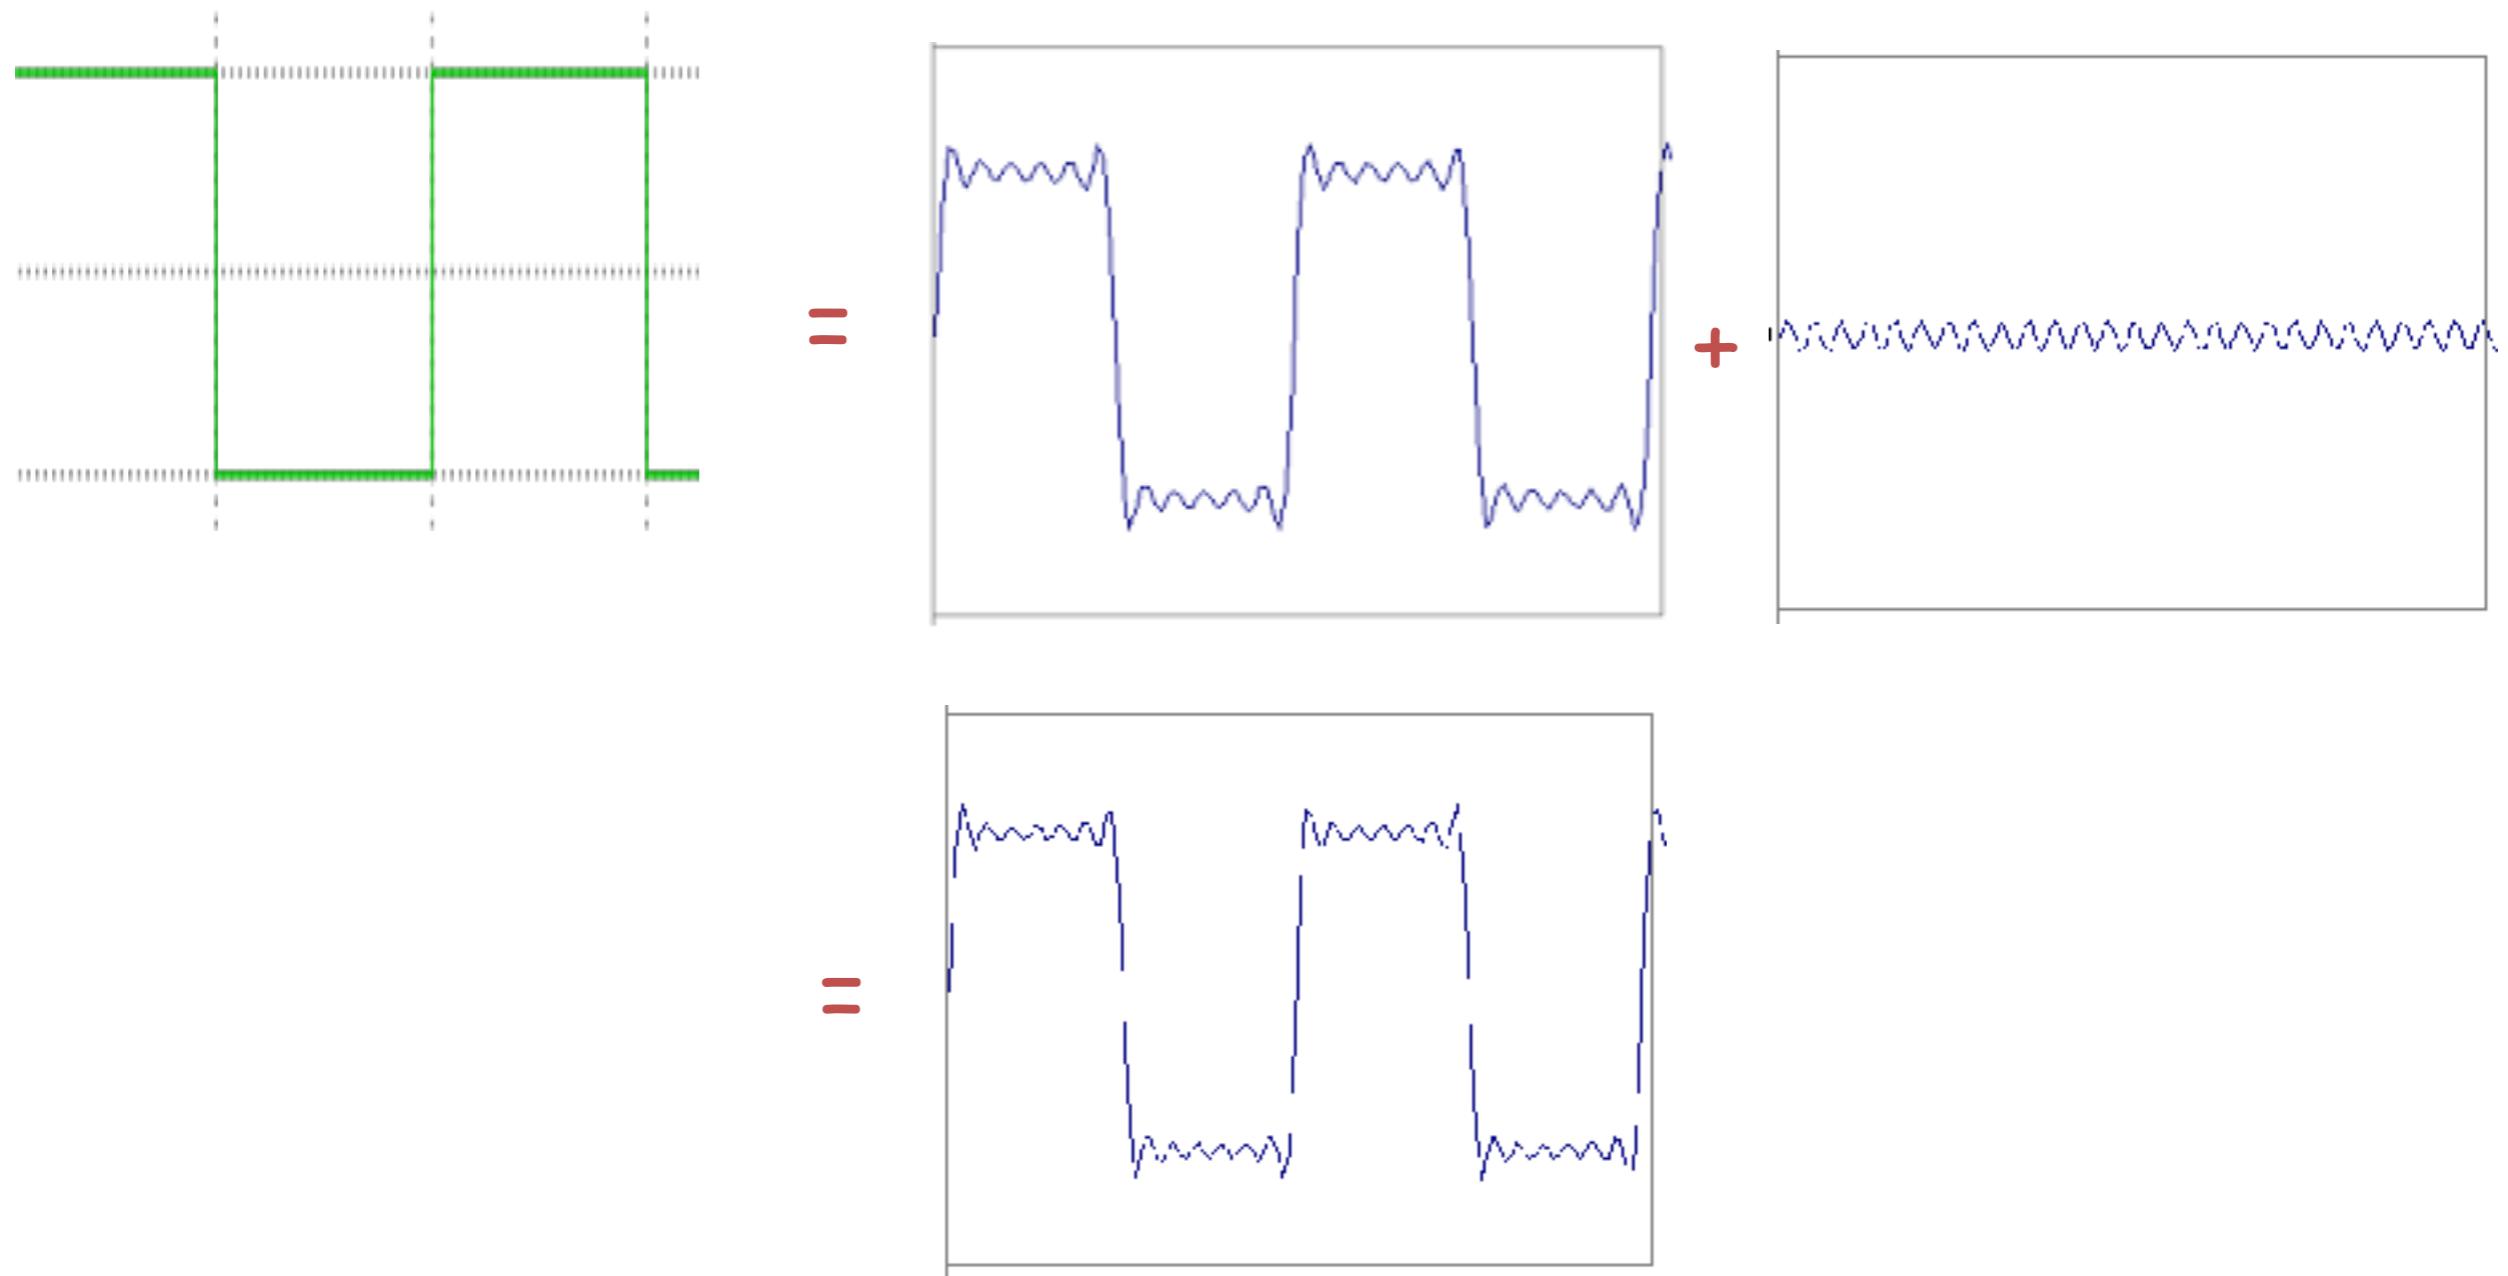
Frequency Spectra



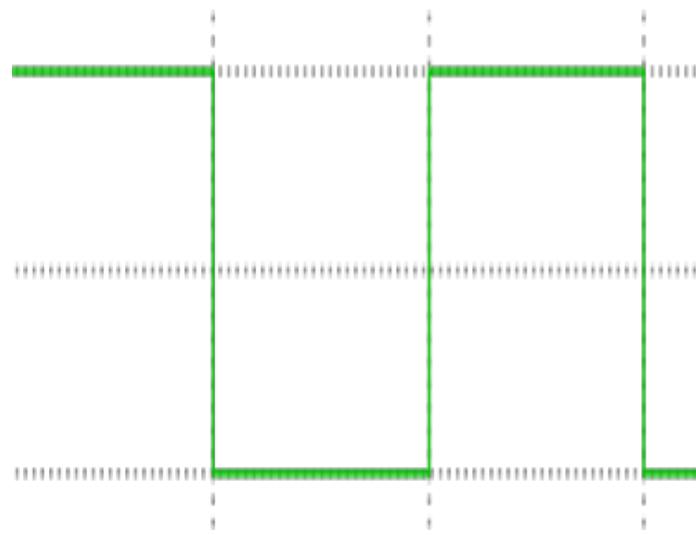
Frequency Spectra



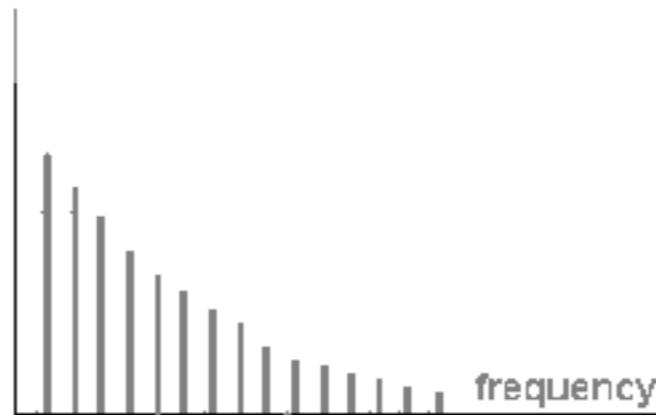
Frequency Spectra



Frequency Spectra



$$= A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$



Formal definition of 1D Fourier transform

Forward mapping Frequency \rightarrow Time

$$F(s) = \int_{-\infty}^{\infty} f(t) e^{-j2\pi st} dt$$

Backward mapping Time \rightarrow Frequency: inverse FT

$$f(t) = \int_{-\infty}^{\infty} F(s) e^{+j2\pi st} ds$$

Fourier series (**Euler's identity**)

$$f(t) = \sum_{n=-\infty}^{\infty} c_n e^{j \frac{2\pi n}{T} t}$$

$$f(t) = a_0 + \sum_{n=1}^{\infty} a_n \sin\left(j2\pi \frac{n}{T} t\right) + \sum_{n=1}^{\infty} b_n \cos\left(j2\pi \frac{n}{T} t\right)$$

Digital, Continuous and Fast Fourier Transform

DFT

CFT

FFT

For now, you only need to know the big picture:

	Continuous Time	Discrete Time
Fourier Series	$x(t) = \sum_{k=-\infty}^{\infty} a_k e^{jk\omega_0 t}$ <p>continuous and periodic in time</p> $a_k = \frac{1}{T} \int_T x(t) e^{-jk\omega_0 t} dt$ <p>discrete and aperiodic in frequency</p>	$x[n] = \sum_{k=<N>} a_k e^{jk\frac{2\pi}{N} n}$ <p>discrete and periodic in time</p> $a_k = \frac{1}{N} \sum_{n=<N>} x[n] e^{-jk\frac{2\pi}{N} n}$ <p>discrete and periodic in frequency</p>
Fourier Transform	$x(t) = \frac{1}{2\pi} \int_{-\infty}^{\infty} X(j\omega) e^{j\omega t} d\omega$ <p>continuous and aperiodic in time</p> $X(j\omega) = \int_{-\infty}^{\infty} x(t) e^{-j\omega t} dt$ <p>continuous and aperiodic in frequency</p>	$x[n] = \frac{1}{2\pi} \int_{2\pi} X(e^{j\omega}) e^{j\omega n} d\omega$ <p>discrete and aperiodic in time</p> $X(e^{j\omega}) = \sum_{n=-\infty}^{\infty} x[n] e^{-j\omega n}$ <p>continuous and periodic in frequency</p>

Take a signal processing course if you ever want to work with images!

Fourier Transform

- Fourier transform stores the magnitude and phase at each frequency
 - Magnitude encodes how much signal there is at a particular frequency
 - Phase encodes spatial information (indirectly)
 - For mathematical convenience, this is often notated in terms of real and complex numbers

$$\text{Amplitude: } A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

$$\text{Phase: } \phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$F[g * h] = F[g]F[h]$$

- The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms

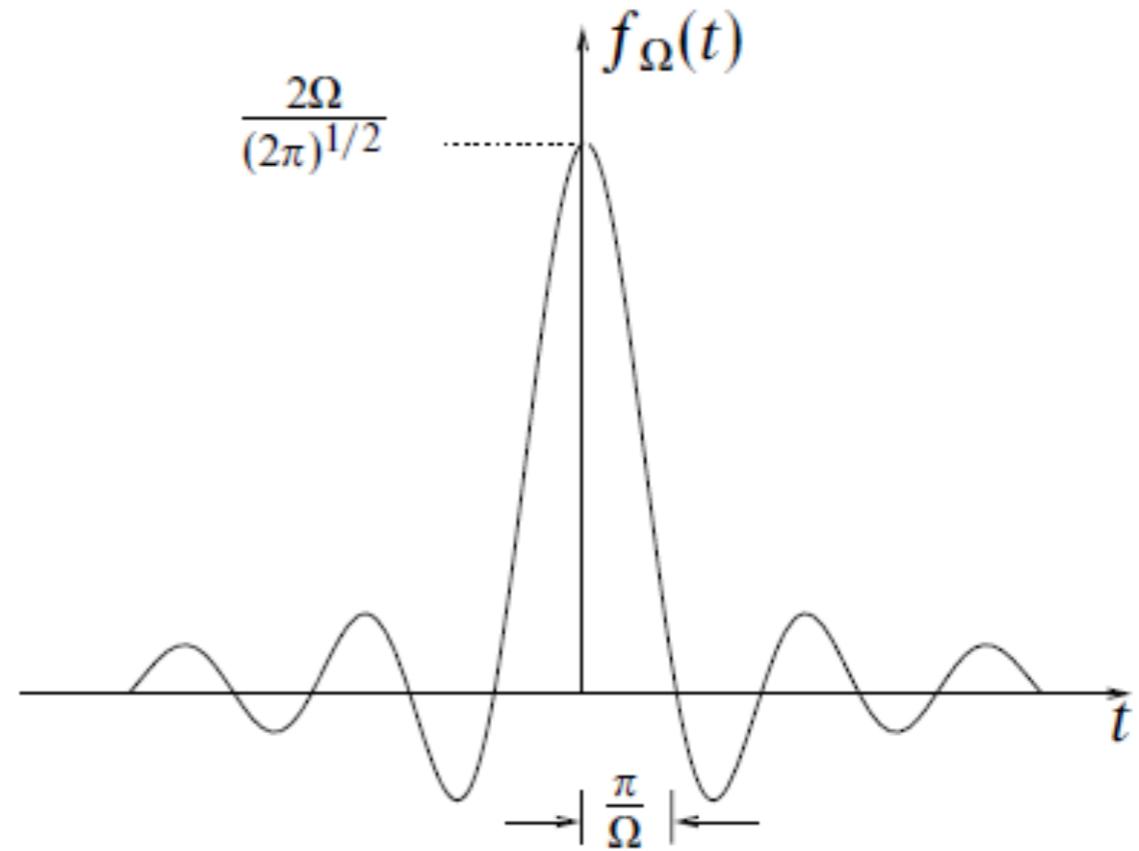
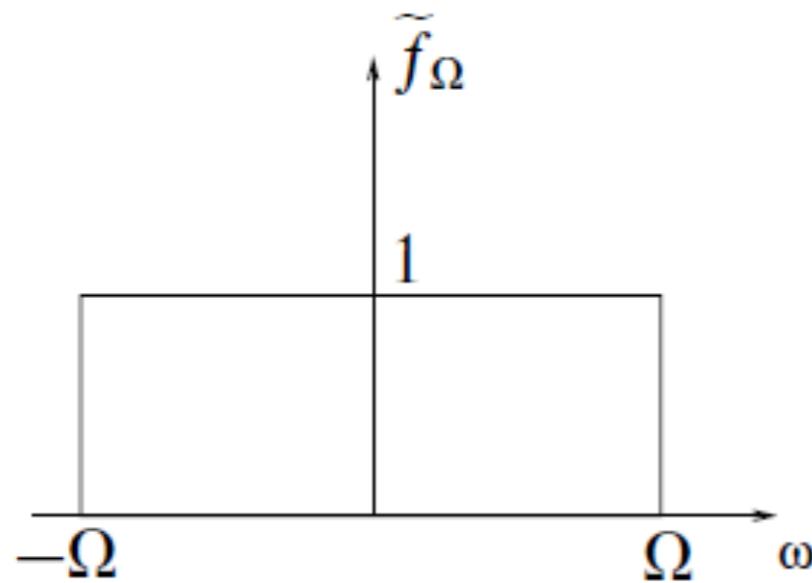
$$F^{-1}[gh] = F^{-1}[g] * F^{-1}[h]$$

- *Convolution in spatial domain is equivalent to multiplication in frequency domain!*

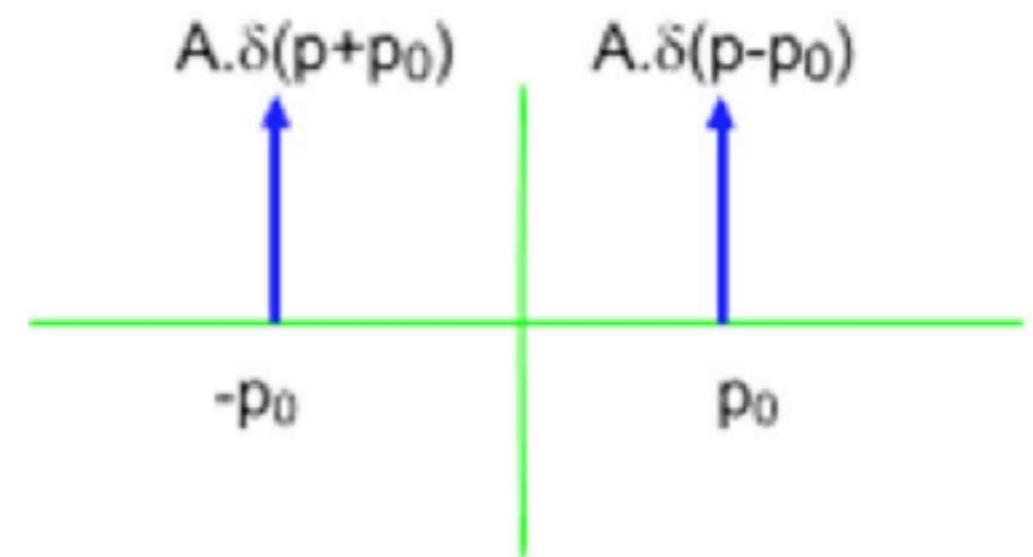
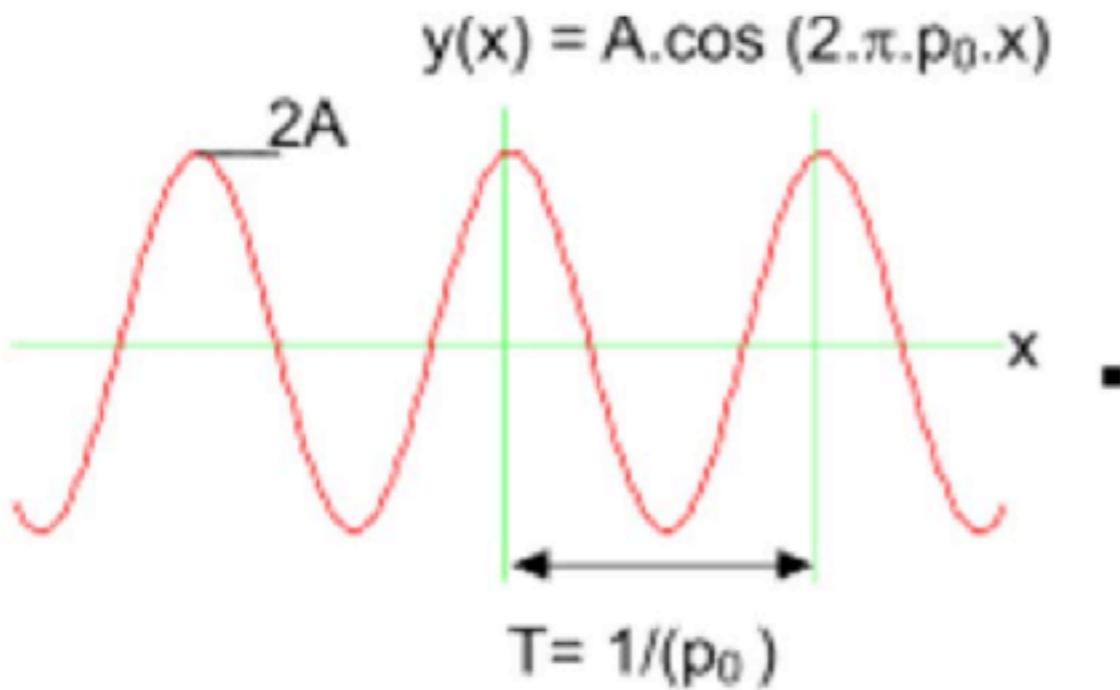
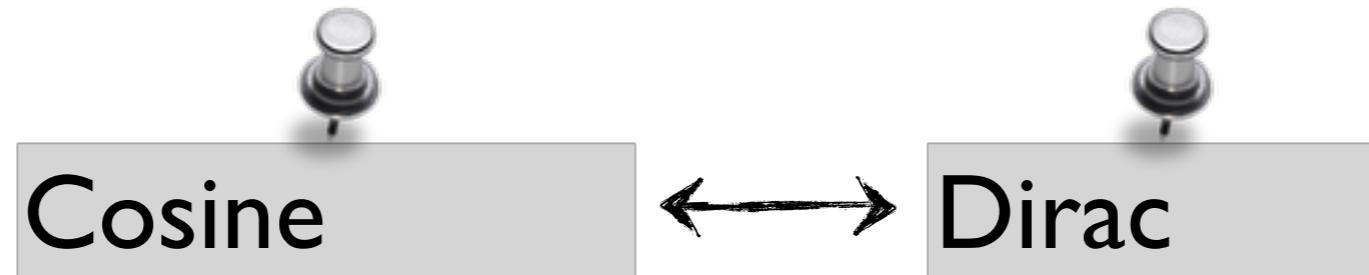
Properties of Fourier Transforms

- Linearity
- Fourier transform of a real signal is symmetric about the origin
- The energy of the signal is the same as the energy of its Fourier transform

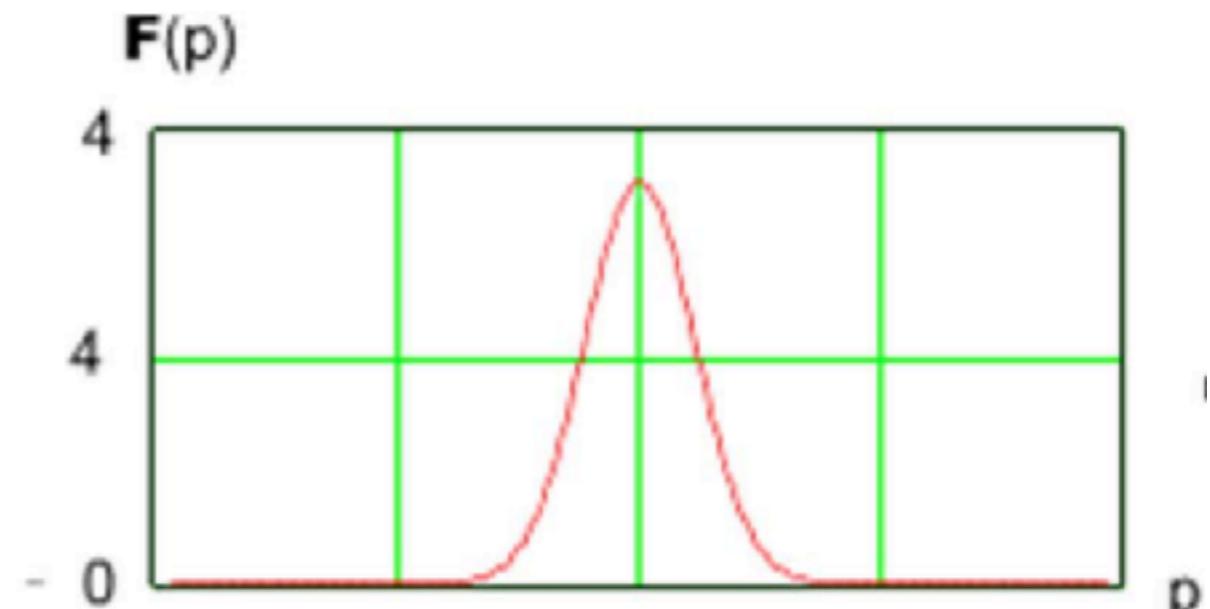
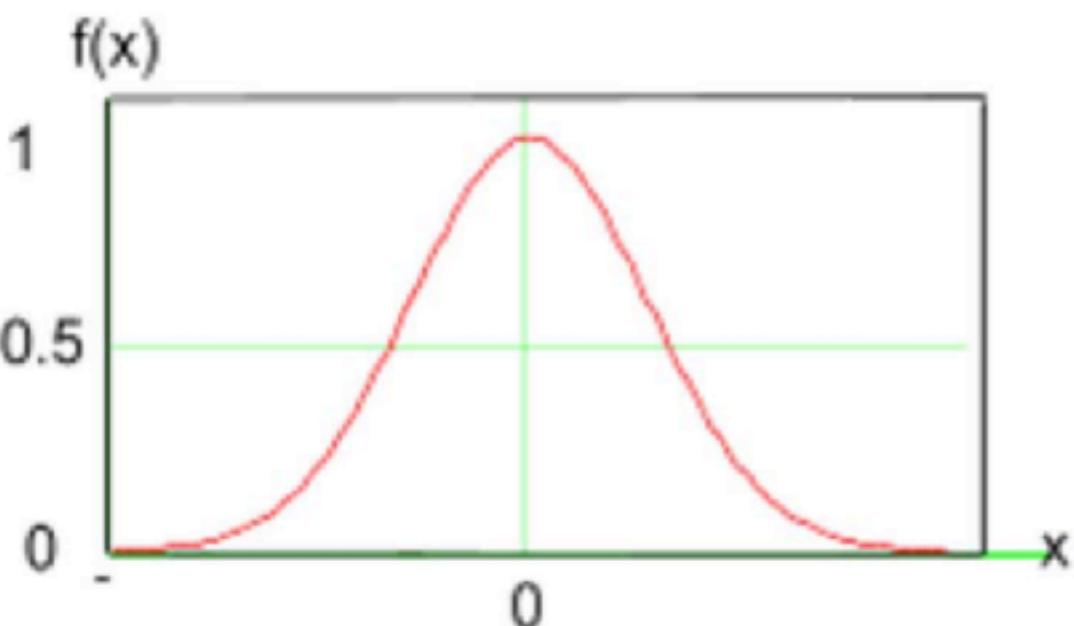
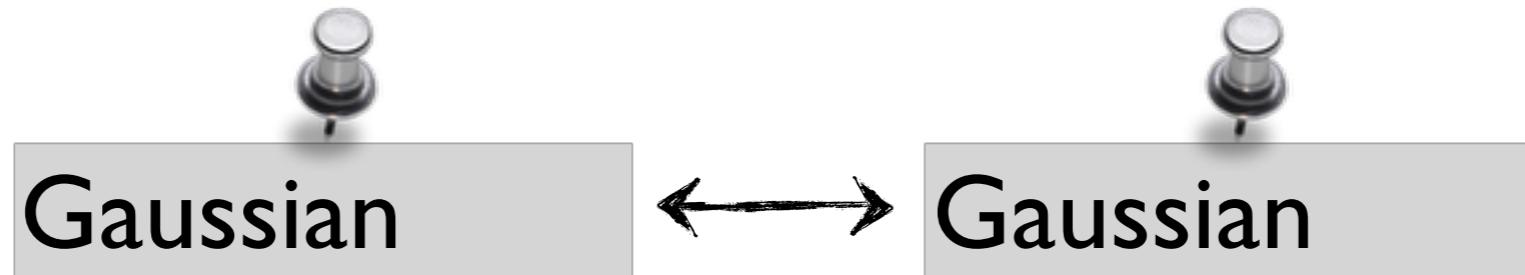
Common Transform pairs



Common Transform pairs

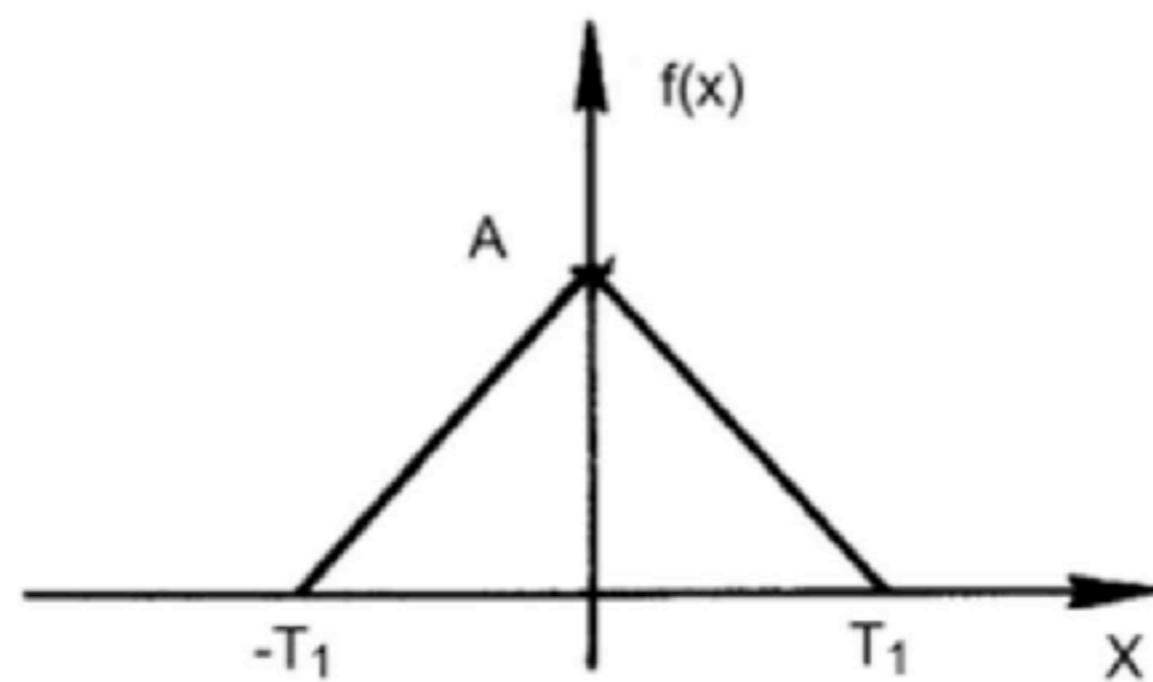


Common Transform pairs



Your turn

What is the FT of a triangle function?

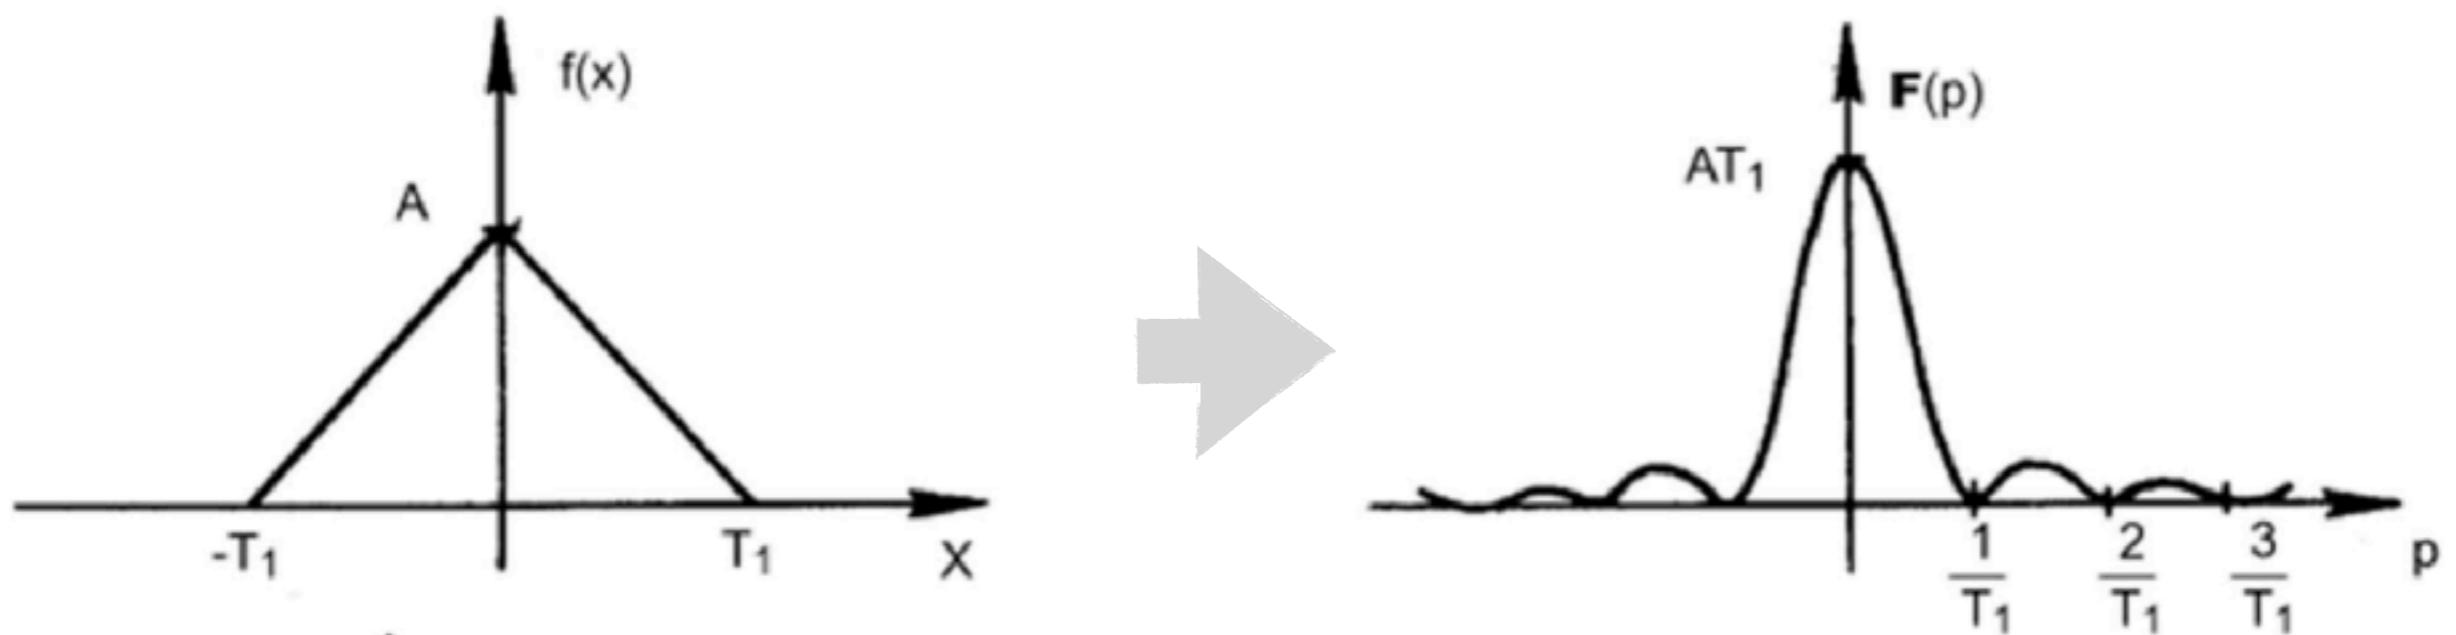


Hint: how do you get triangle function from the functions shown so far?

Convolution and Multiplication

Triangle = box convolved with box

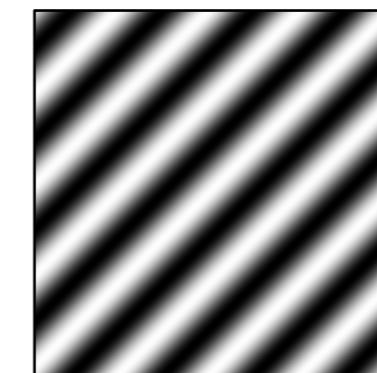
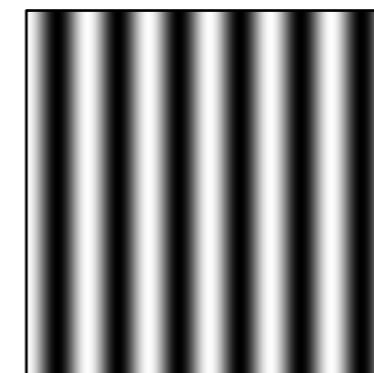
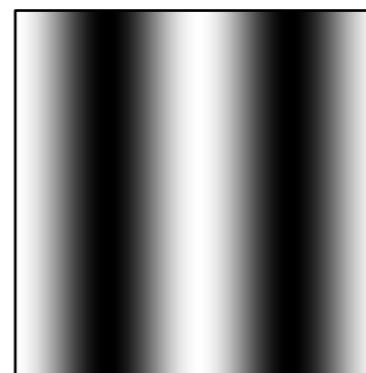
So its FT is sinc \times sinc



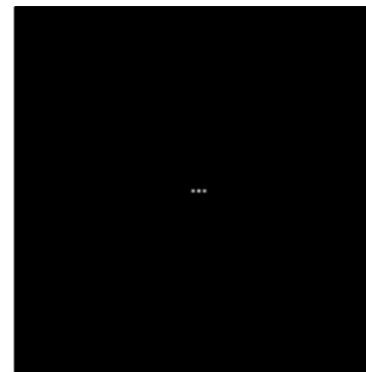
1D → 2D

2D Fourier analysis in images

Intensity Image
(spatial)



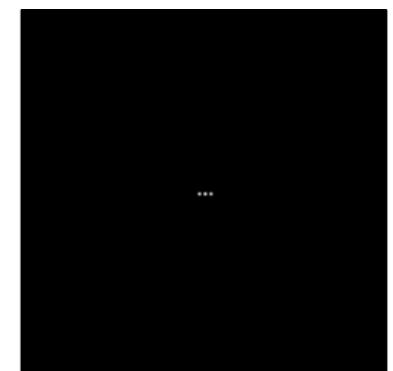
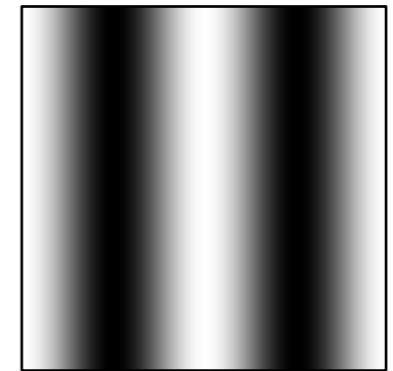
Fourier Image
(Frequency)



2D Fourier transform Plot

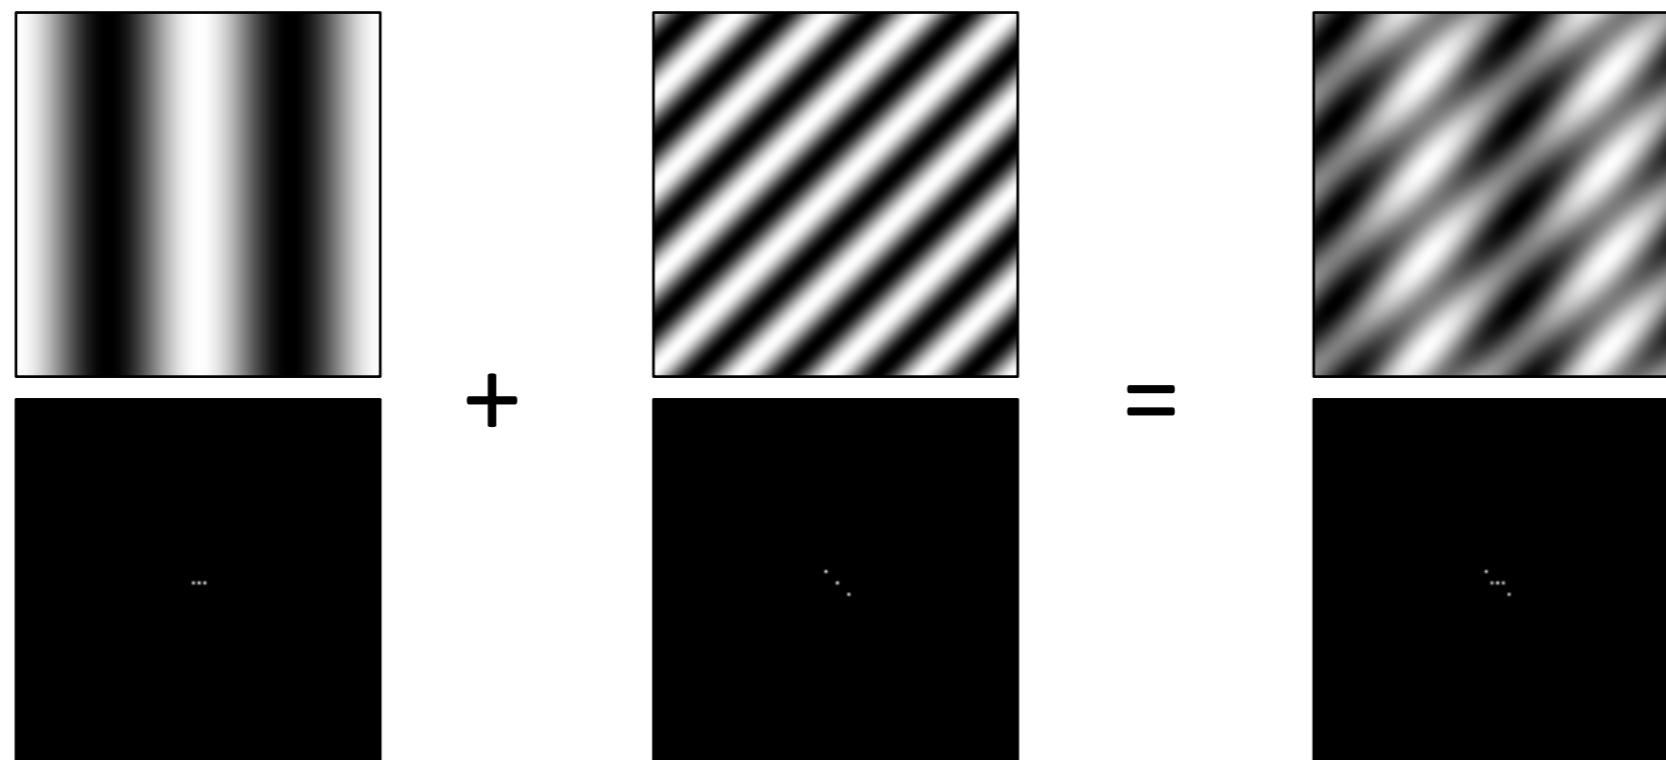
```
import cv2
import numpy as np
from matplotlib import pyplot as plt

img = cv2.imread('messi5.jpg',0)
f = np.fft.fft2(img)
fshift = np.fft.fftshift(f)
magnitude_spectrum = 20*np.log(np.abs(fshift))
```



- Only plot magnitude for visualization
 - Need phase for reconstruction
- Image center —> zero frequency
 - fftshift
 - log scale

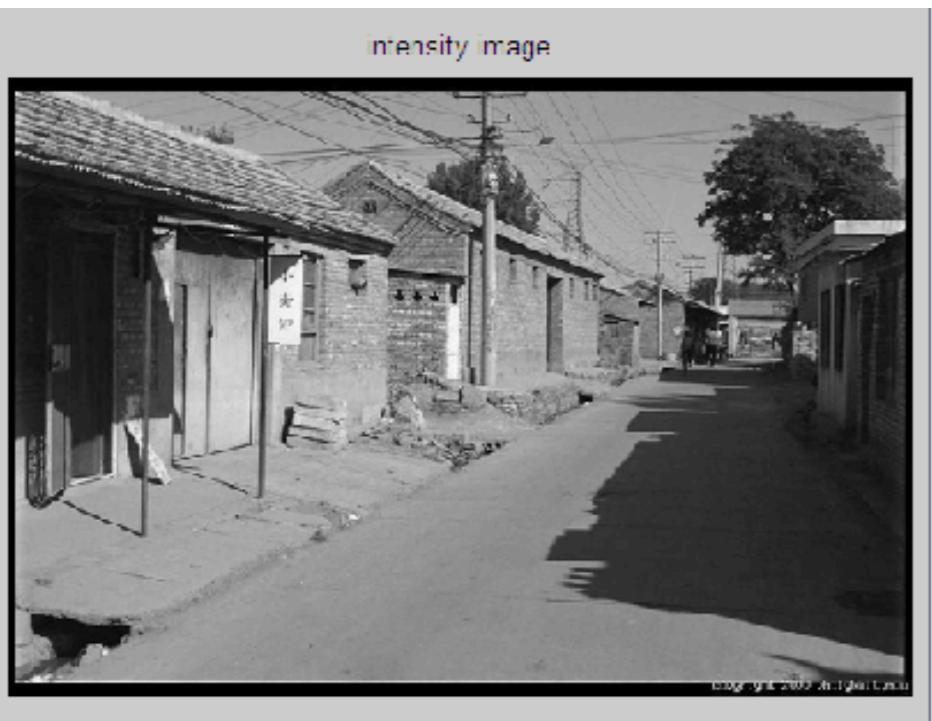
Signals can be (de)composed



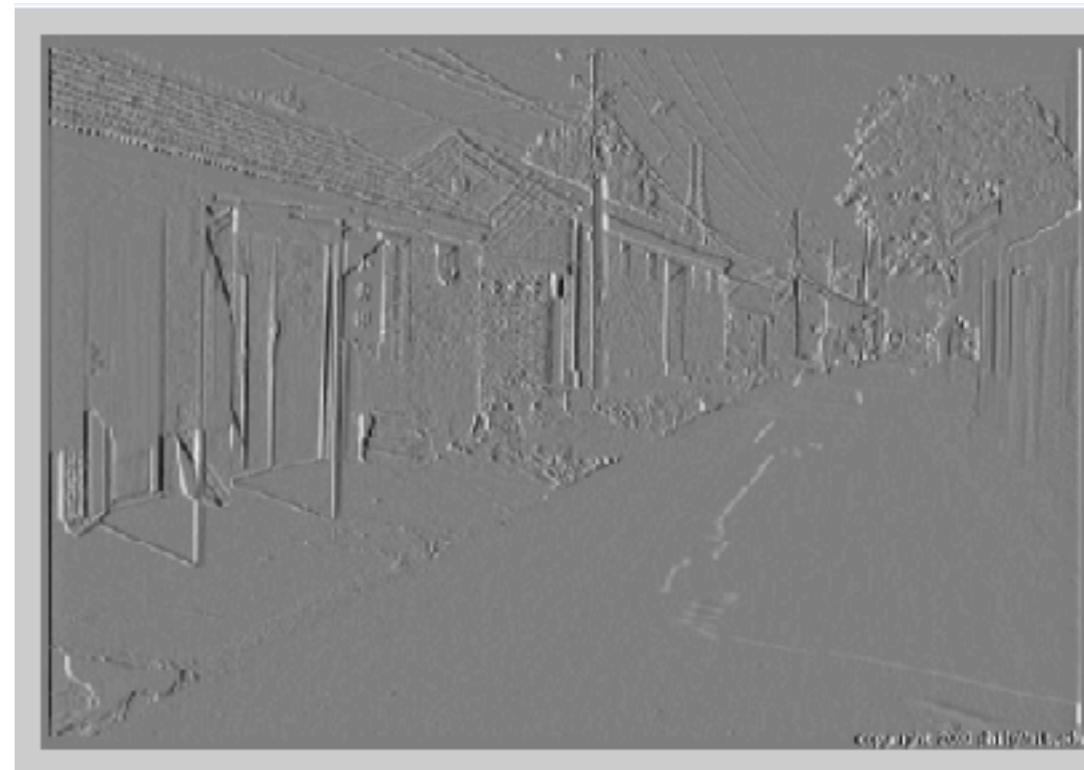
<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering>
More: <http://www.cs.unm.edu/~brayer/vision/fourier.html>

Remember?

Filtering in spatial domain

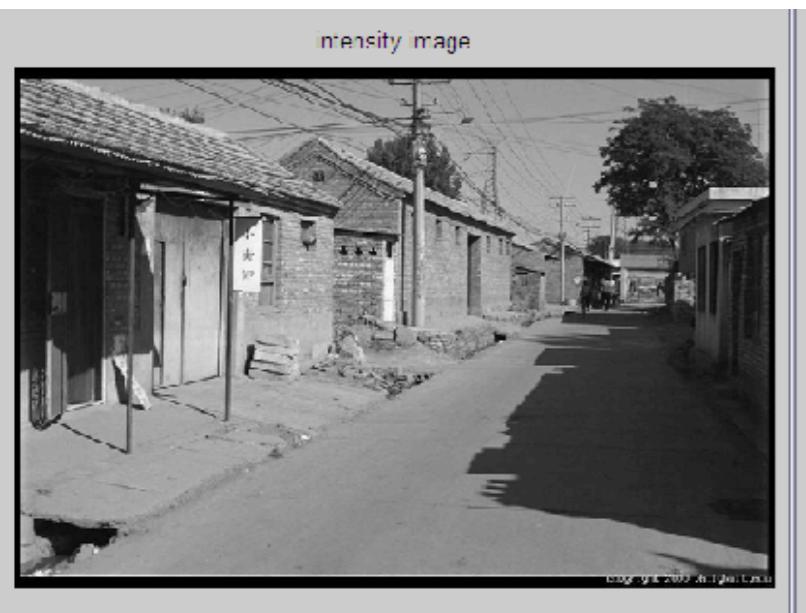


$$* \begin{bmatrix} \text{dark} & \text{medium} & \text{light} \\ \text{medium} & \text{light} & \text{dark} \end{bmatrix} =$$

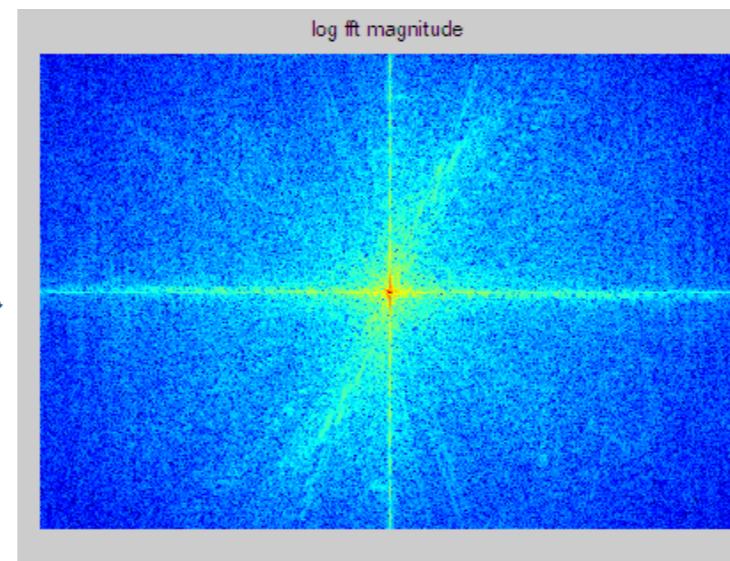


1	0	-1
2	0	-2
1	0	-1

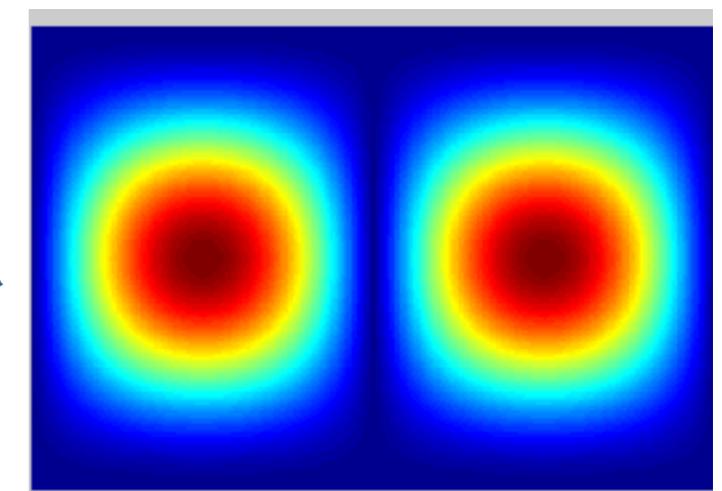
Filtering in frequency domain



FFT

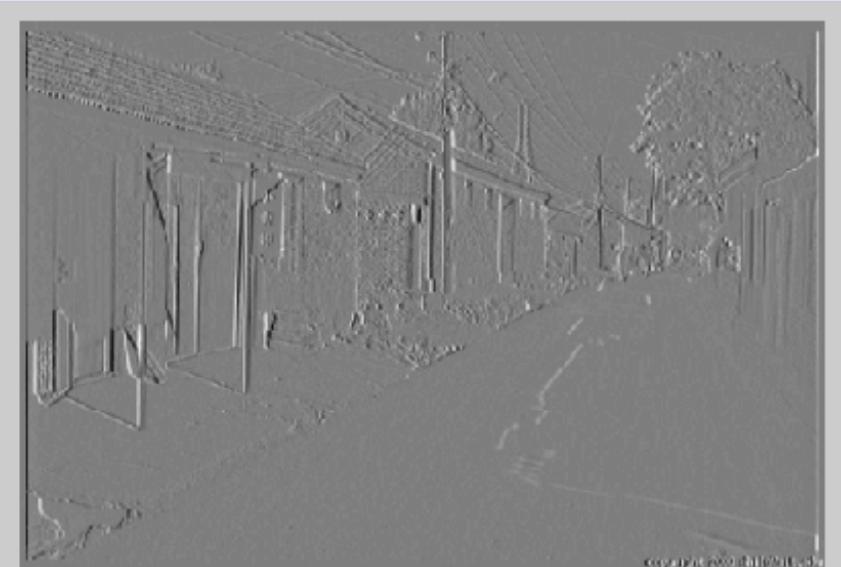
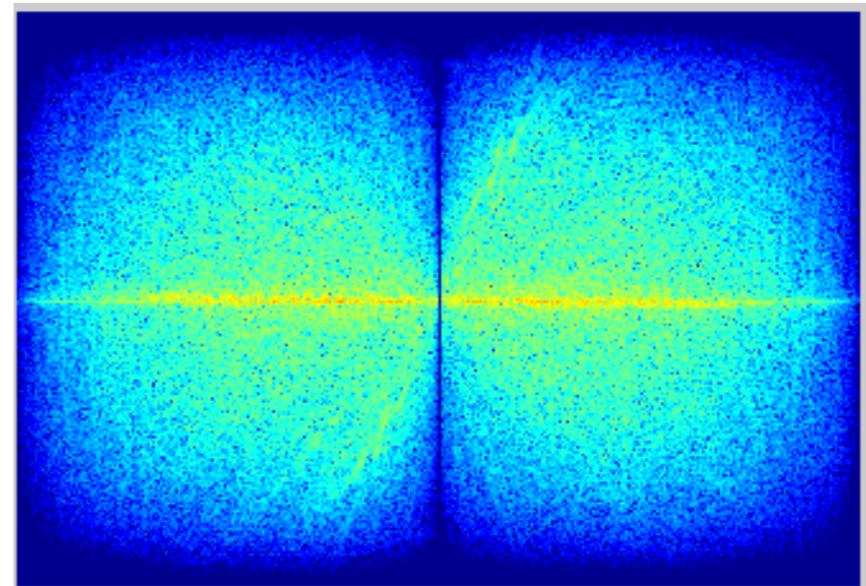


\times



||

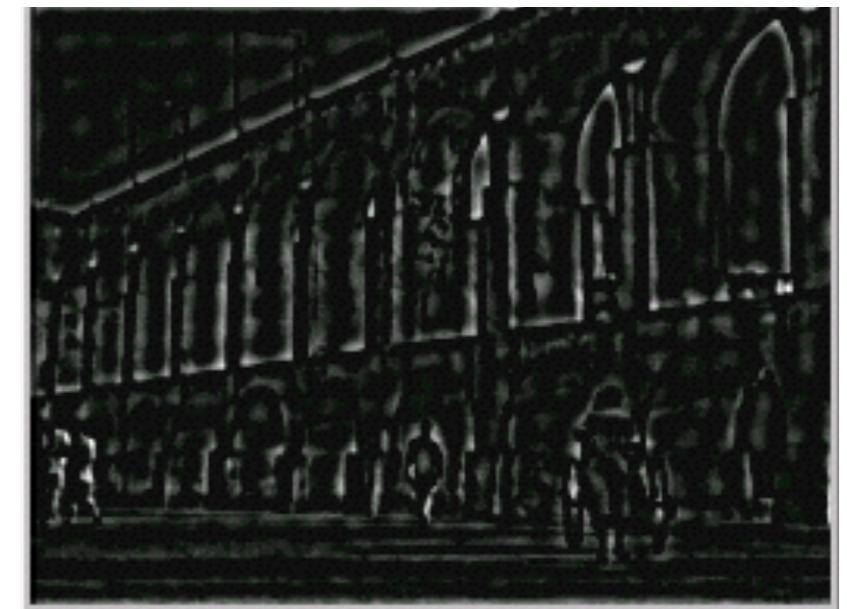
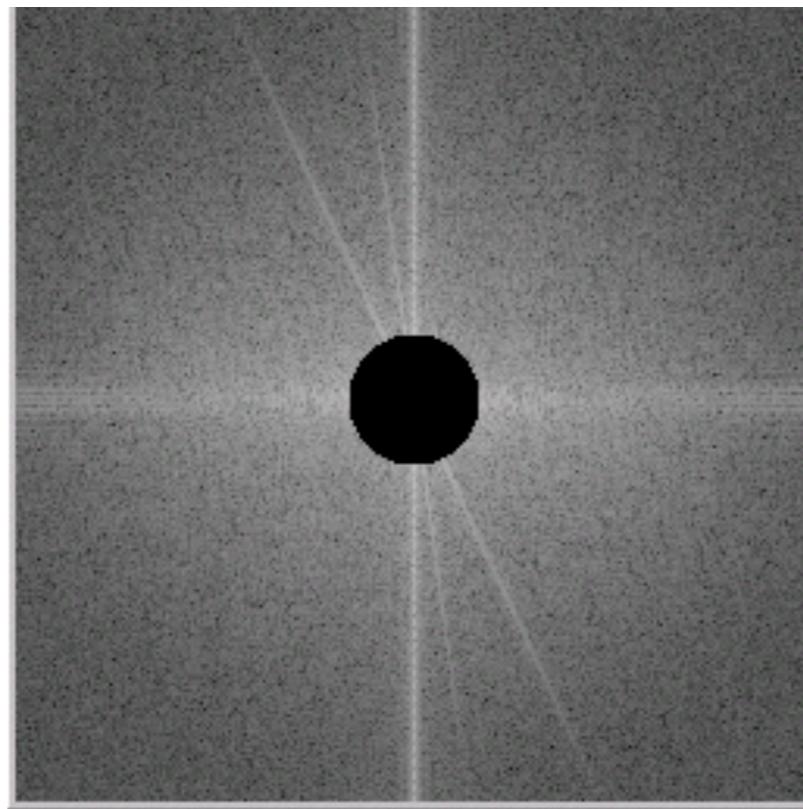
Inverse FFT



FFT



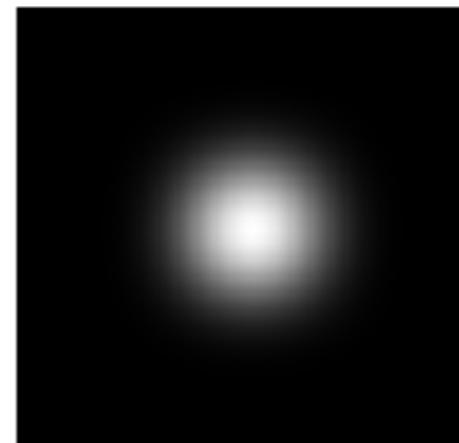
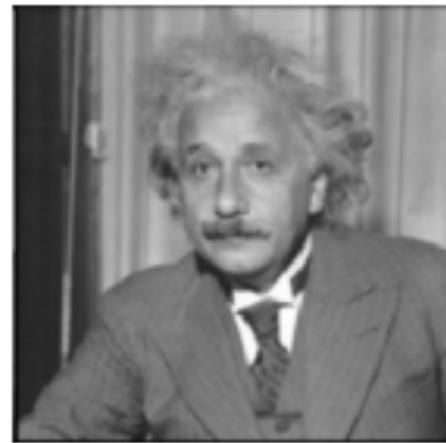
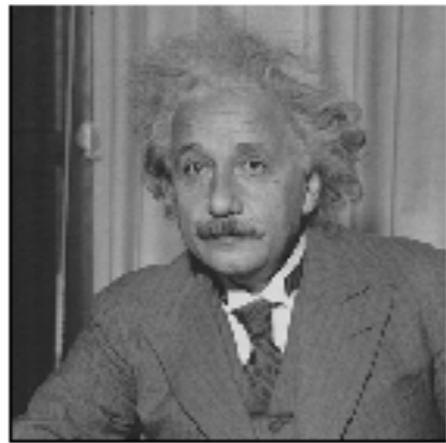
Edges in images



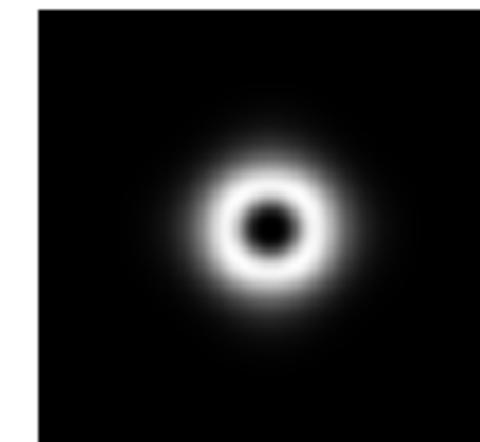
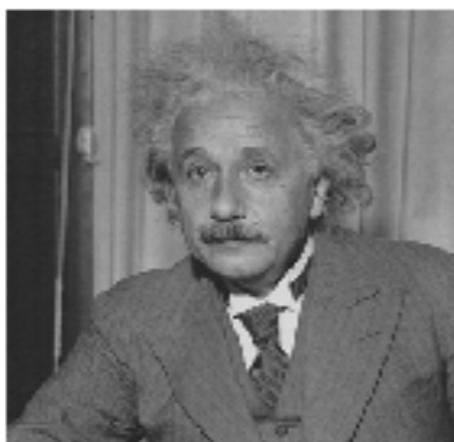
High-pass filter

Low-pass, Band-pass, High-pass filters

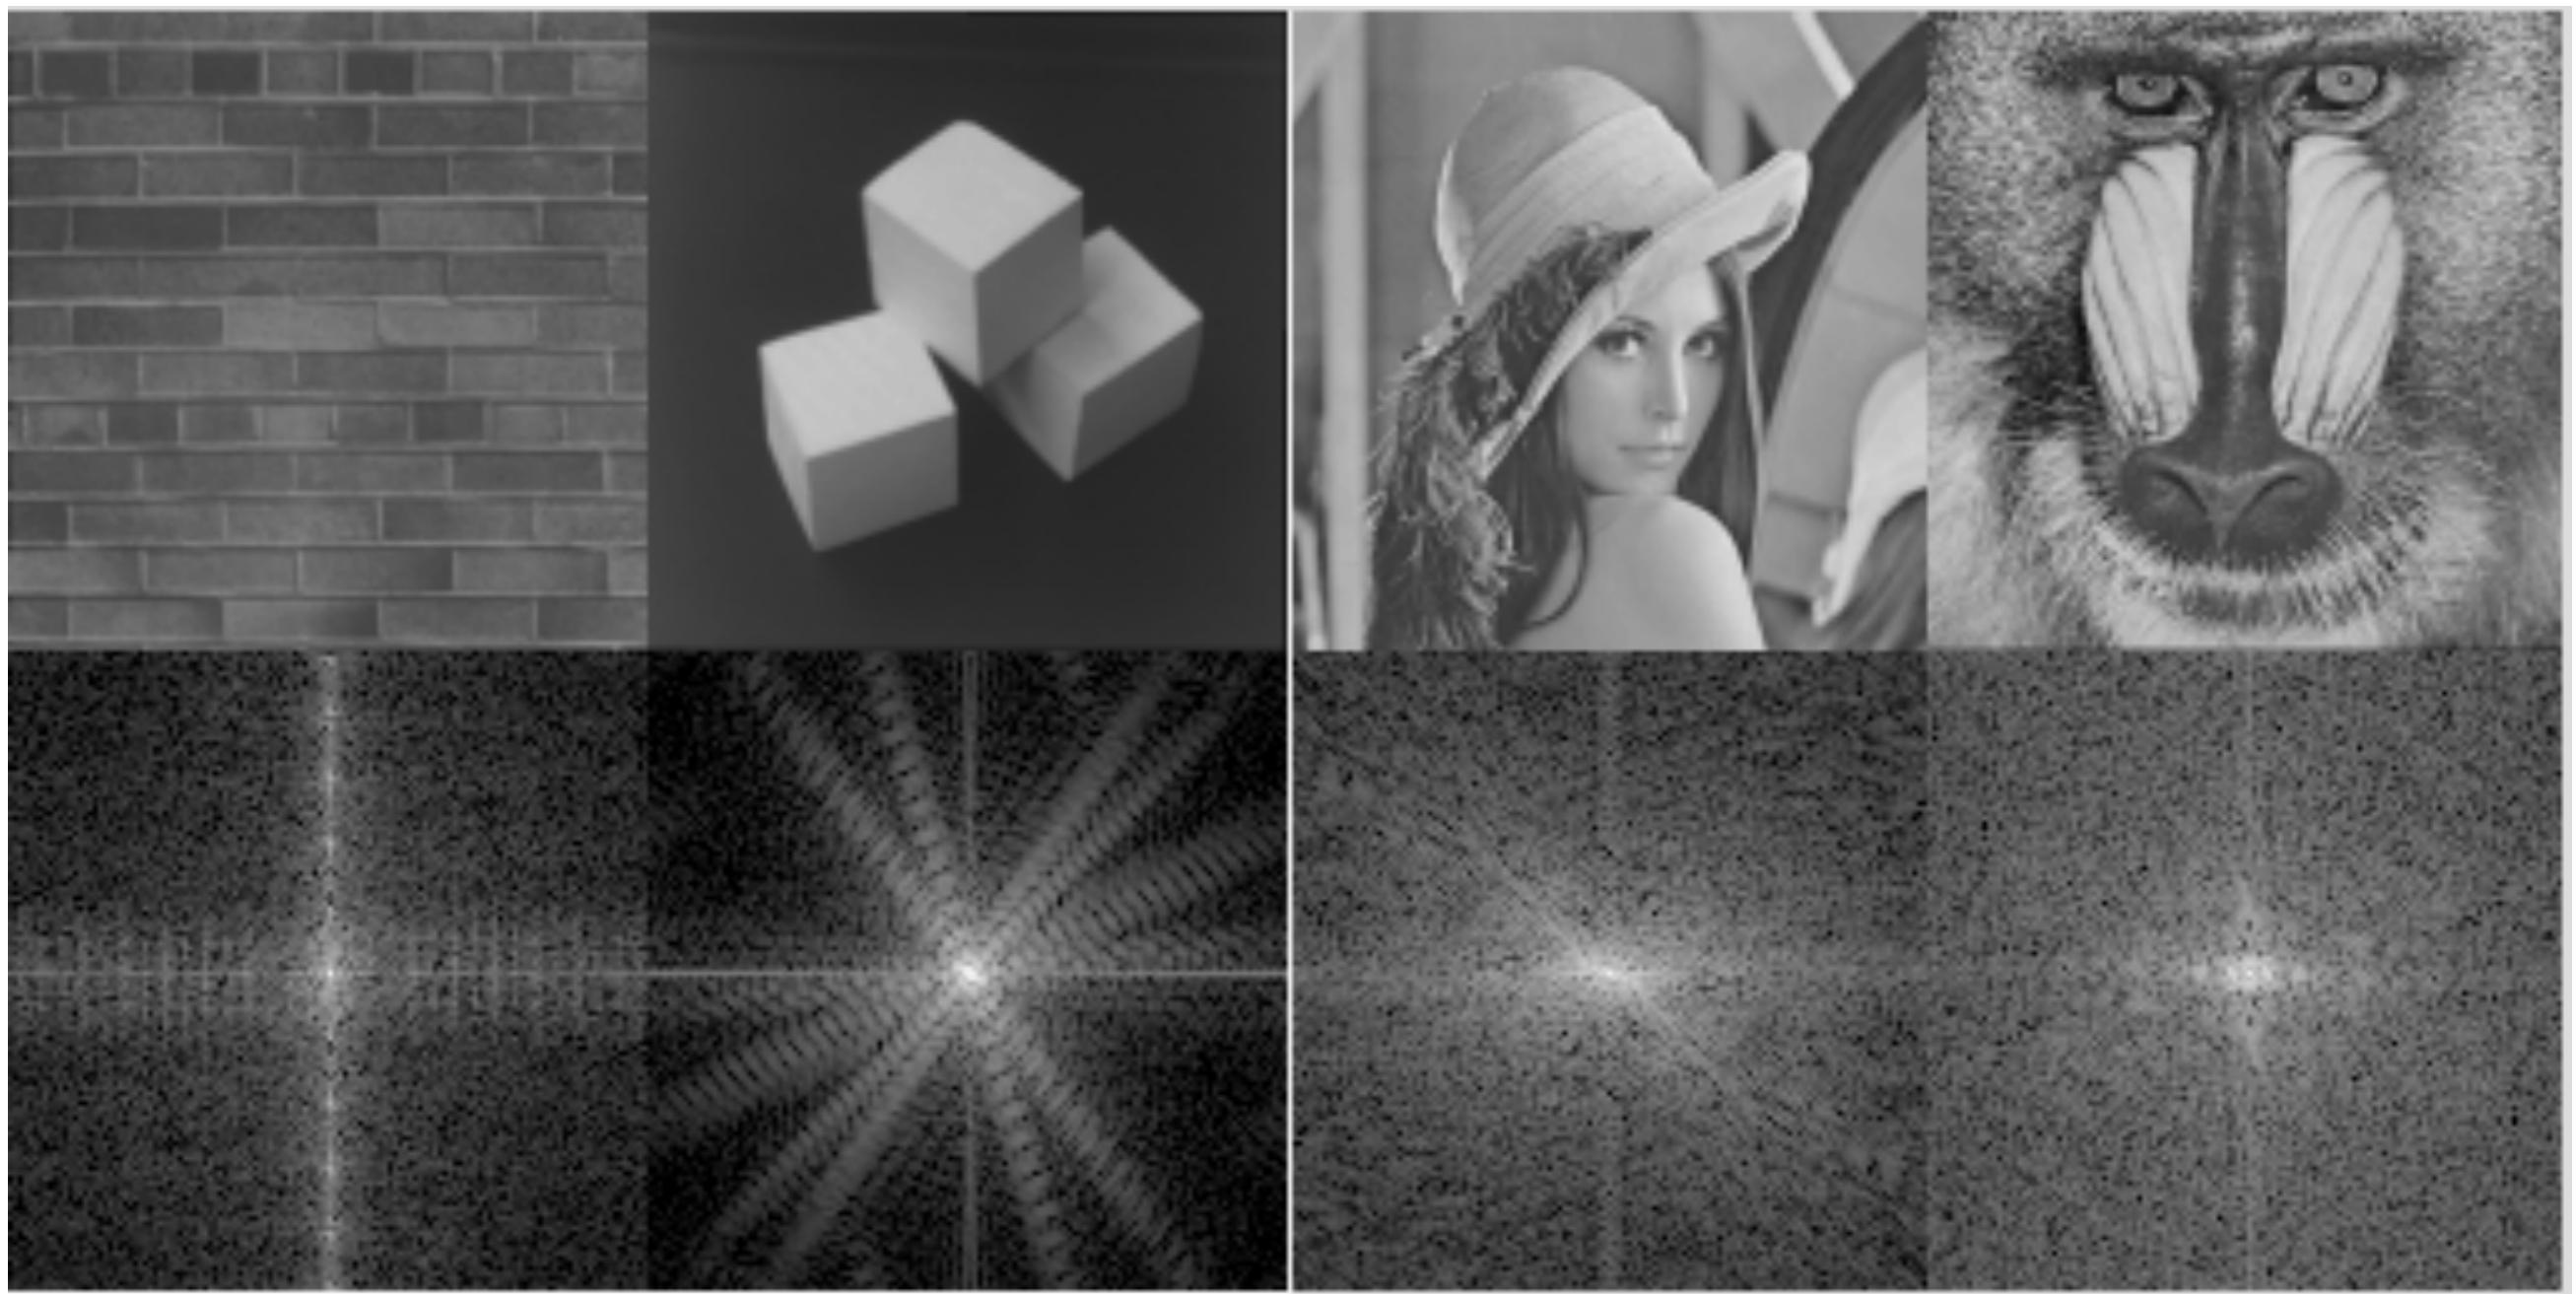
Low-pass:



High-pass / band-pass:



2D Fourier transform examples



Sampling

Sampling

Why does a lower resolution image still make sense to us? What do we lose?



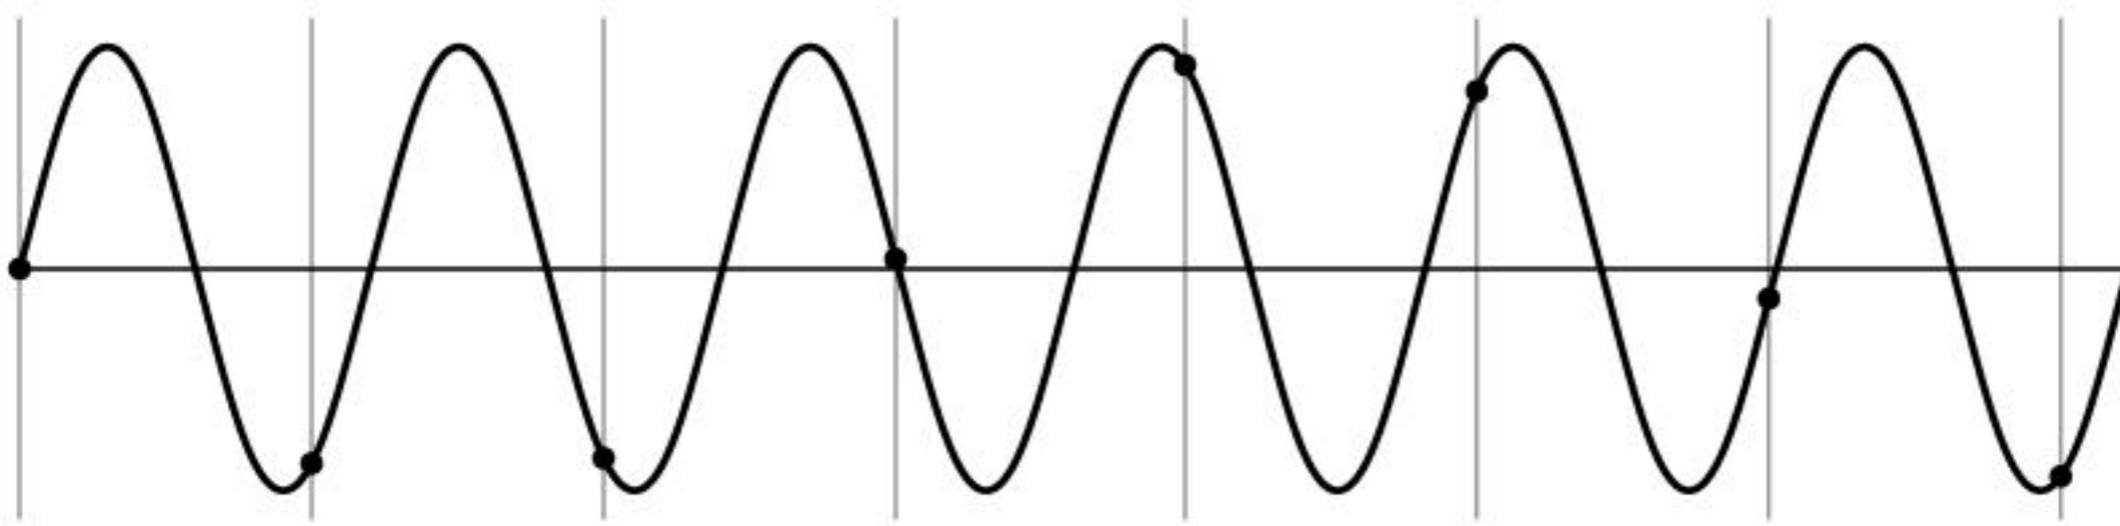
Subsampling by a factor of 2



Throw away every other row and column
to create a 1/2 size image

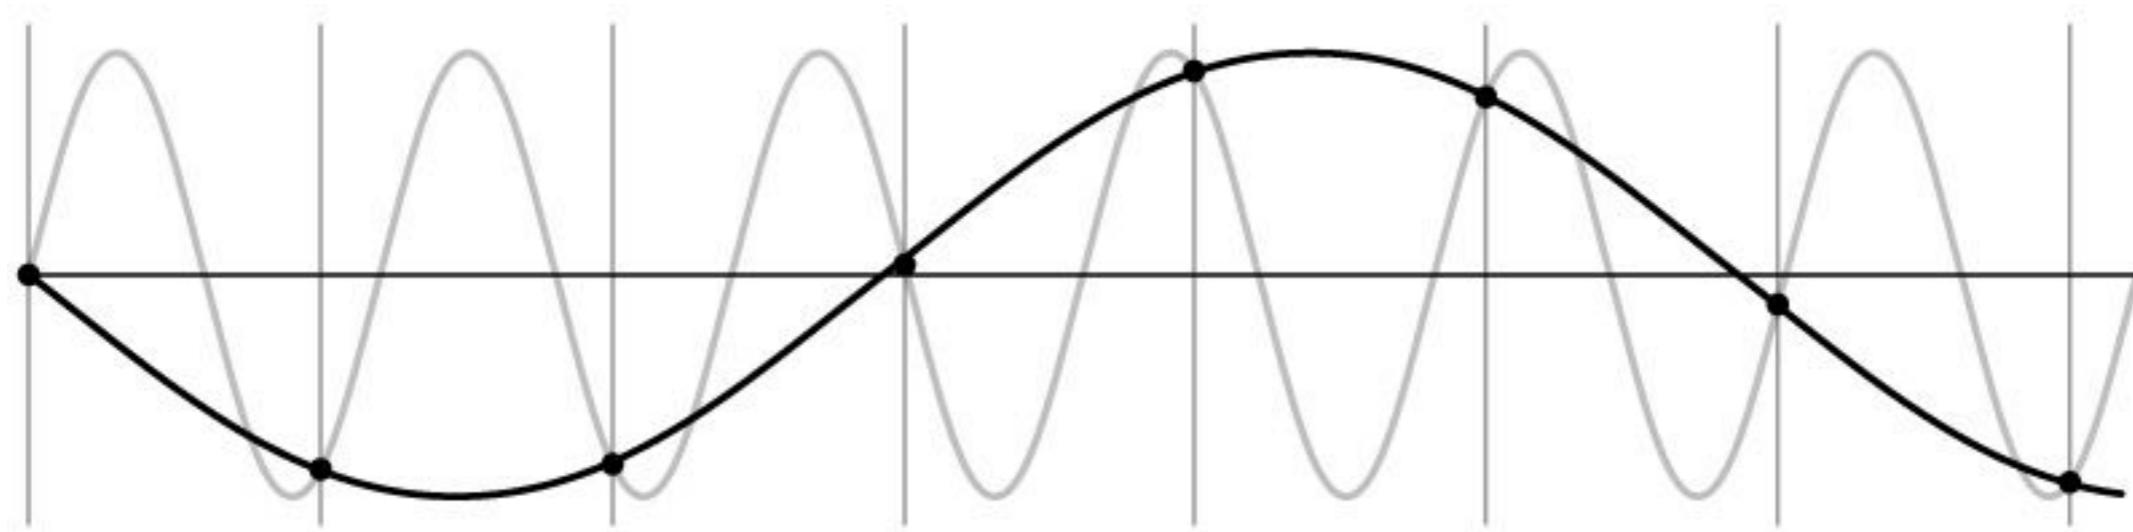
Aliasing problem

- 1D example (sinewave):



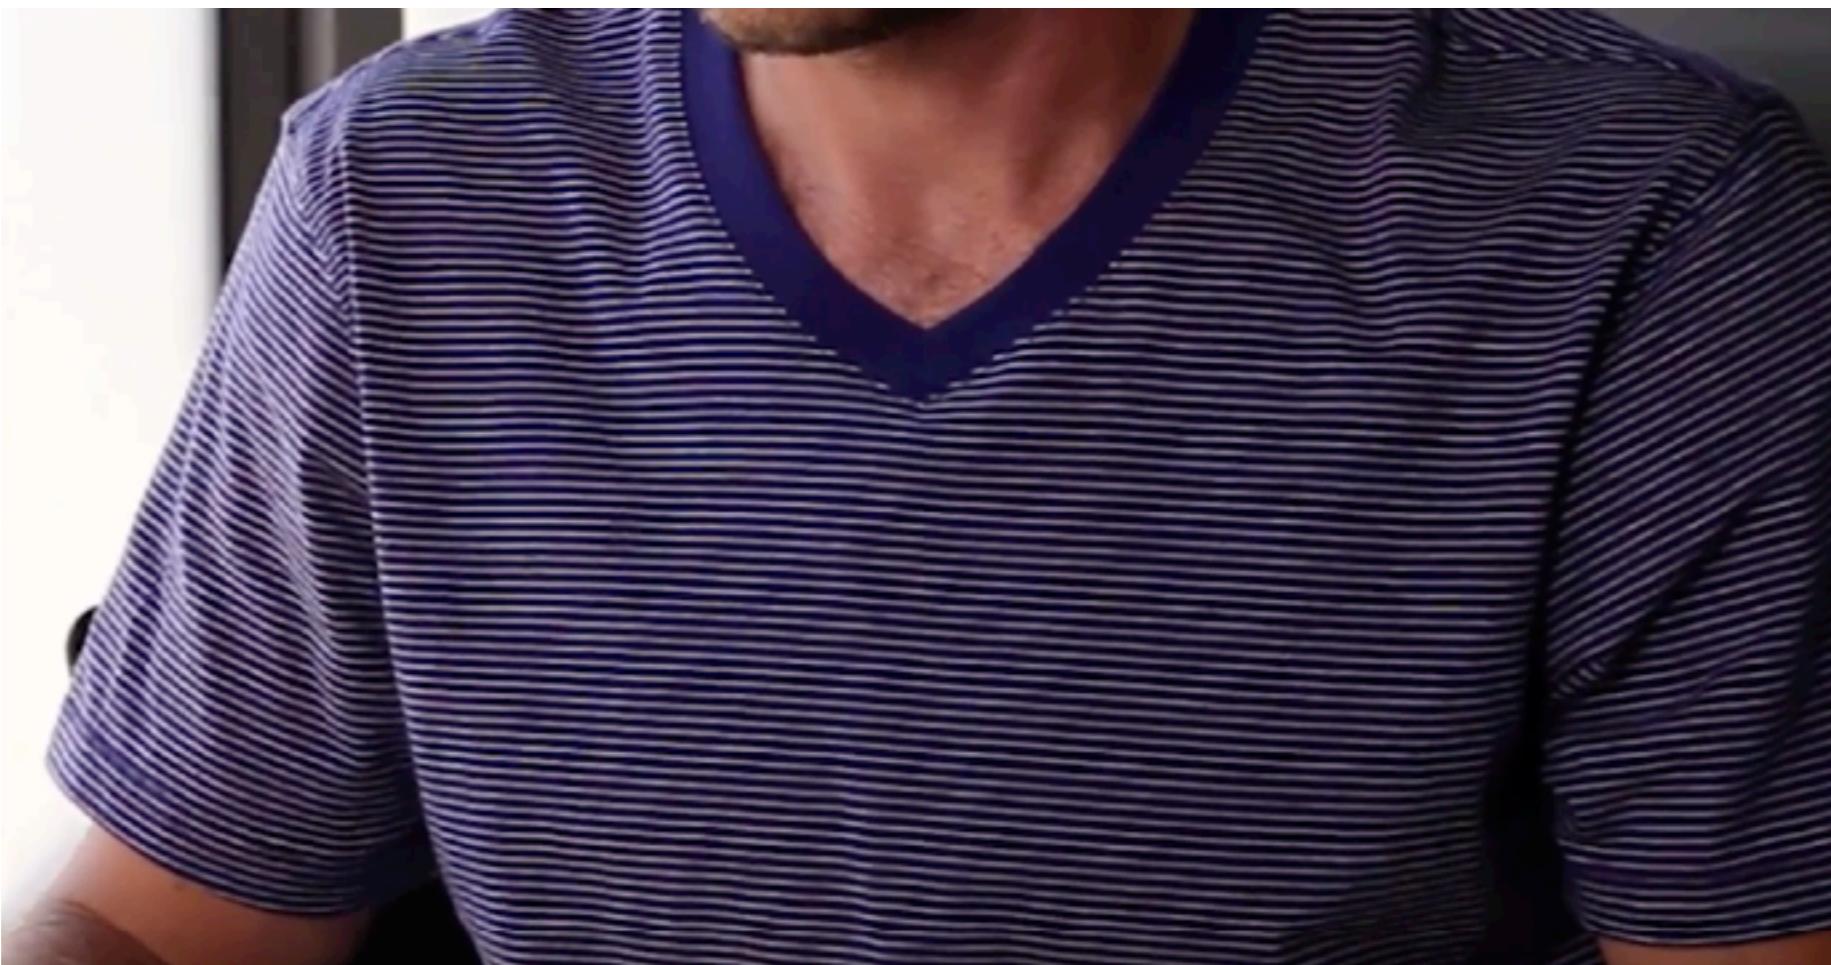
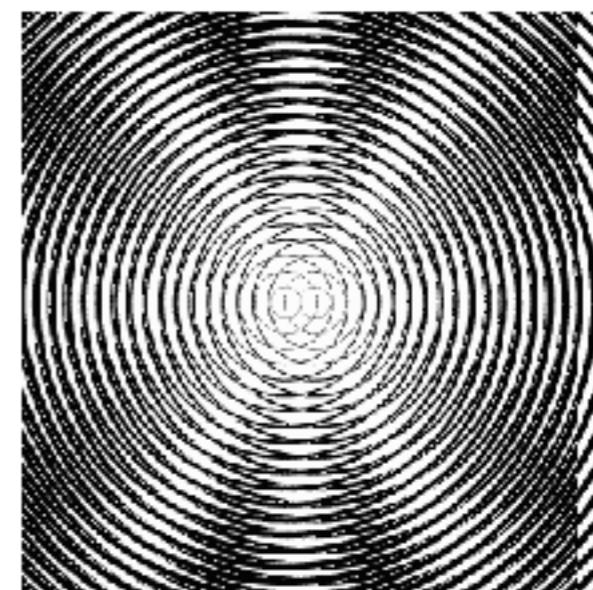
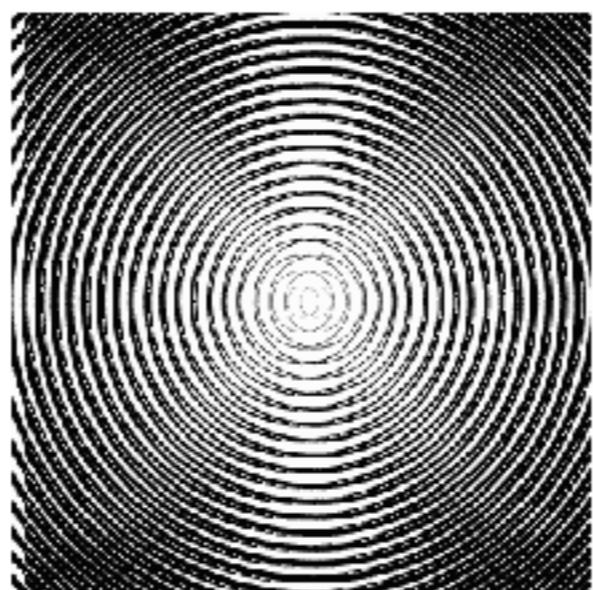
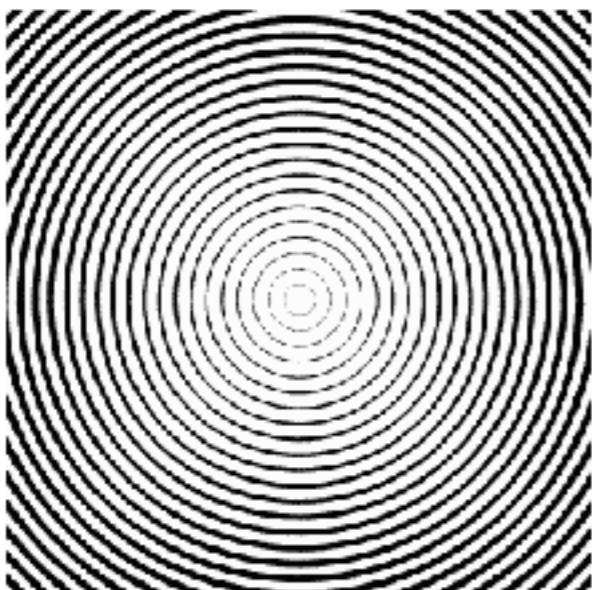
Aliasing problem

- 1D example (sinewave):



**Aliasing happens when
choosing the wrong numbers to
fill in during sampling**

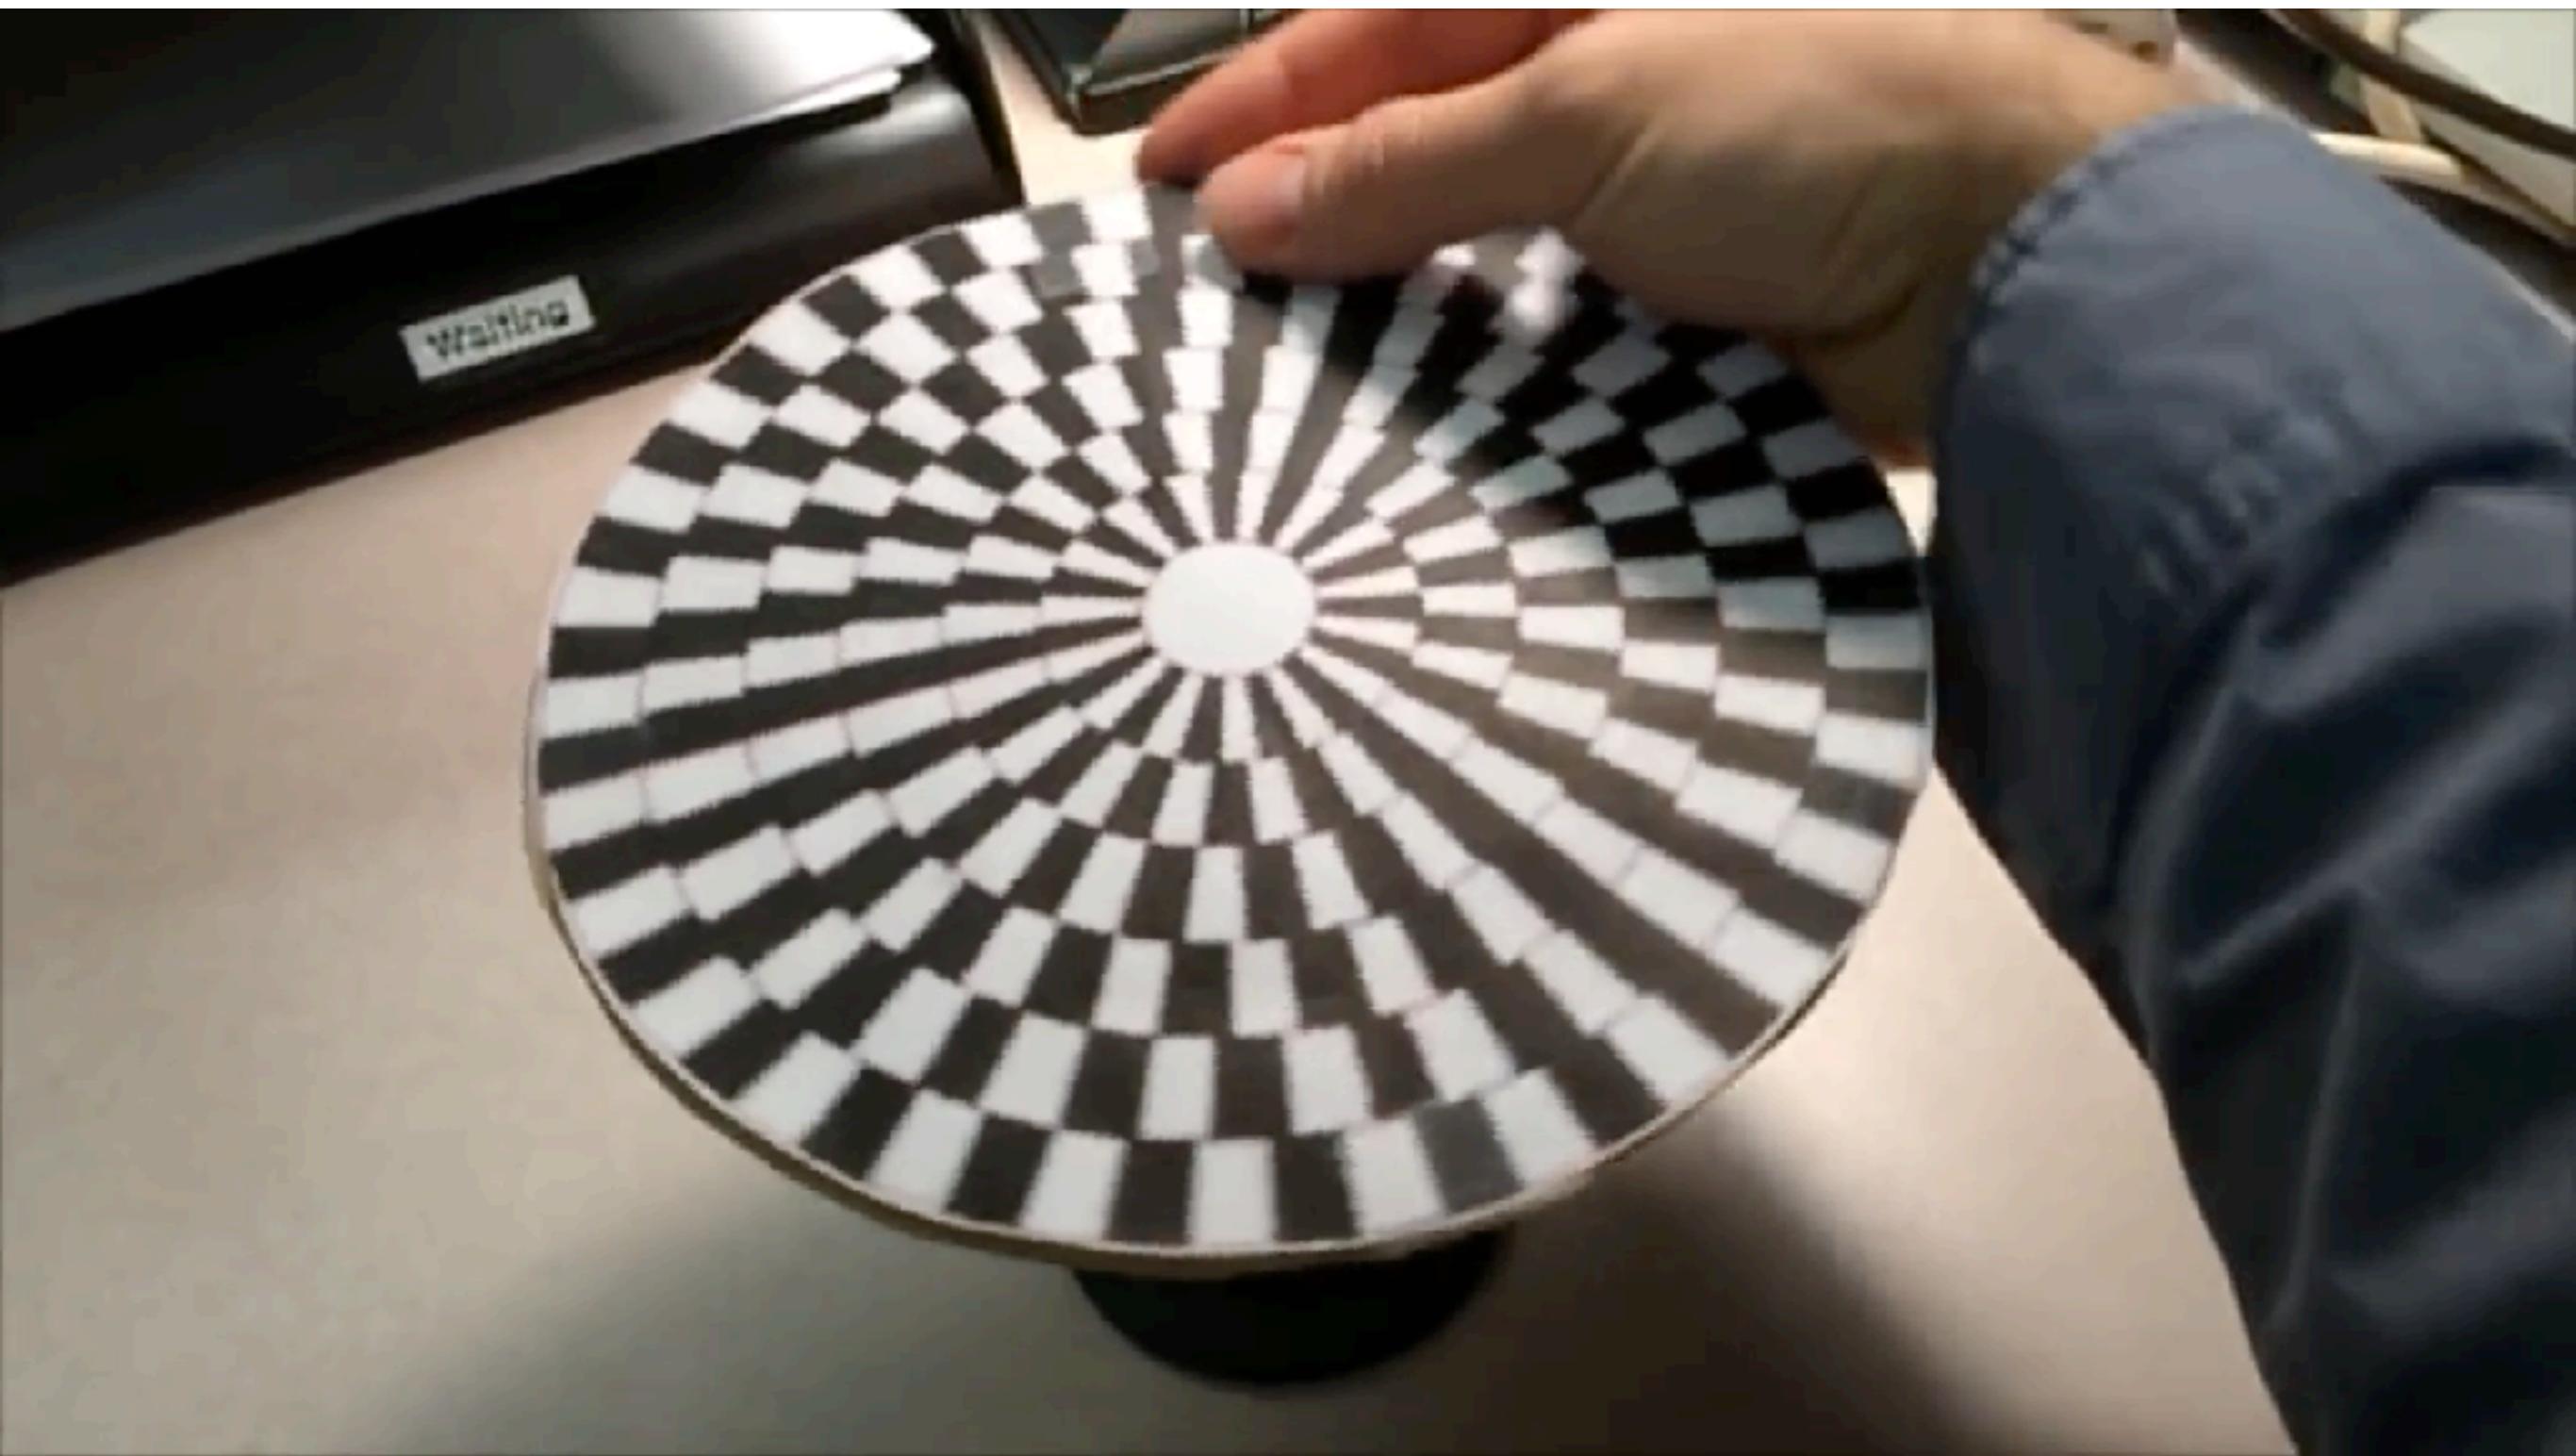
Moire pattern



Aliasing in texture mapping



Aliasing in a video – wagon wheel effect

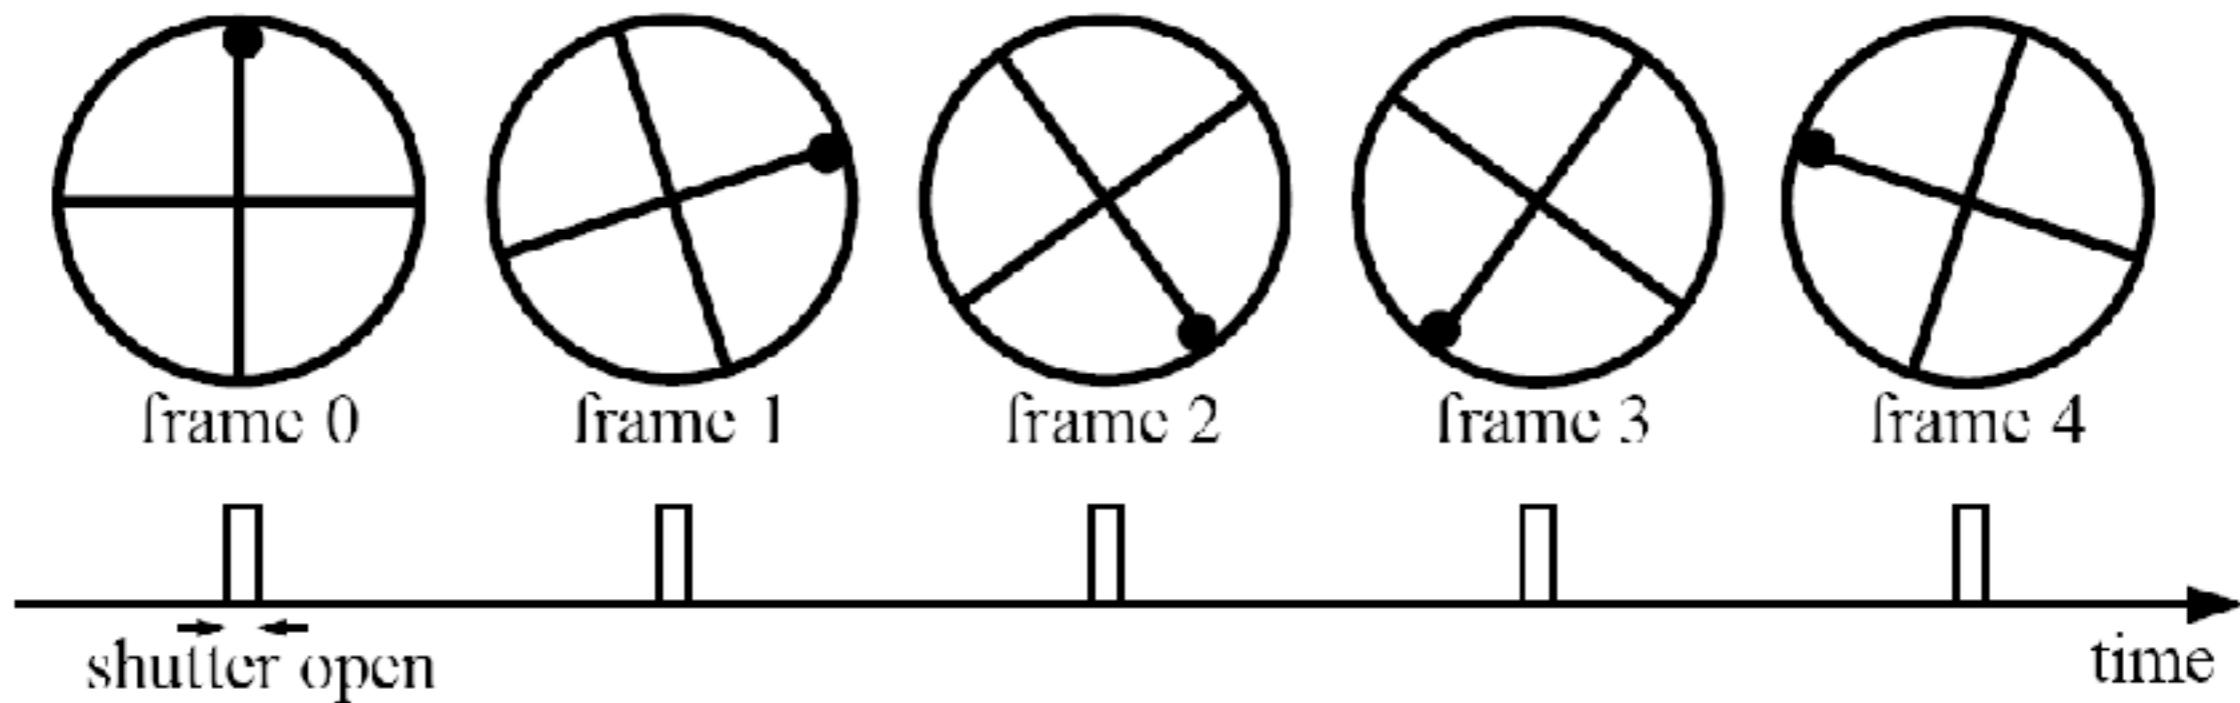


Aliasing in a video

Imagine a spoked wheel moving to the right (rotating clockwise).

Mark wheel with dot so we can see what's happening.

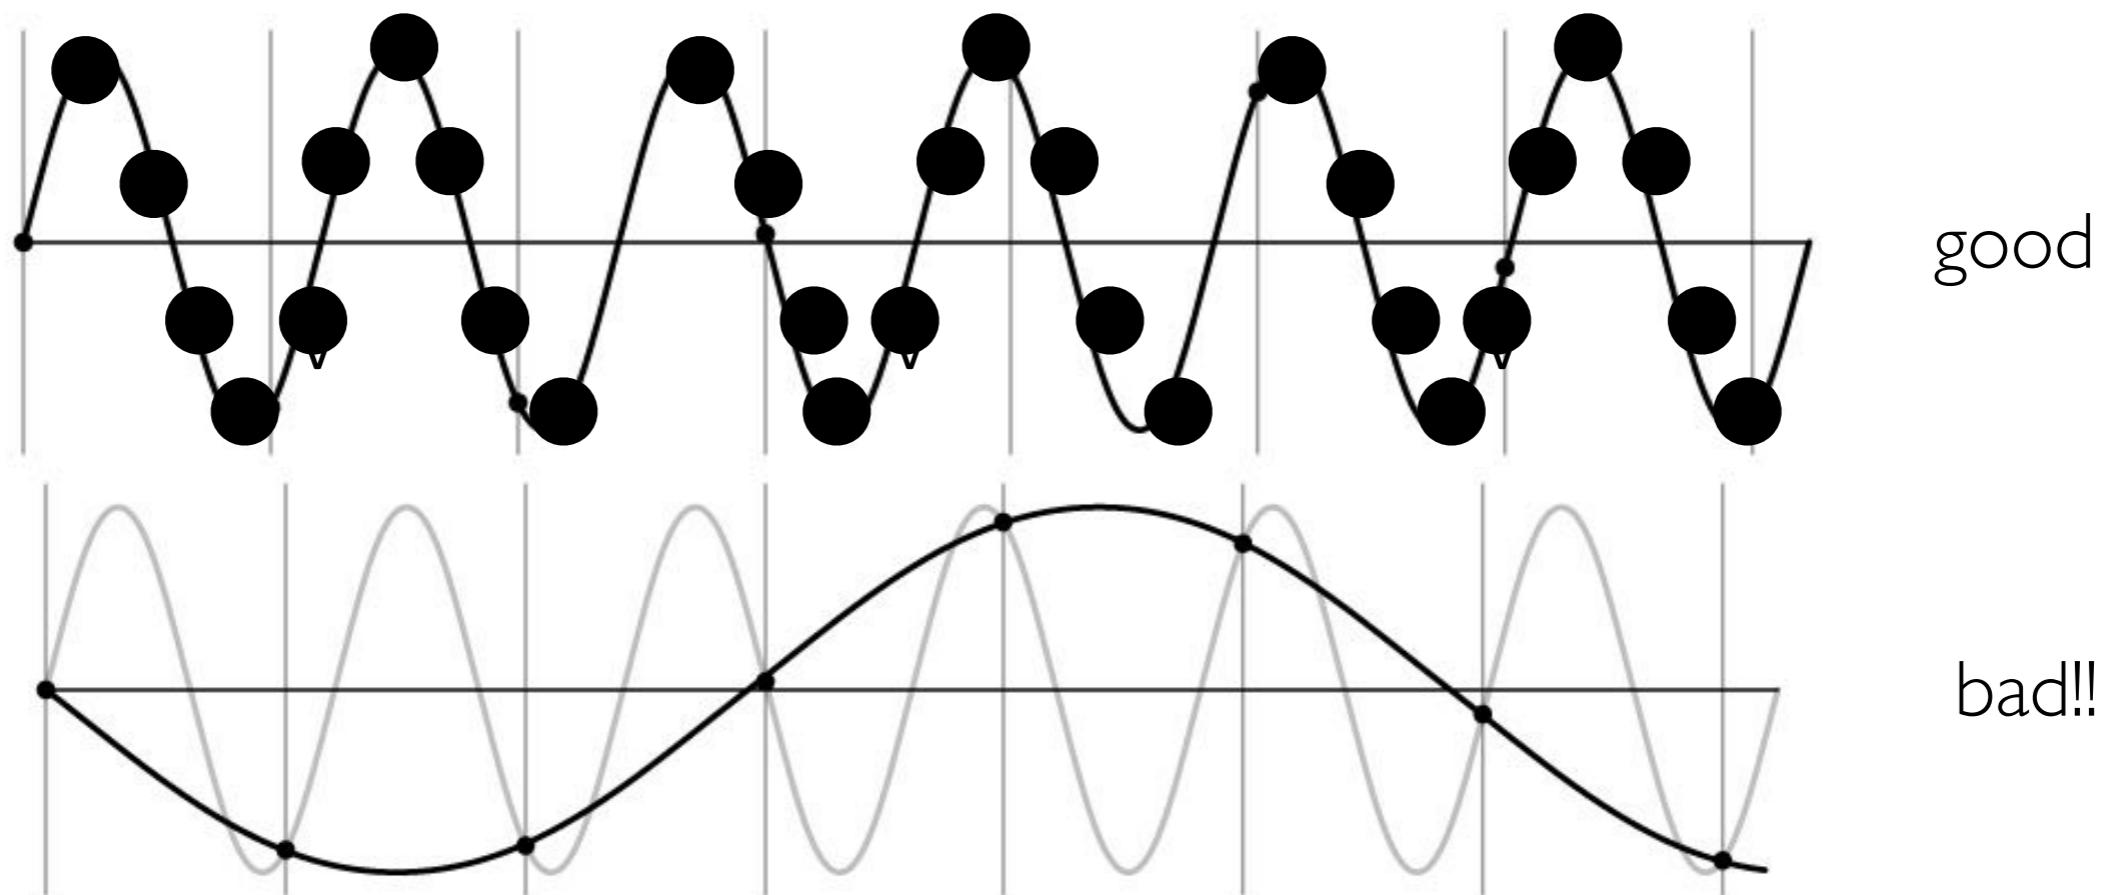
If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):



Without dot, wheel appears to be rotating slowly backwards!
(counterclockwise)

Nyquist-Shannon Sampling Theorem

- When sampling a signal at discrete intervals, the sampling frequency must be $\geq 2 \times f_{\max}$
- f_{\max} = max frequency of the input signal
- This will allow to reconstruct the original perfectly from the sampled version



Anti-aliasing

Two Solutions:

- Sample more often
- Get rid of all frequencies that are greater than half the new sampling frequency
 - Will lose information
 - But it's better than aliasing
 - Apply a smoothing filter

Algorithm for downsampling by factor of 2

1. Start with $\text{image}(h, w)$

2. Apply low-pass filter

```
im.blur = imfilter(image, fspecial('gaussian', 7, 1))
```

3. Sample every other pixel

```
im.small = im.blur(1:2:end, 1:2:end);
```

Sampling and Aliasing

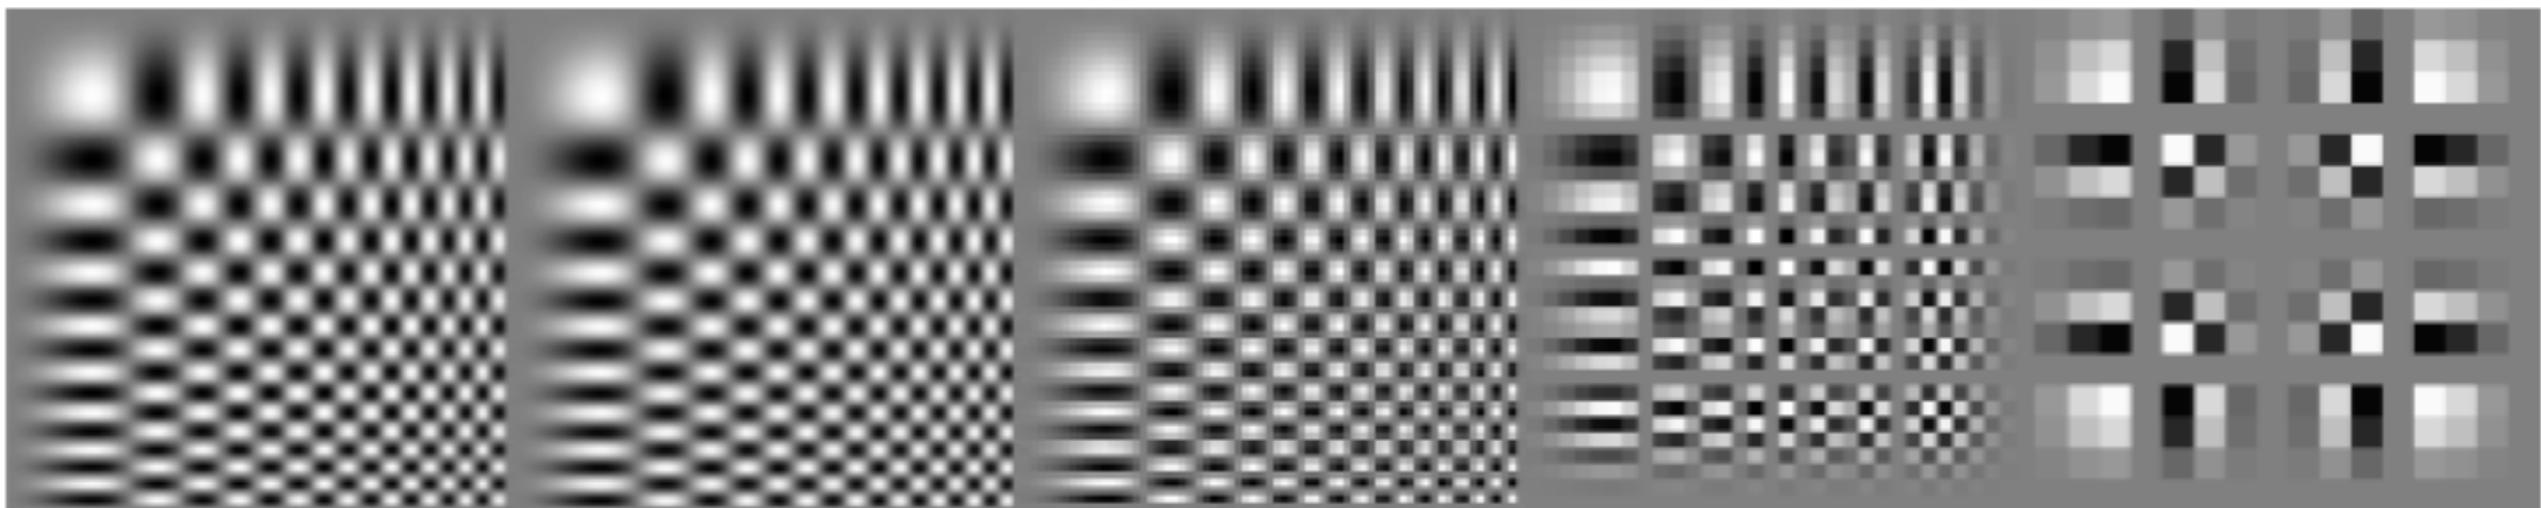
256x256

128x128

64x64

32x32

16x16



Anti-aliasing

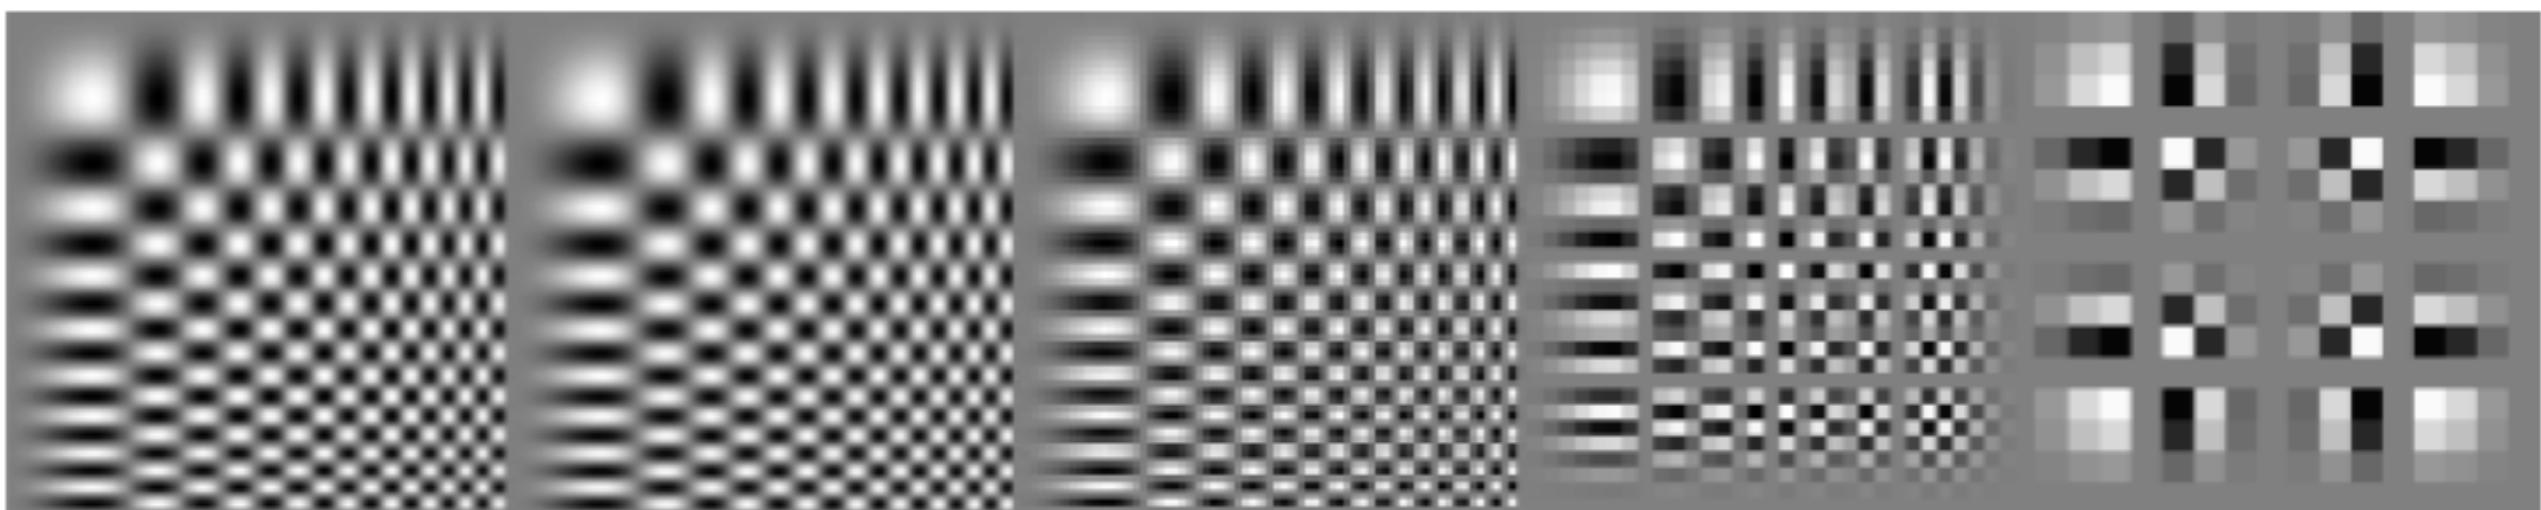
256x256

128x128

64x64

32x32

16x16



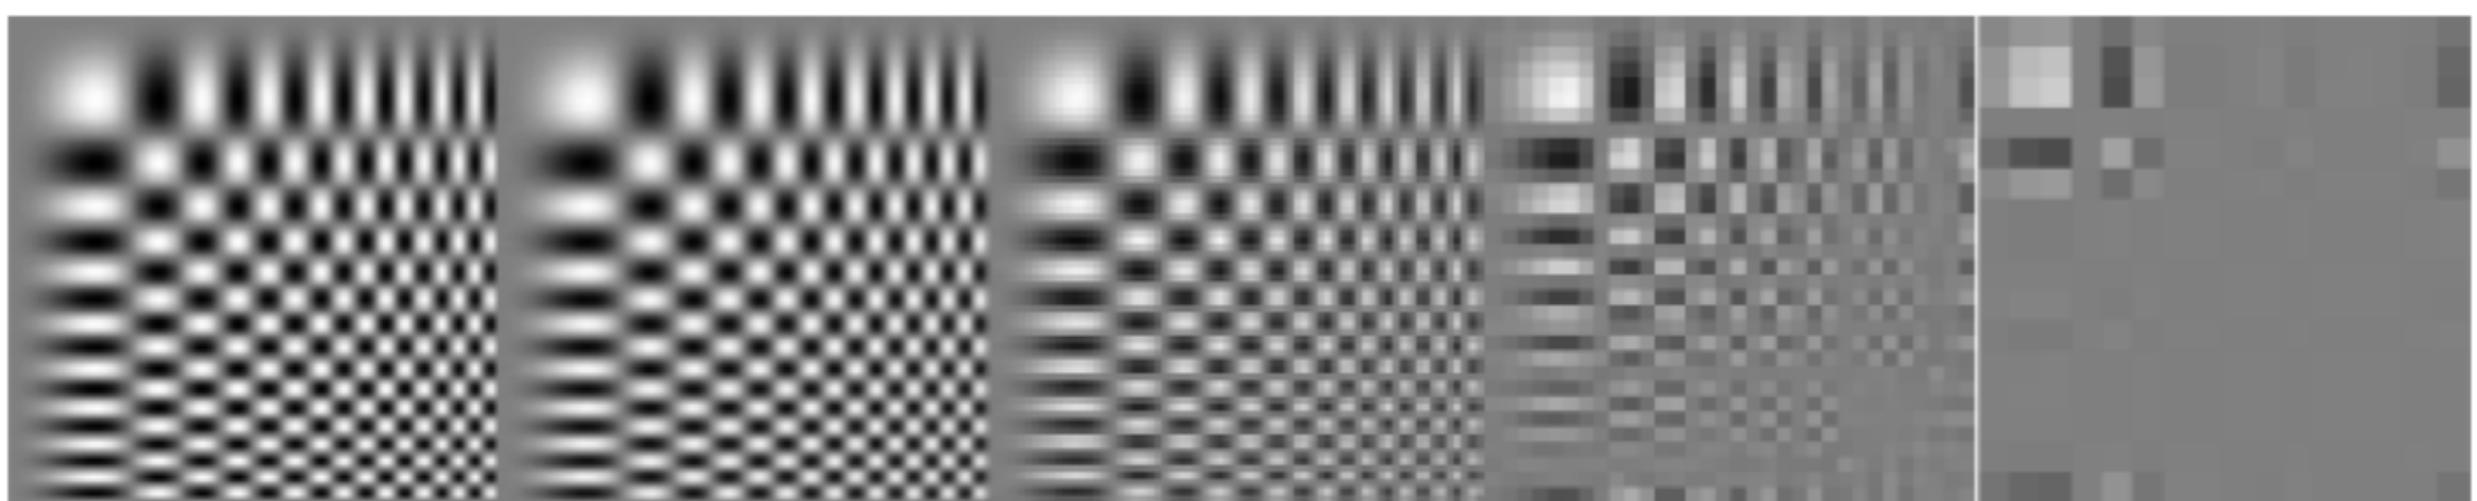
256x256

128x128

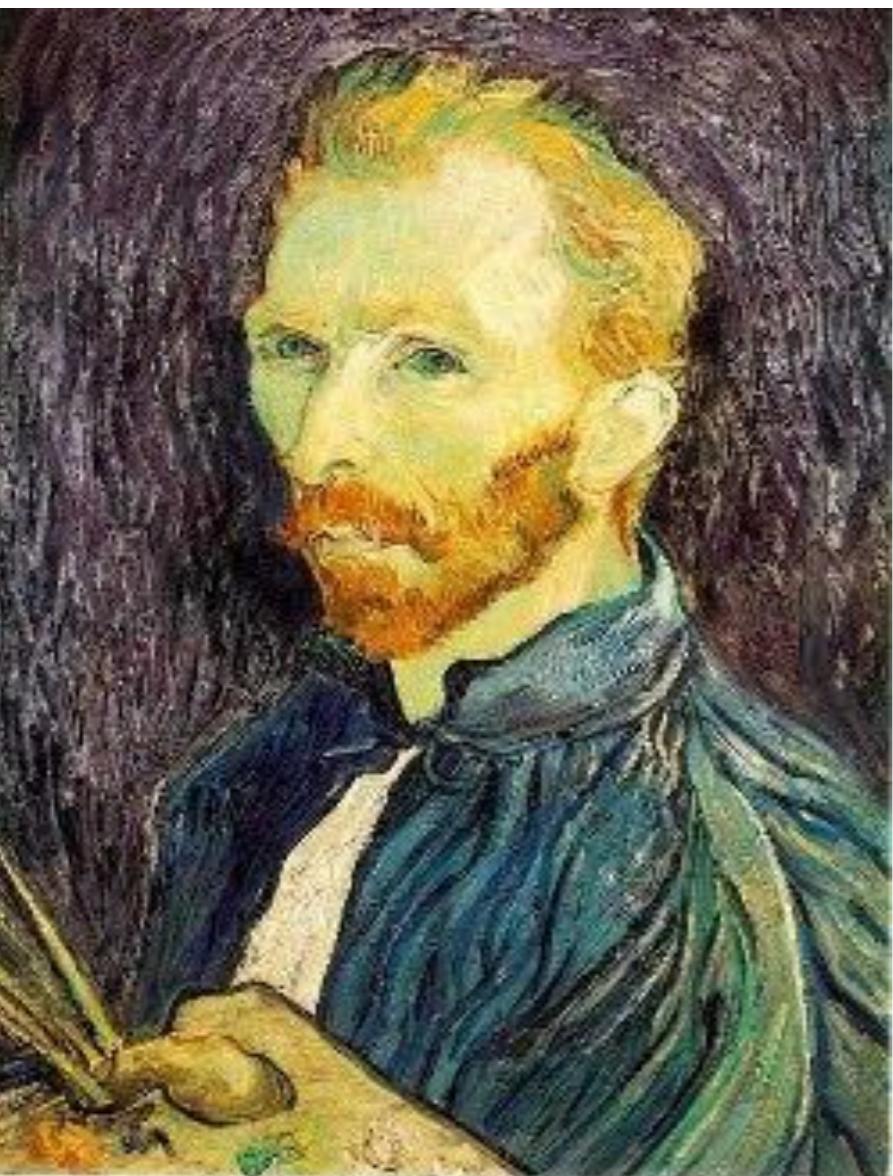
64x64

32x32

16x16



Subsampling without pre-filtering



1/2



1/4 (2x zoom)

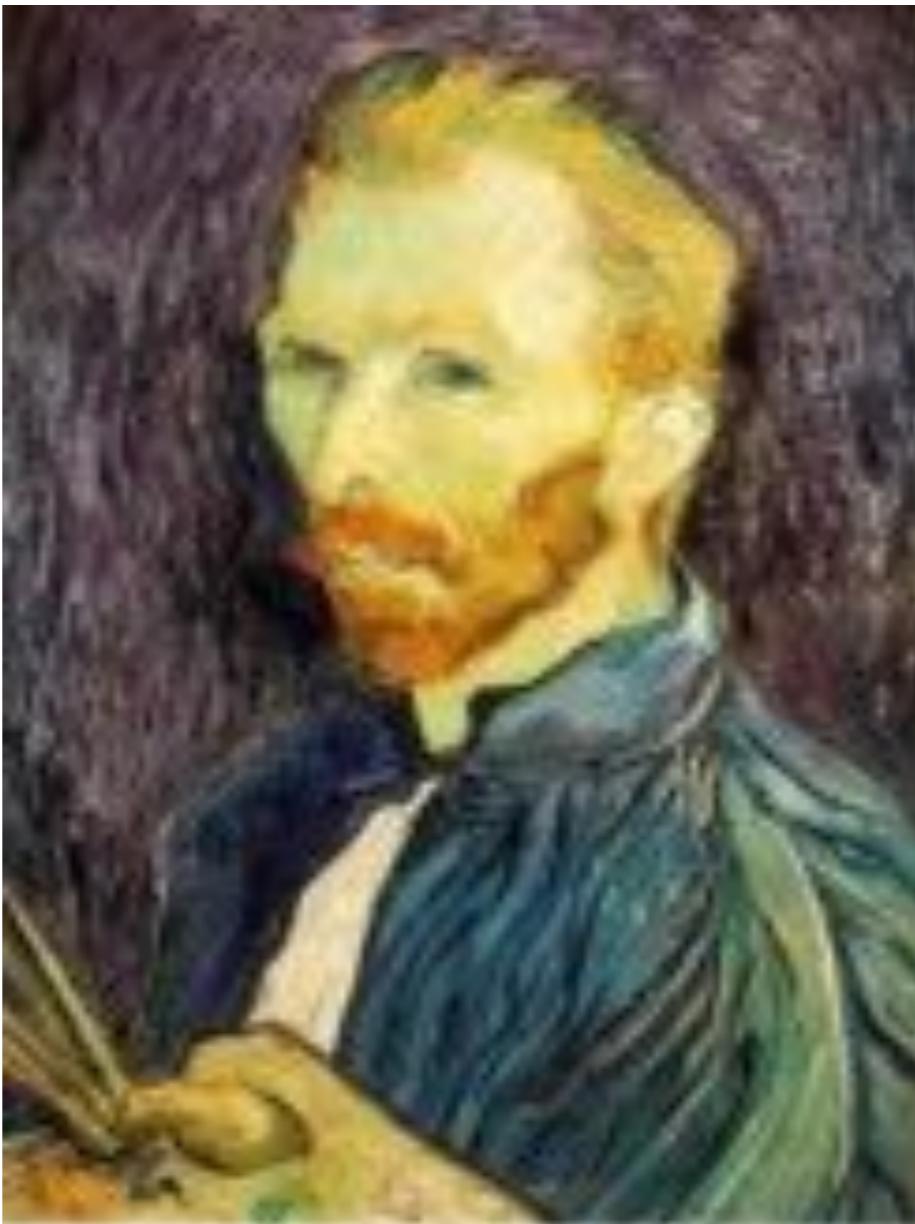


1/8 (4x zoom)

Subsampling with Gaussian pre-filtering



Gaussian 1/2



G 1/4

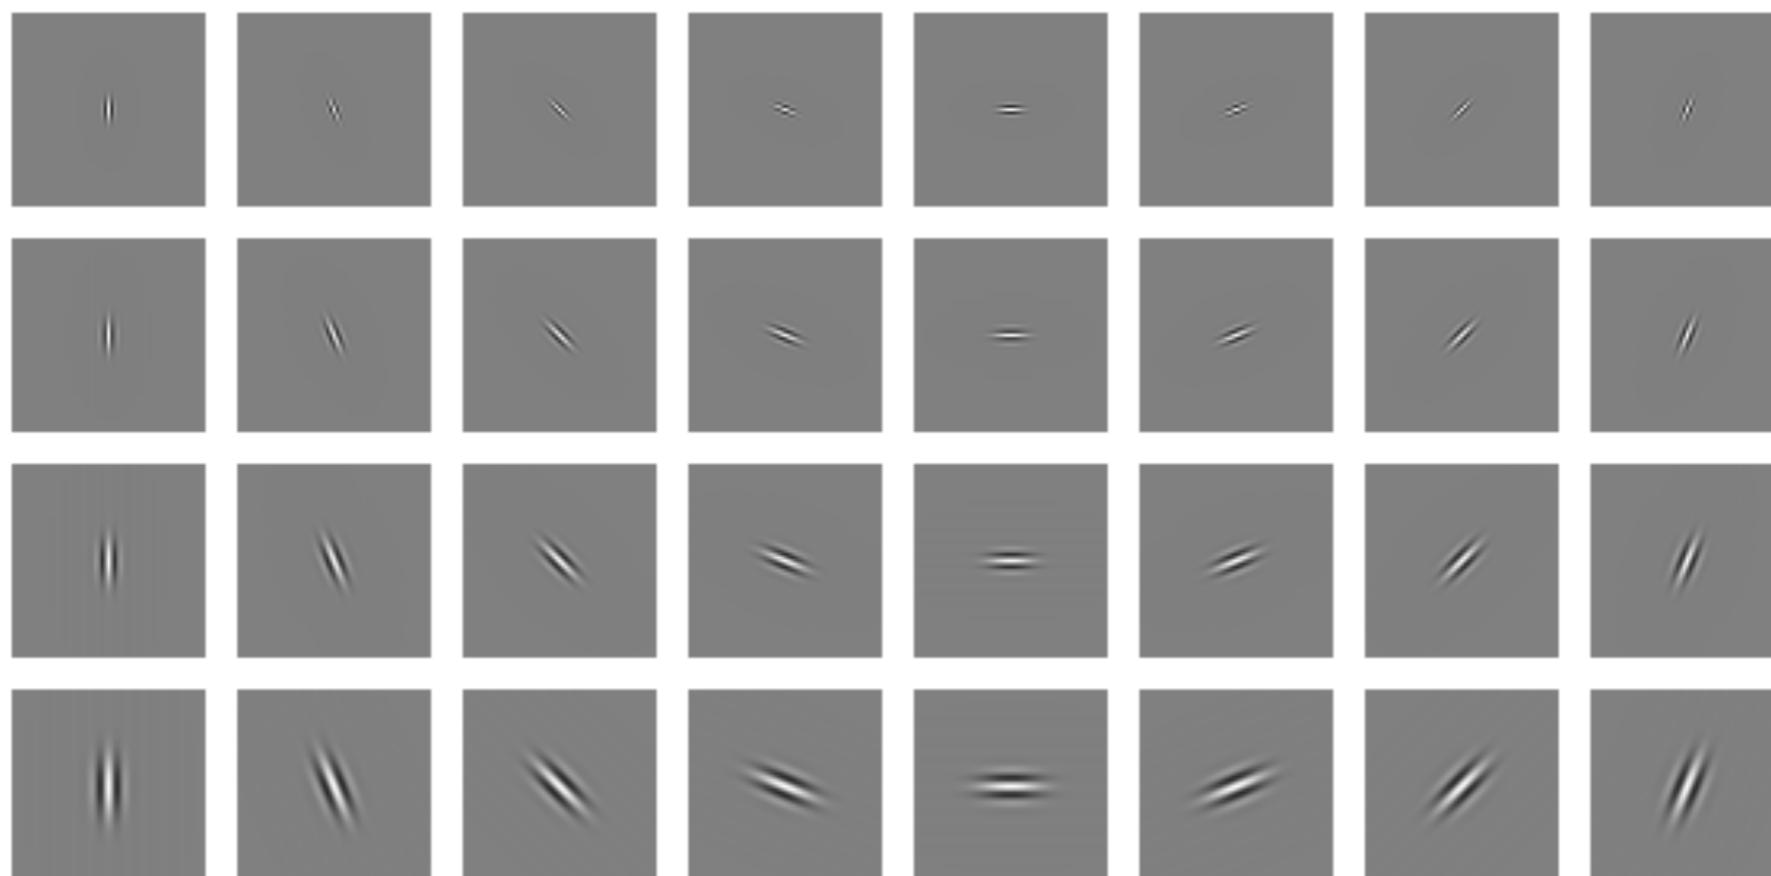


G 1/8

**Always filtering and
then sampling!**

Subsampling with Gaussian pre-filtering

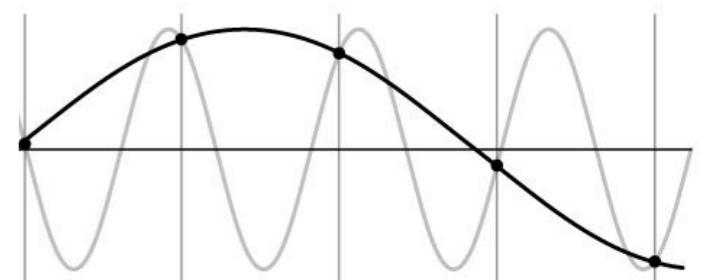
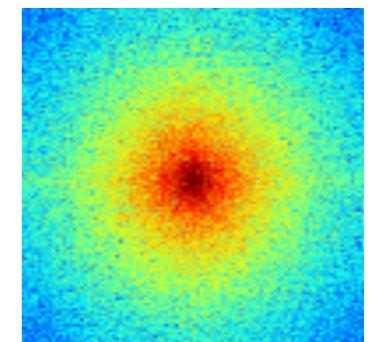
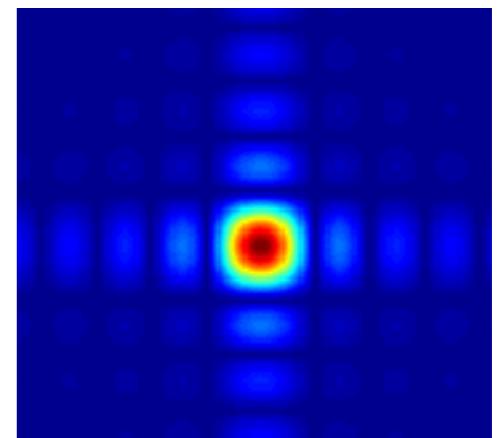
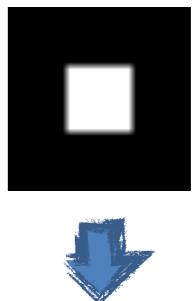
- Early processing in humans filters for various orientations and scales of frequency
- Perceptual cues in the mid-high frequencies dominate perception
- When we see an image from far away, we are effectively subsampling it



Early Visual Processing: Multi-scale edge and blob filters

Things to Remember

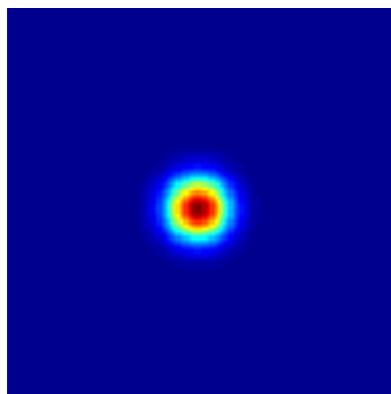
- Sometimes it makes sense to think of images and filtering in the frequency domain
 - Fourier analysis
- Can be faster to filter using FFT for large images ($N \log N$ vs. N^2)
- Images are mostly smooth
 - Basis for compression
- Remember to low-pass before sampling



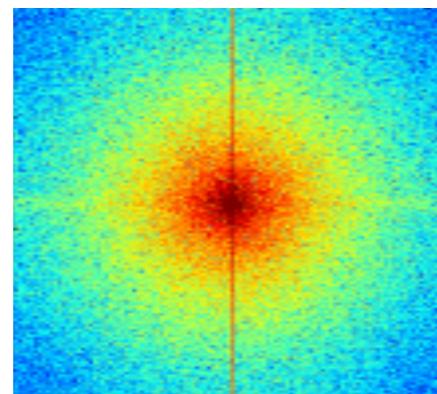
Practice question

Match the spatial domain image to the Fourier magnitude image

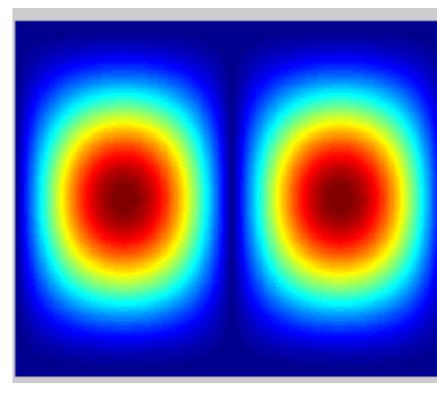
1



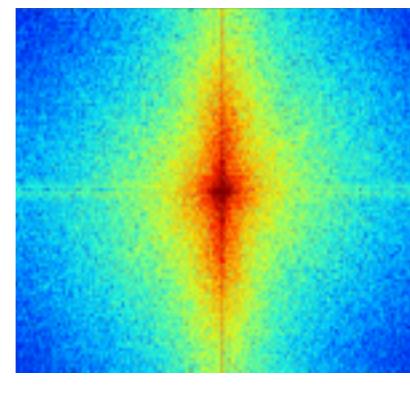
2



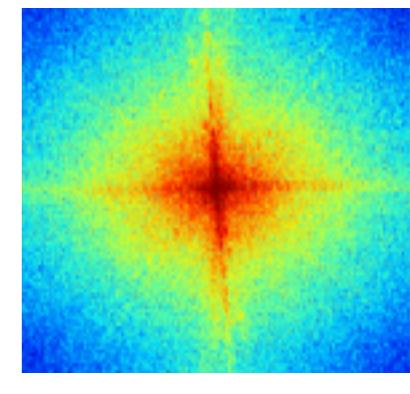
3



4



5



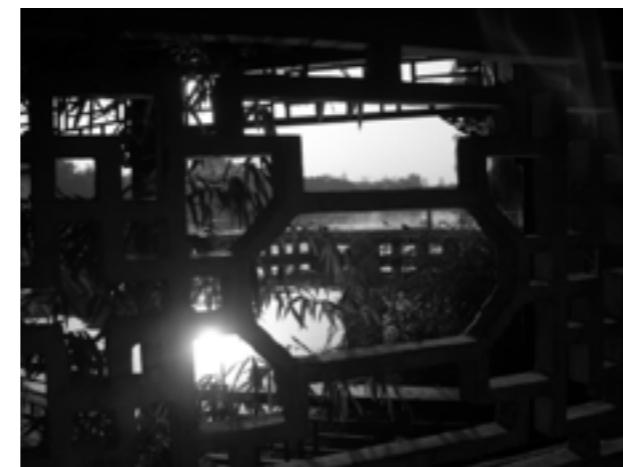
A



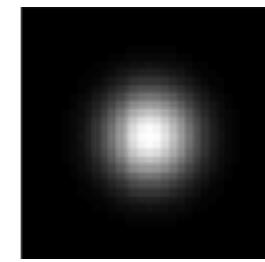
B



C



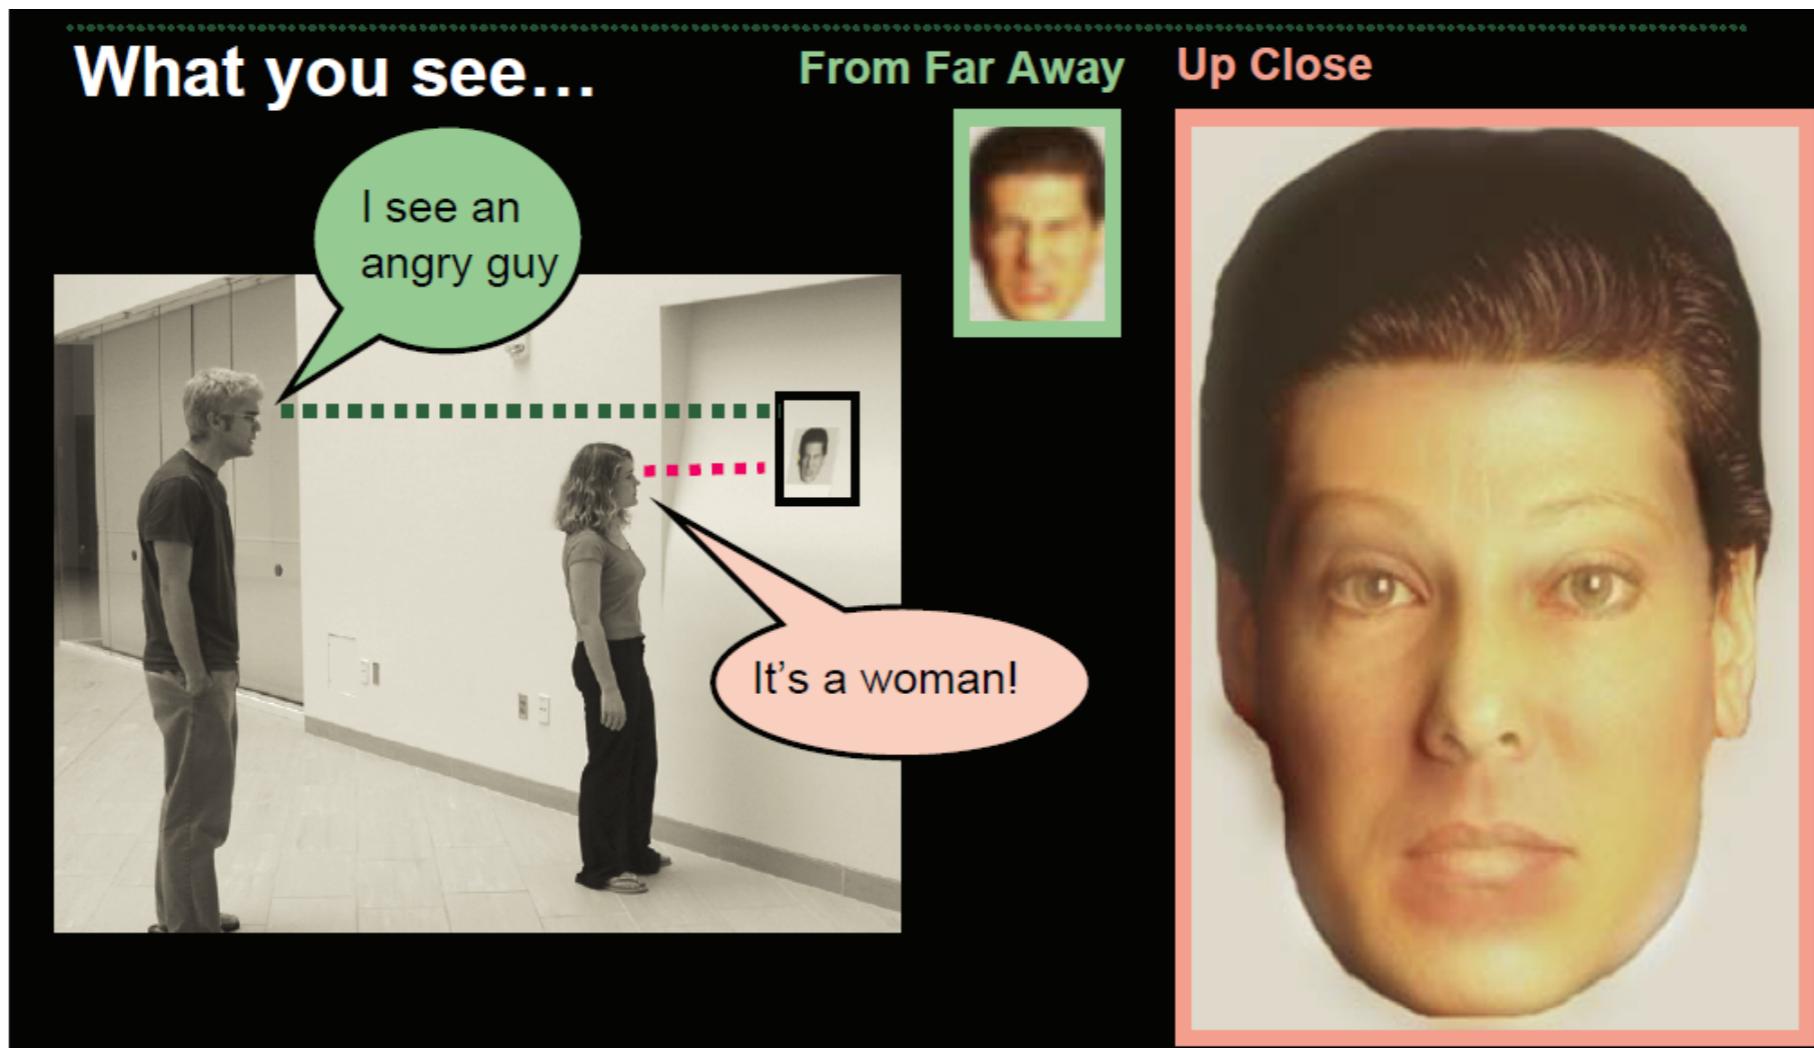
D



E



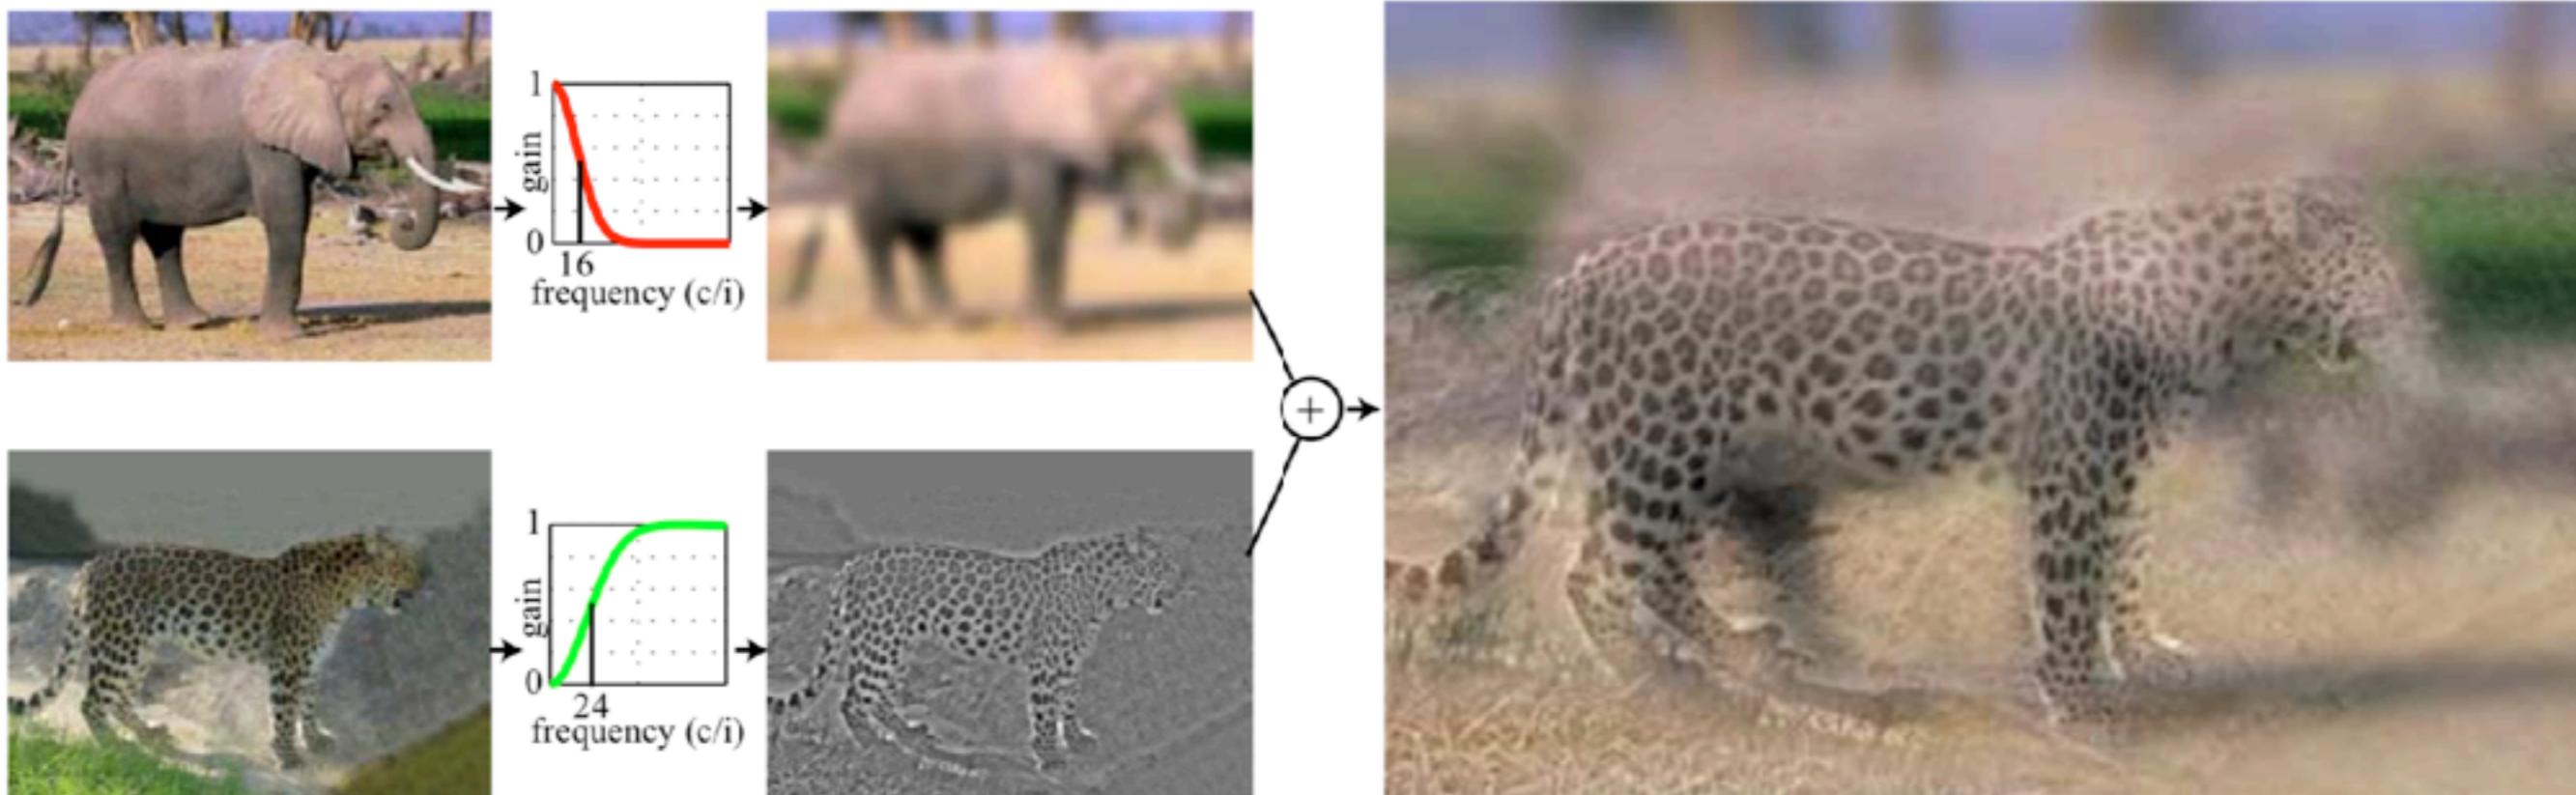
Application: hybrid image



Oliva et al [SIGGRAPH 2016]

Hybrid image: main idea

- Generated by superimposing two images at two different spatial scales
- The low-spatial scale is obtained by filtering one image with a low-pass filter
- The high spatial scale is obtained by filtering a second image with a high-pass filter
- The final hybrid image is composed by adding these two filtered images



Changing expression



Sad

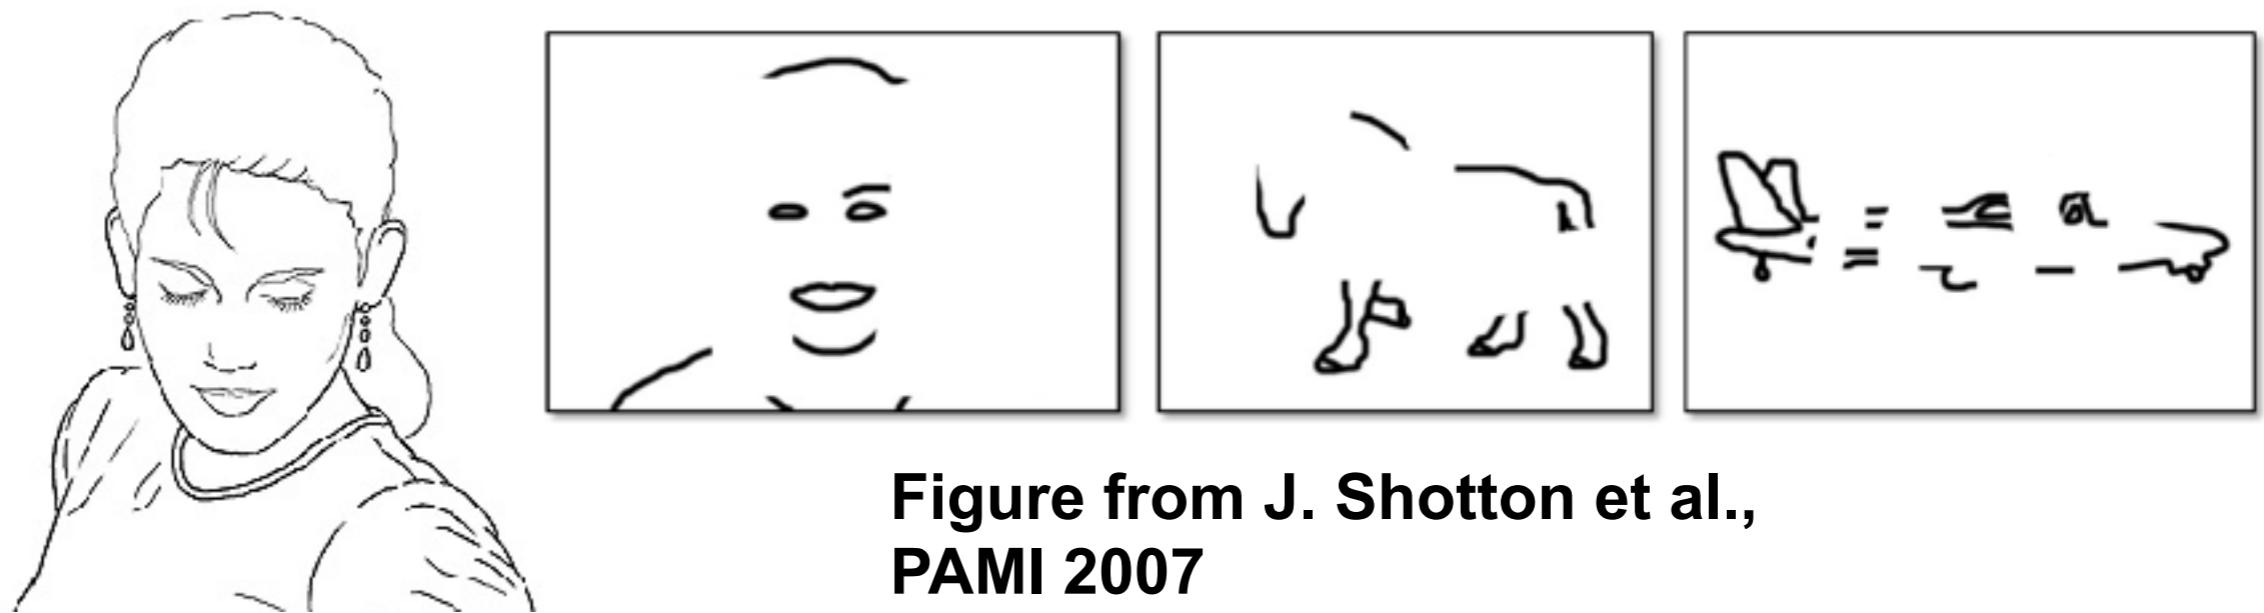
Surprised



Edge Detection

Edge detection

- Goal: map image from 2d array of pixels to a set of curves or line segments or contours.
- Why?



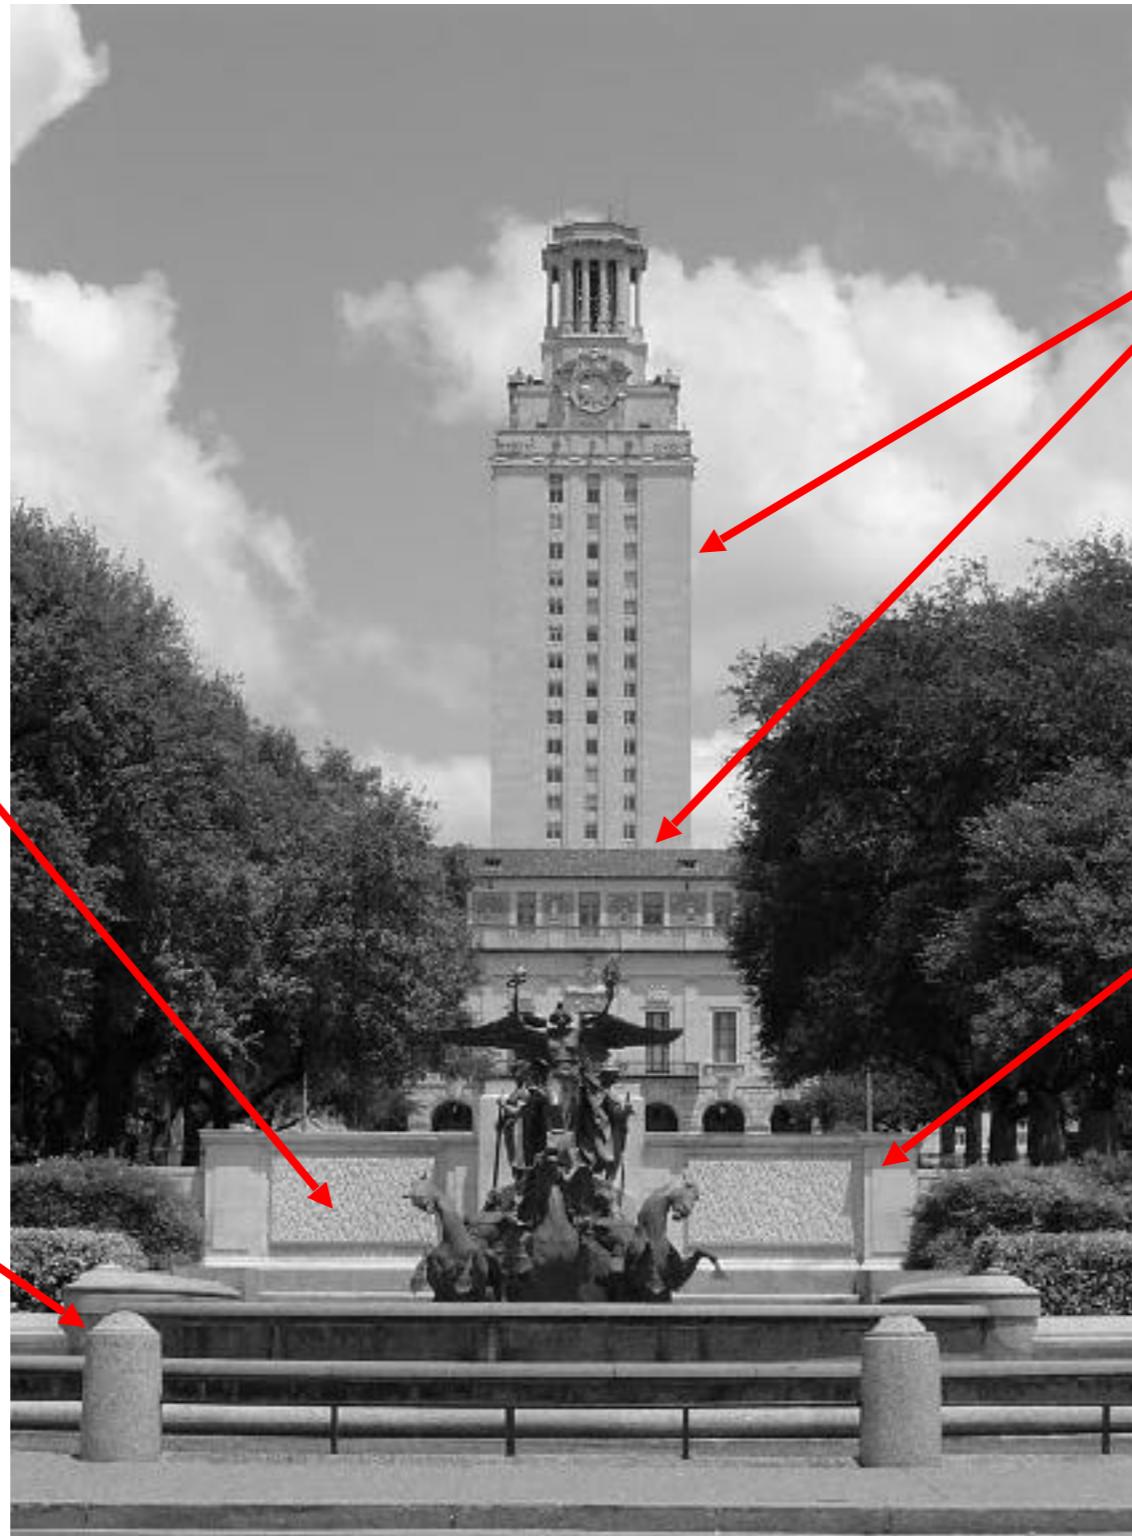
**Figure from J. Shotton et al.,
PAMI 2007**

- Main idea: look for strong gradients, post-process

What causes an edge?

Reflectance change:
appearance information,
texture

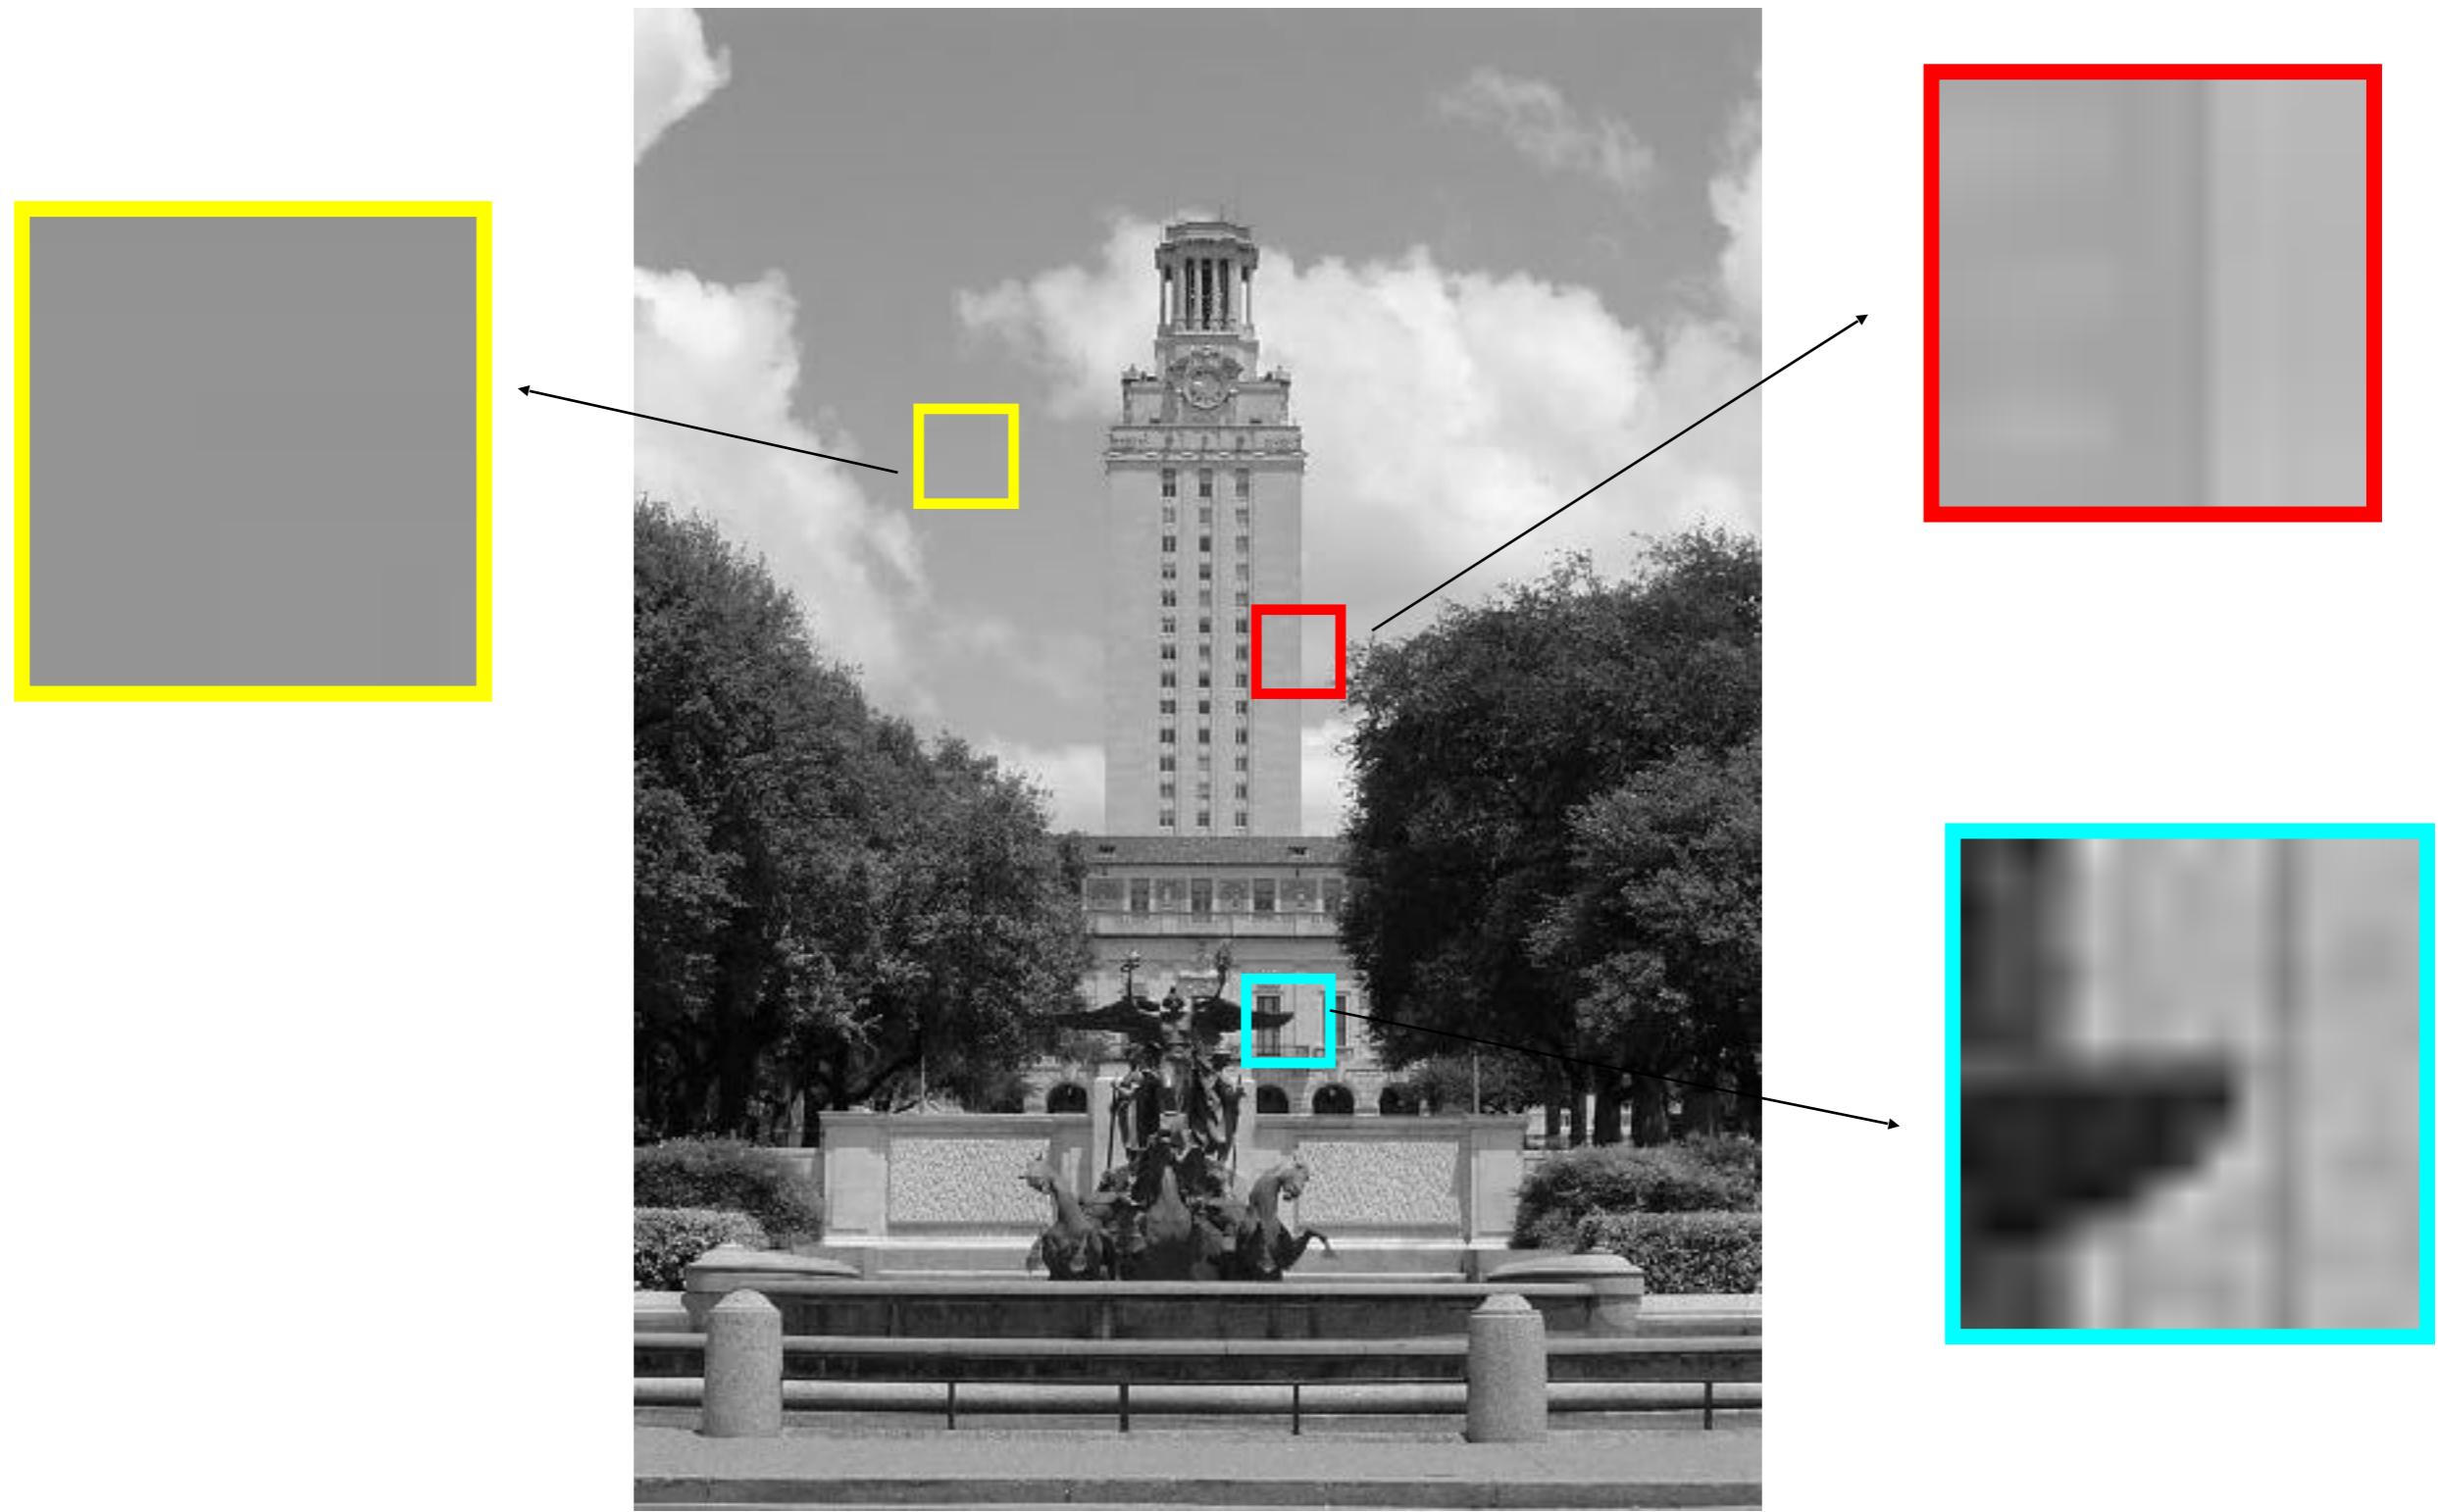
Change in surface
orientation: shape



Depth discontinuity:
object boundary

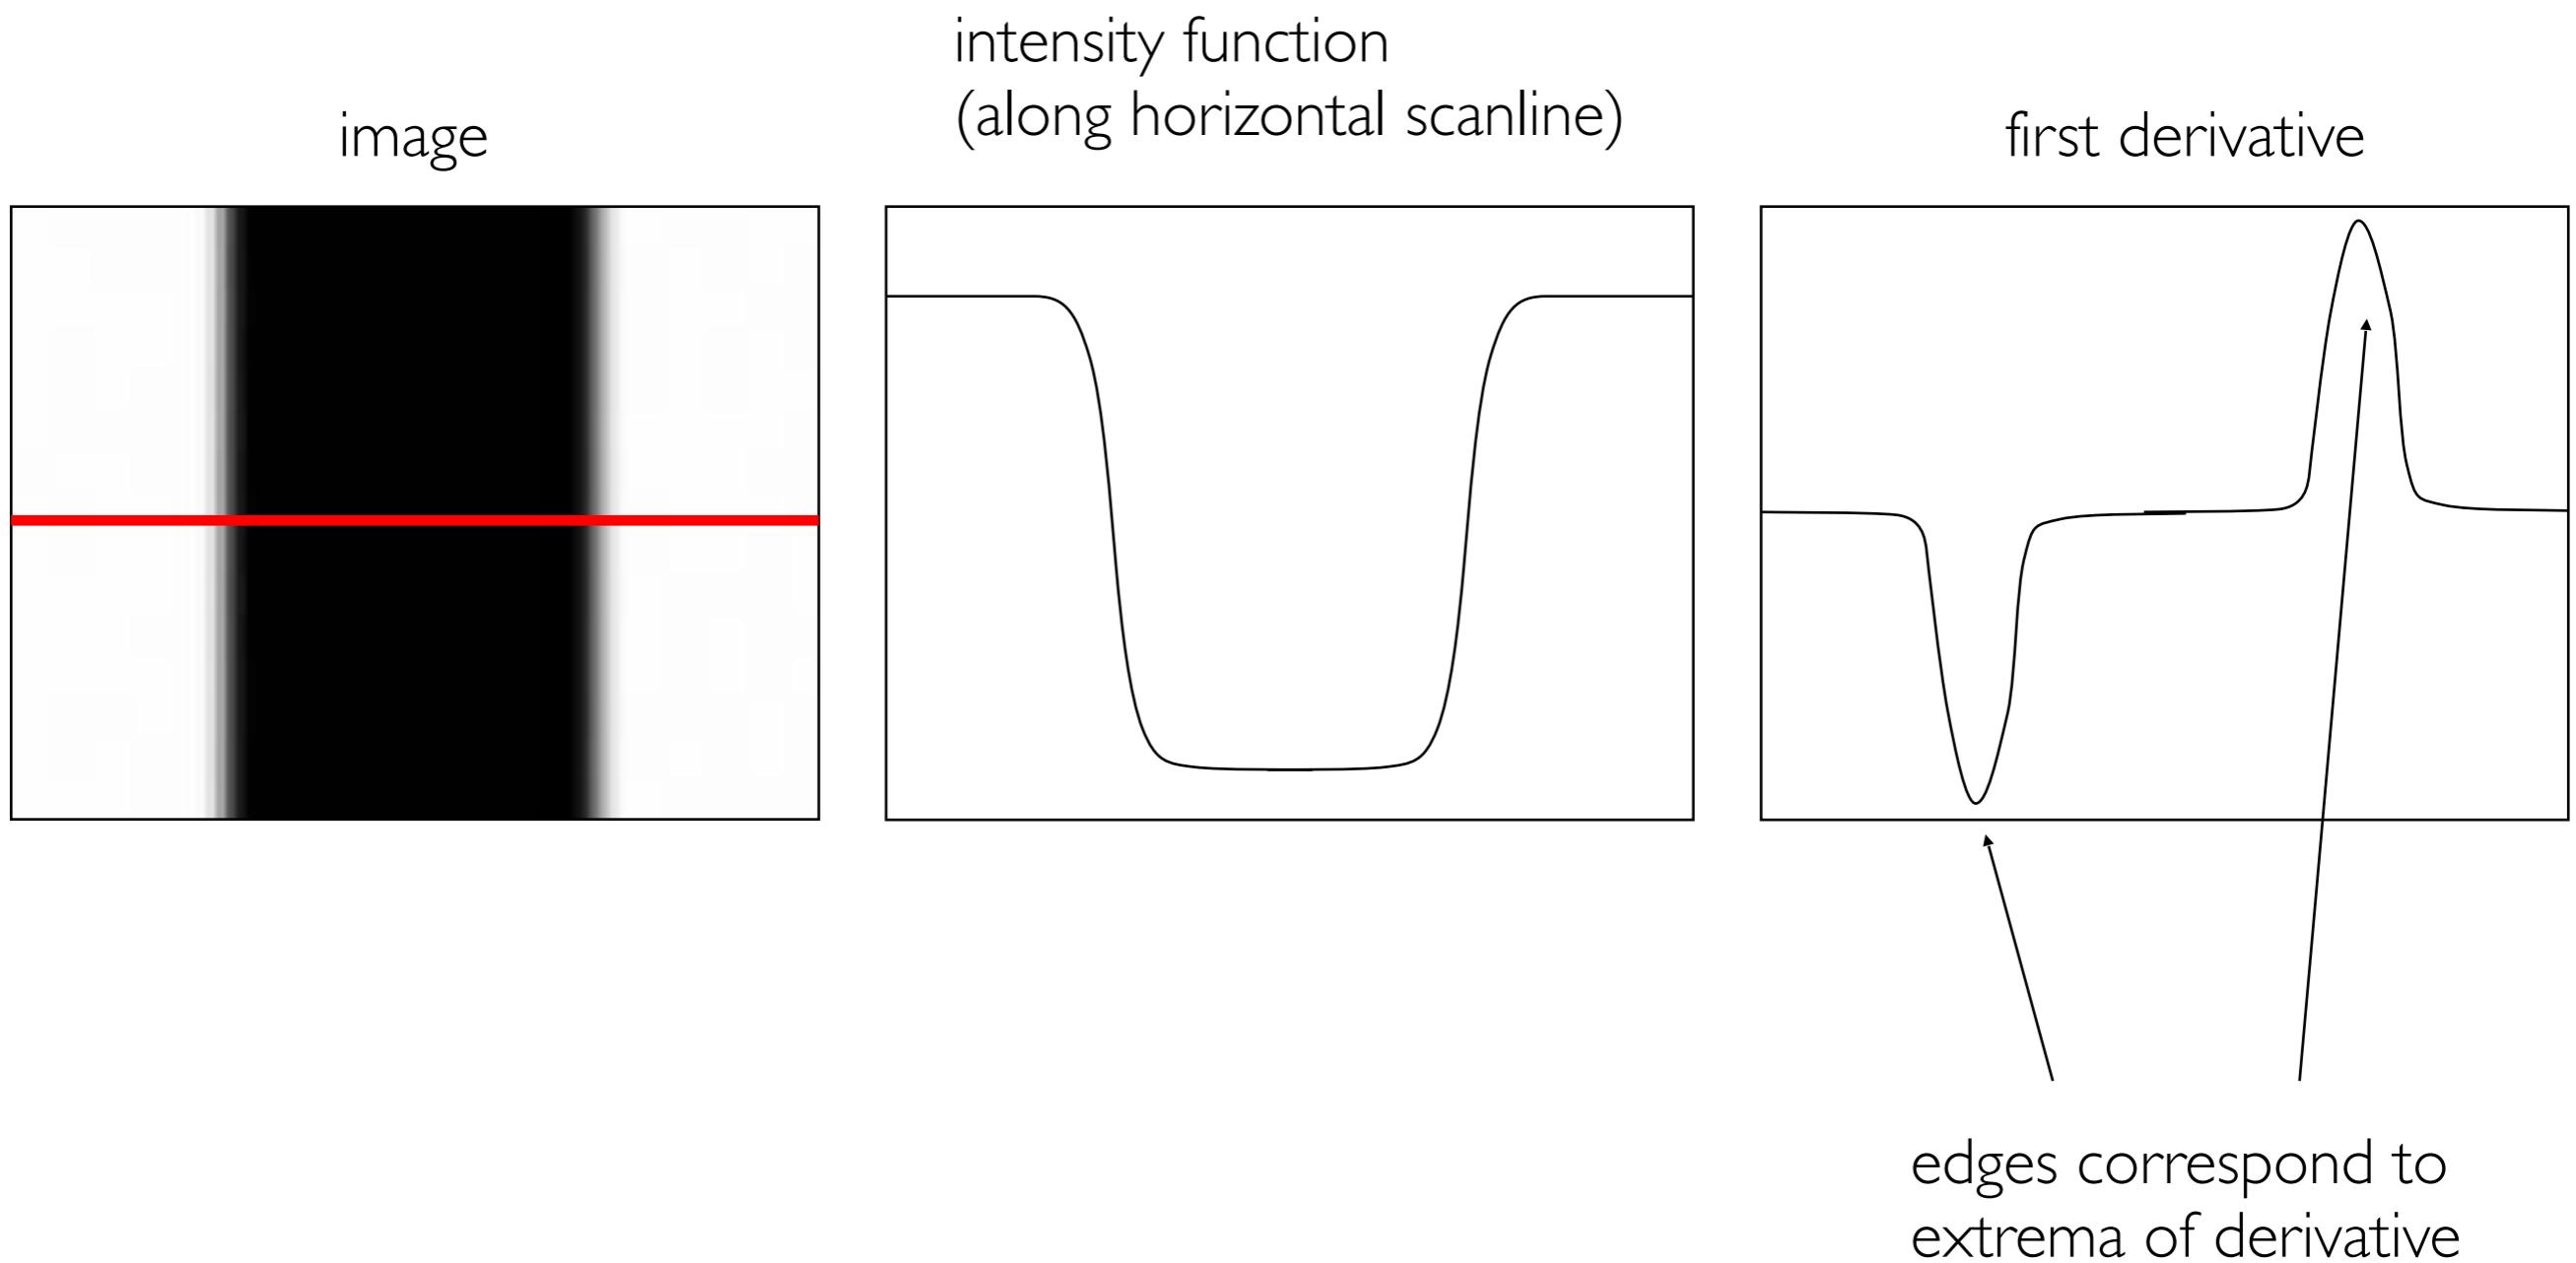
Cast shadows

Edges/gradients and invariance



Derivatives and edges

An edge is a place of rapid change in the image intensity function.



Derivatives with convolution

For 2D function, $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

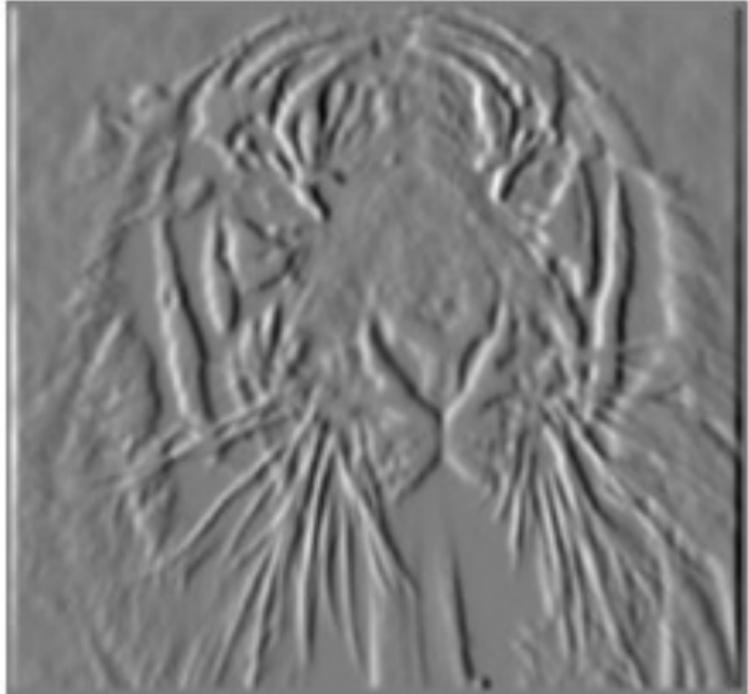
To implement above as convolution, what would be the associated filter?

Partial derivatives of an image

$$\frac{\partial f(x, y)}{\partial x}$$

∂x

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

∂y

-1
1

?

or

1
-1

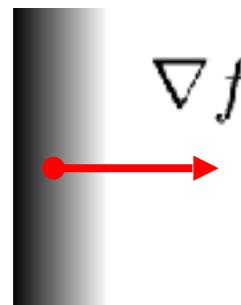
Which shows changes with respect to x?
(showing filters for correlation)

Image gradient

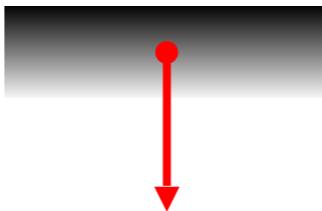
The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

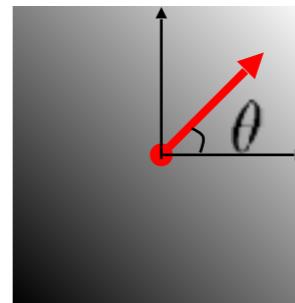
The gradient points in the direction of most rapid change in intensity



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$



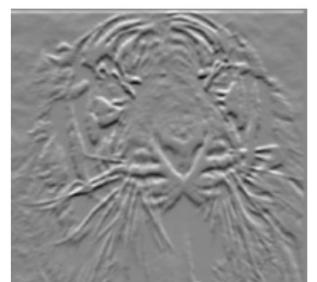
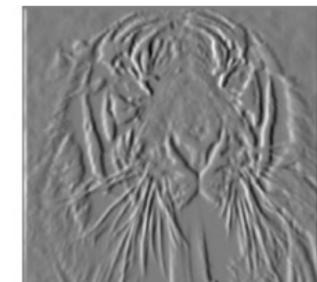
$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient direction (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The edge strength is given by the gradient magnitude

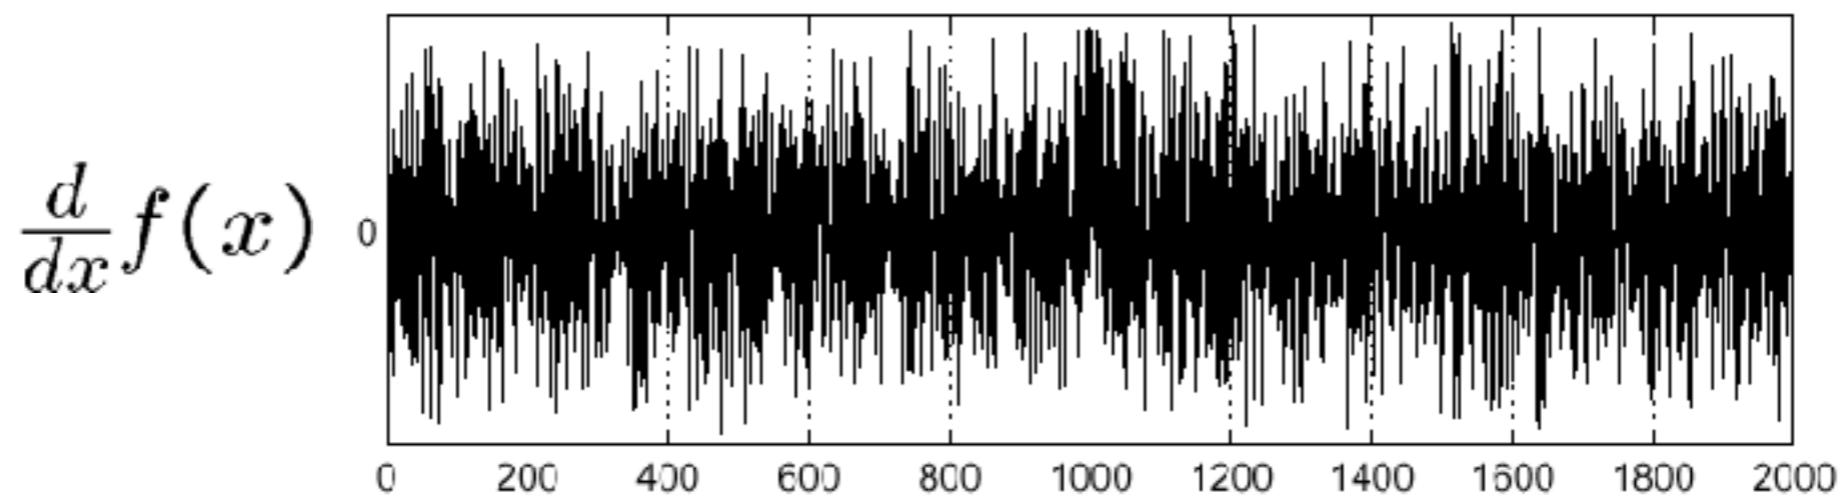
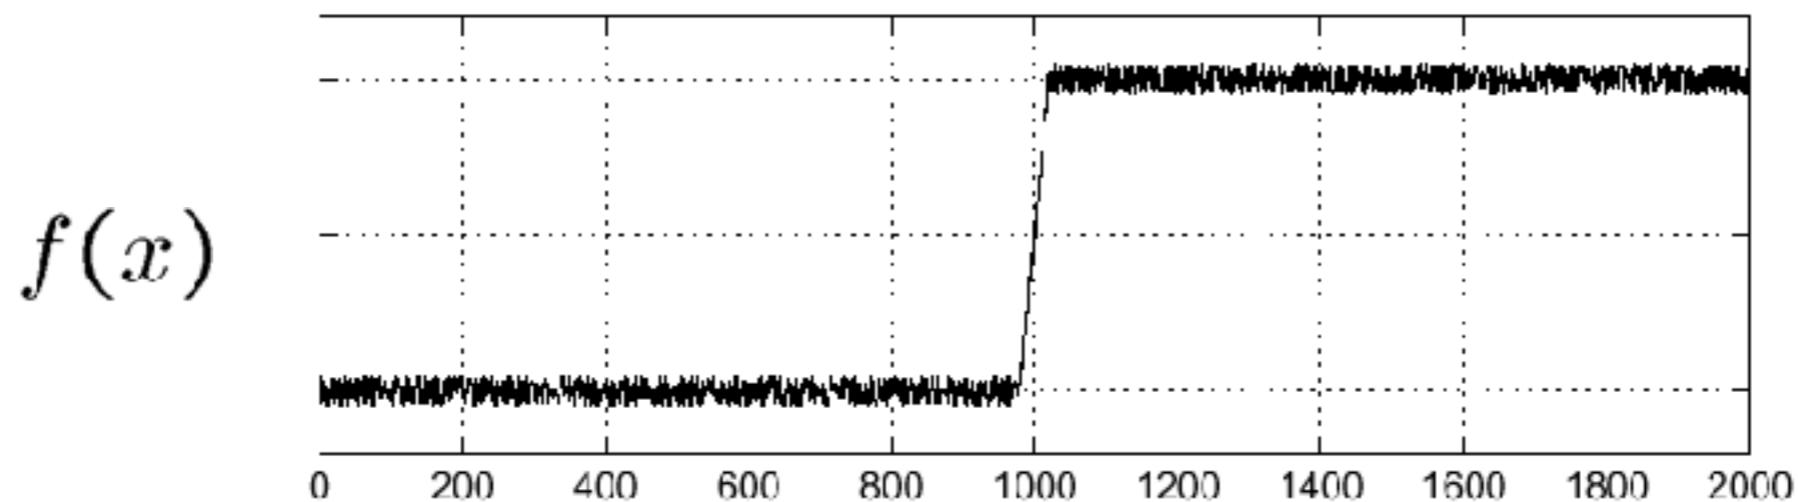
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$



Effects of noise

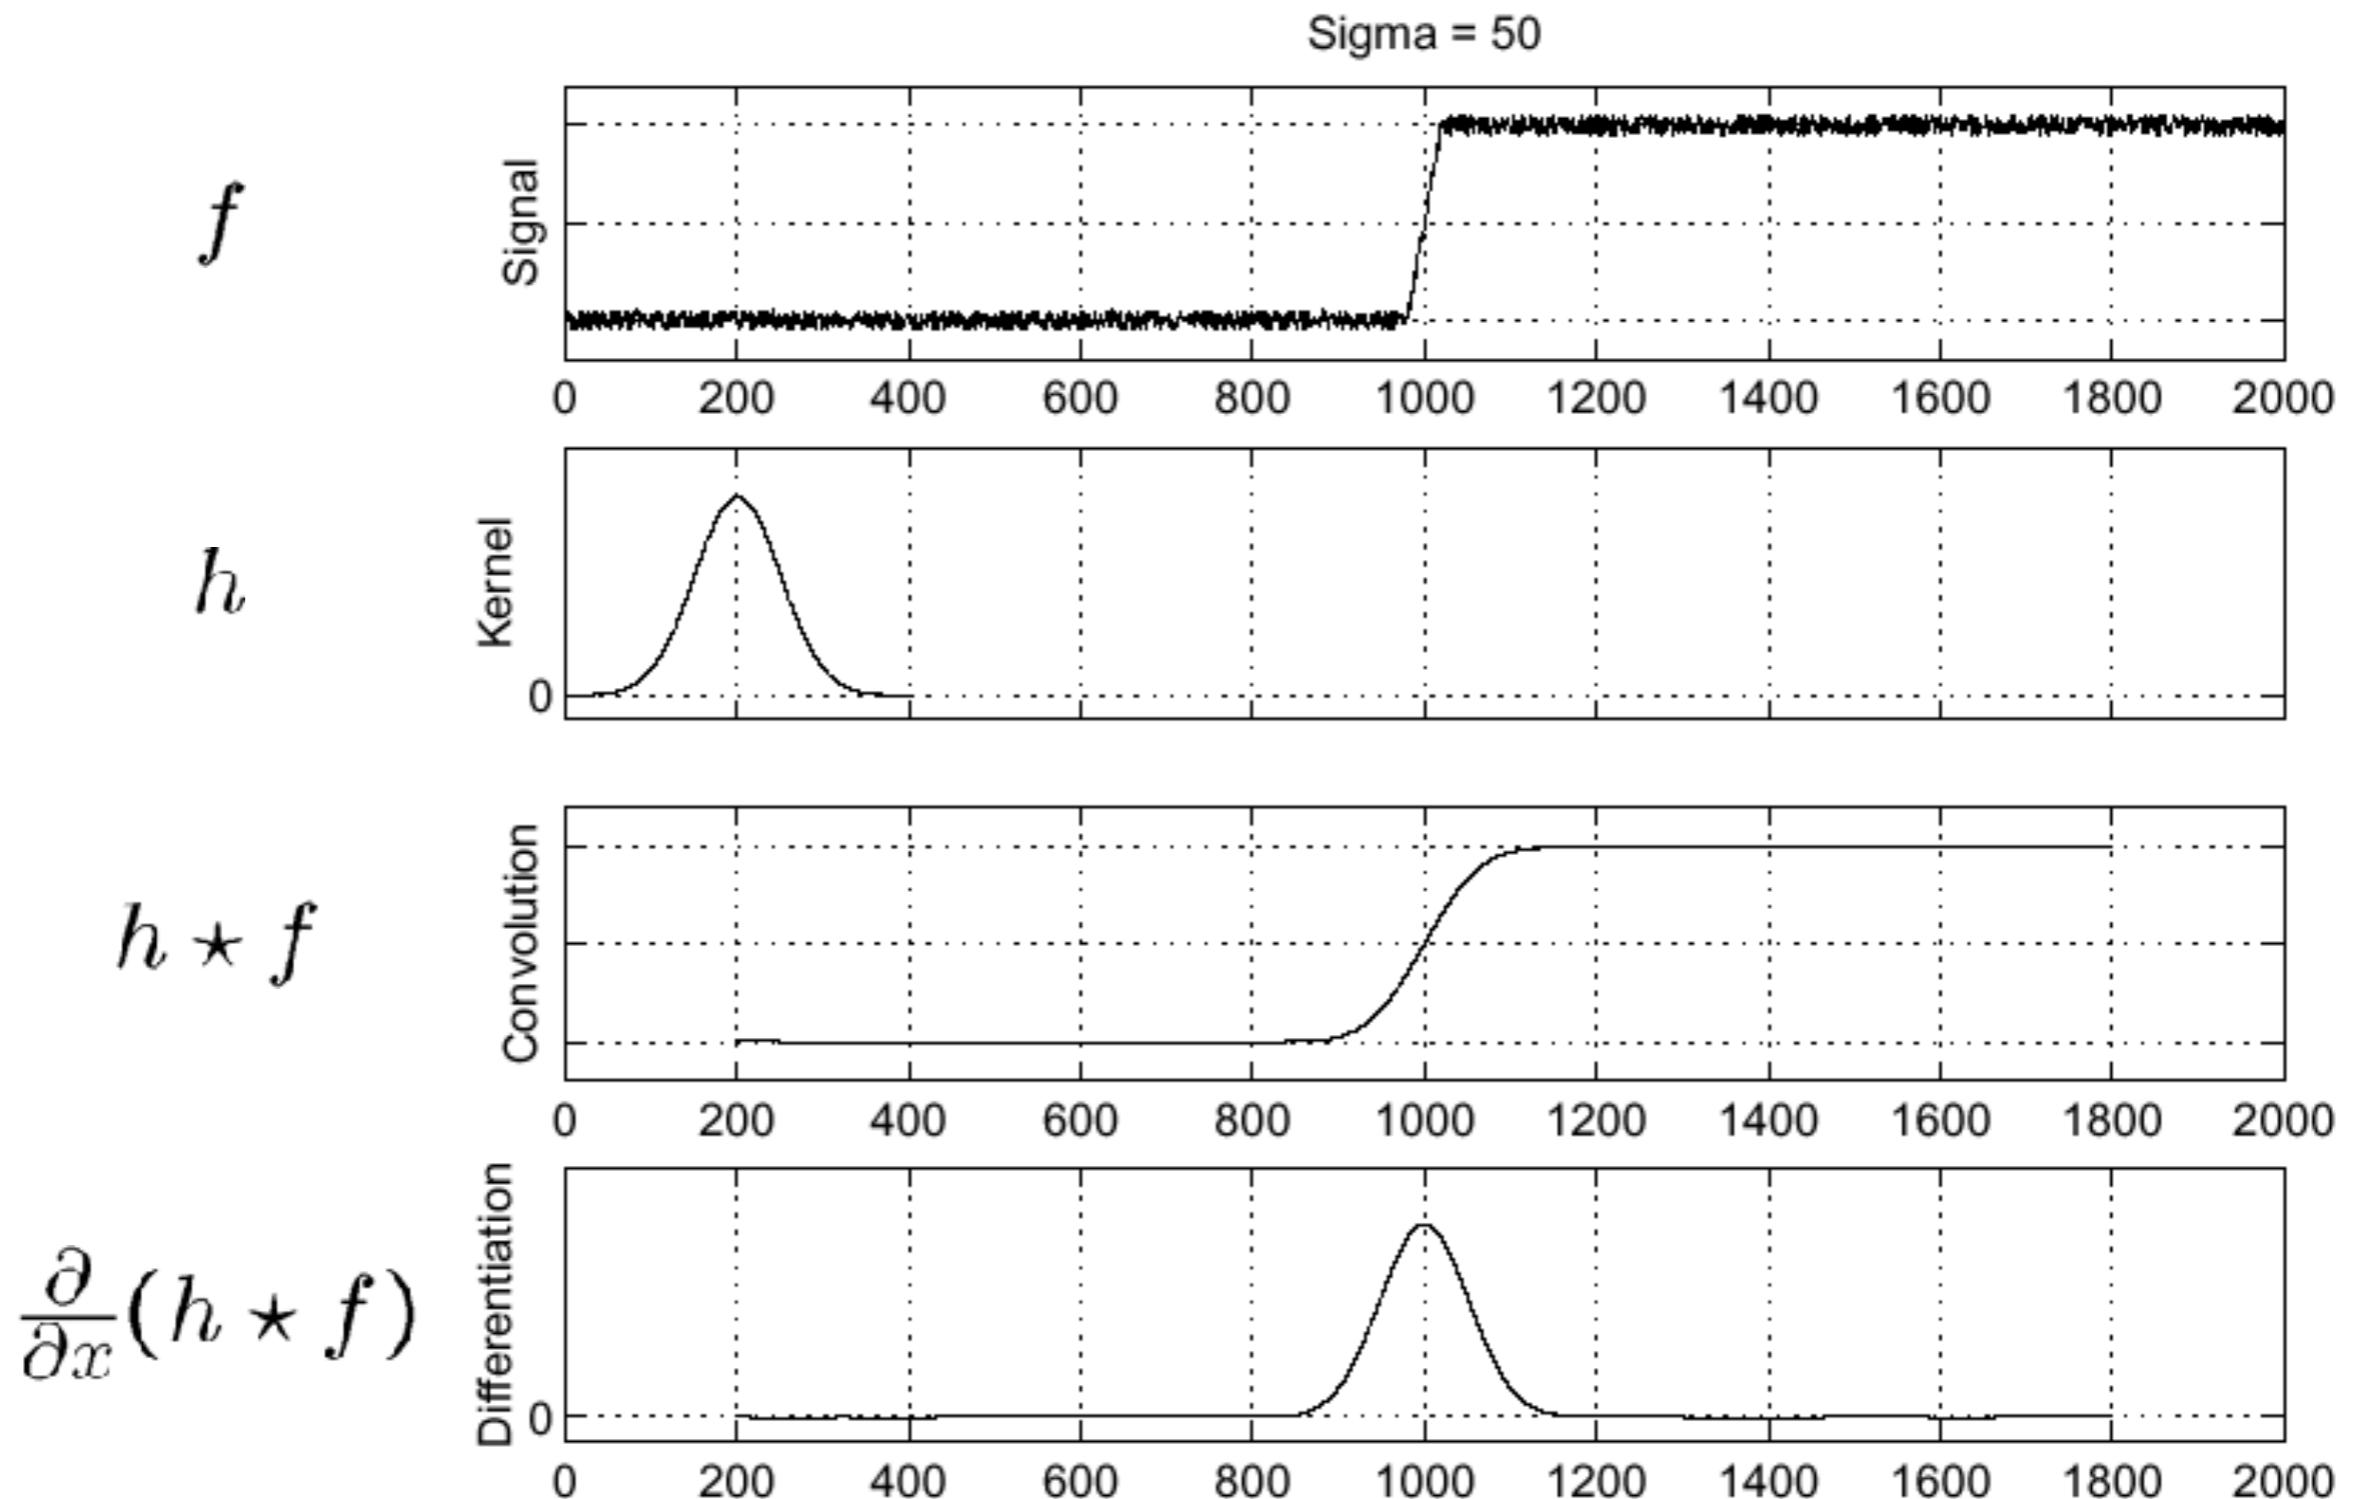
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



Where is the edge?

Solution: smooth first!



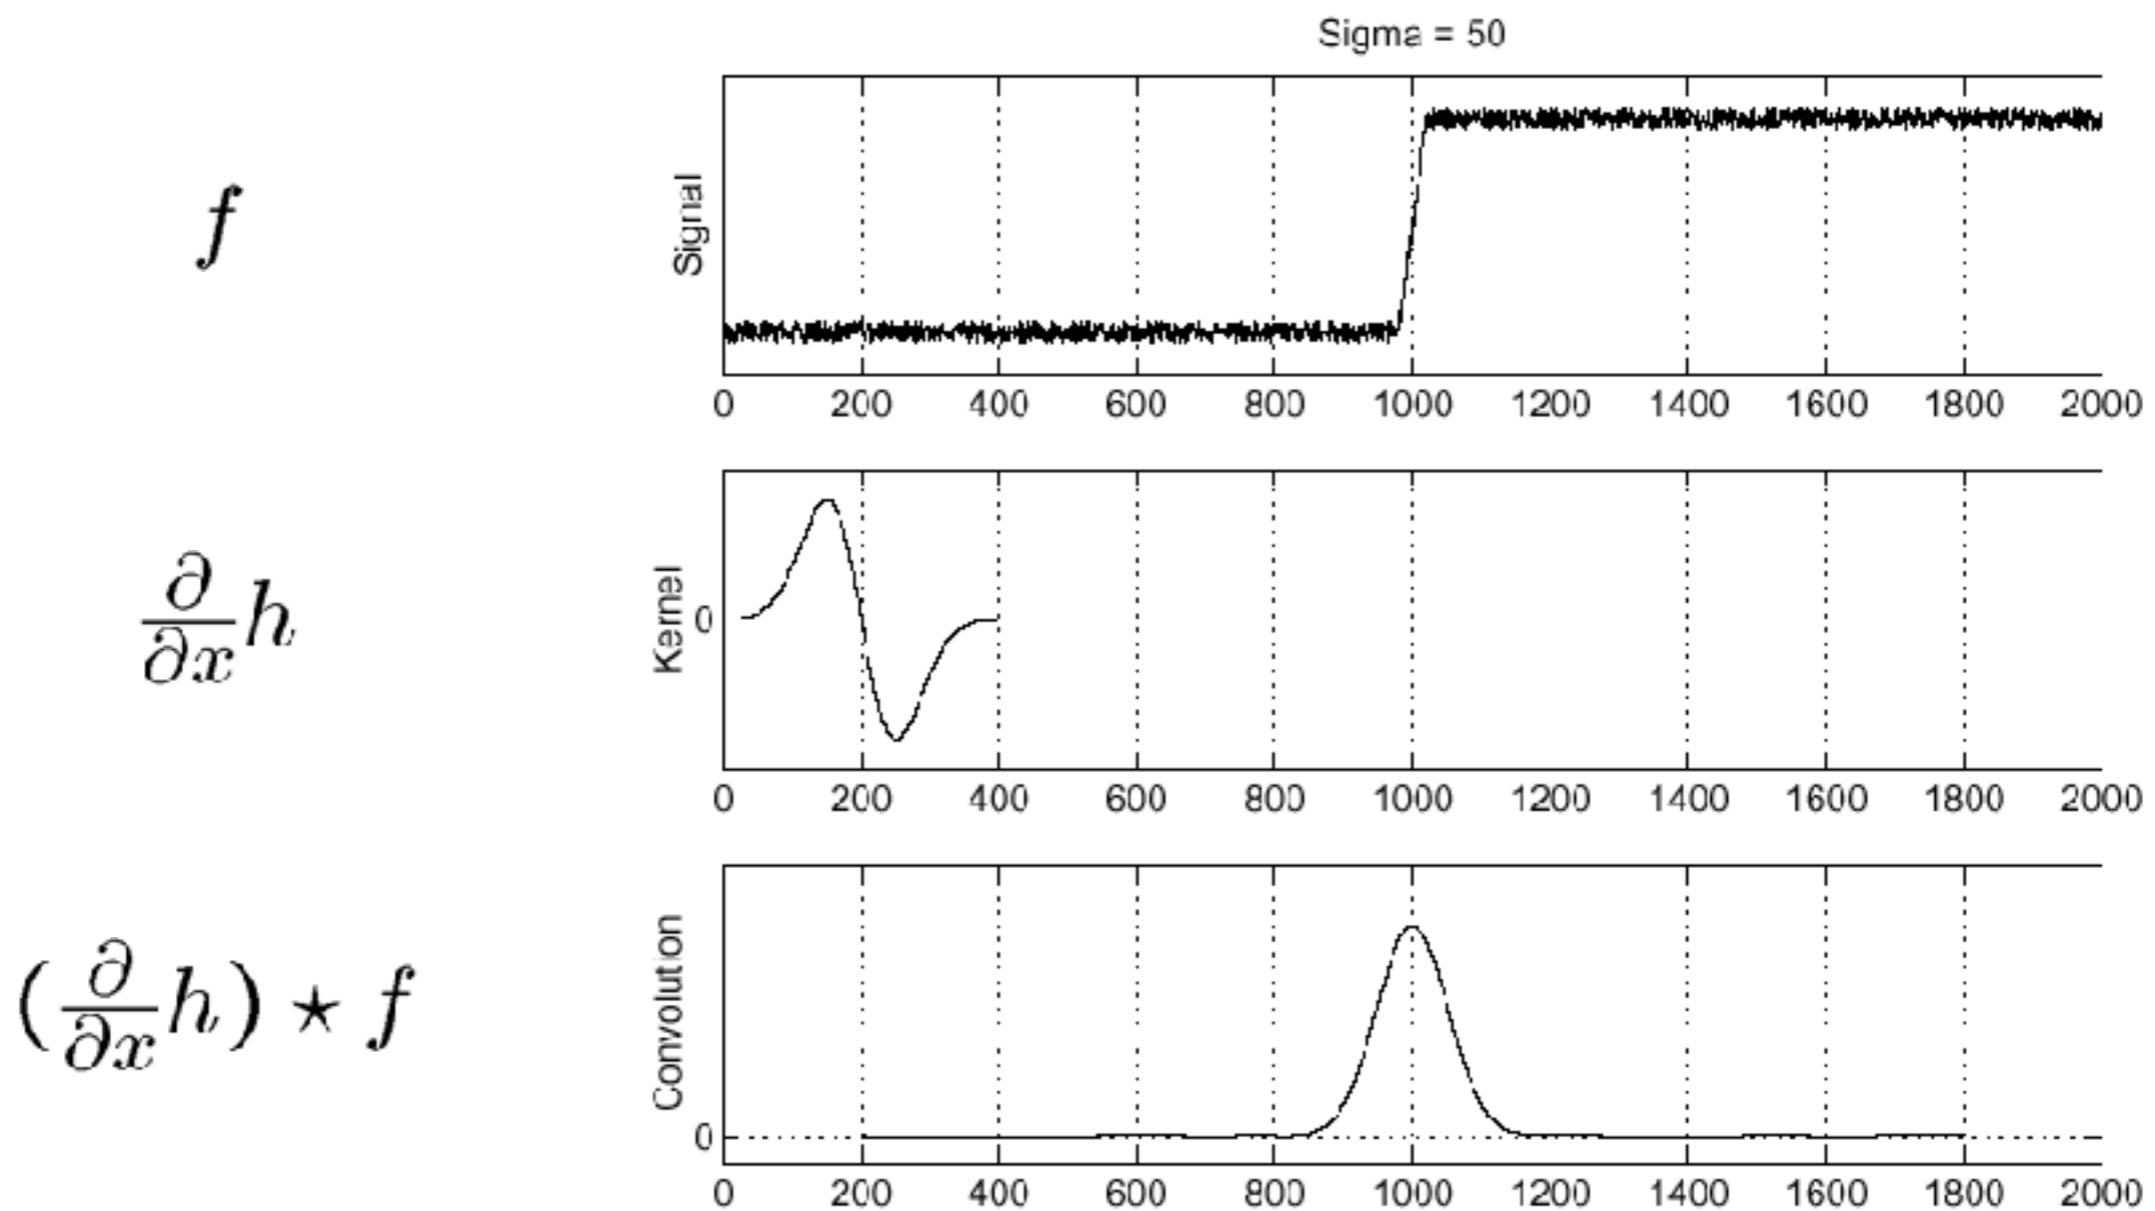
Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

Differentiation property of convolution



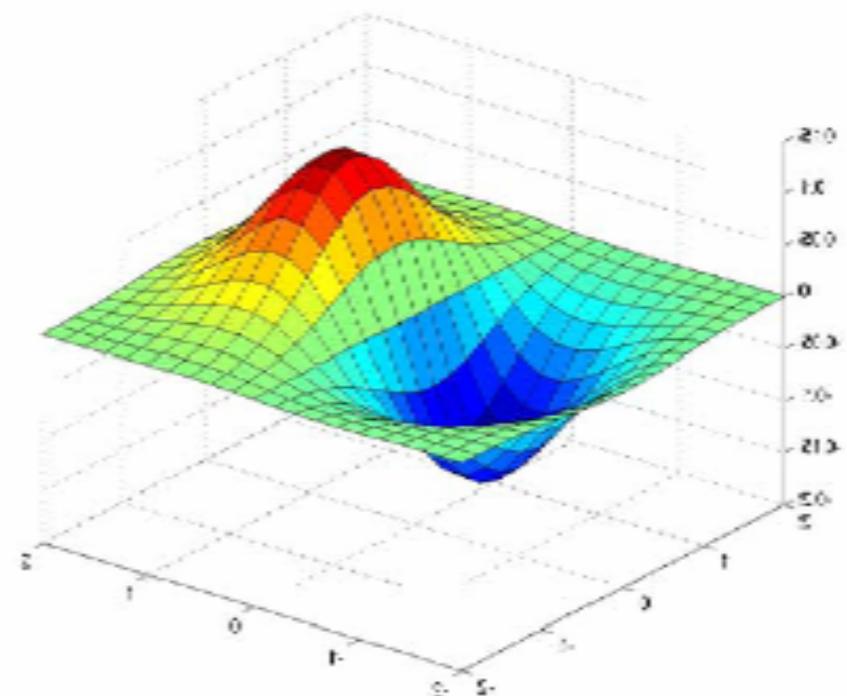
Derivative of Gaussian filters

$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$

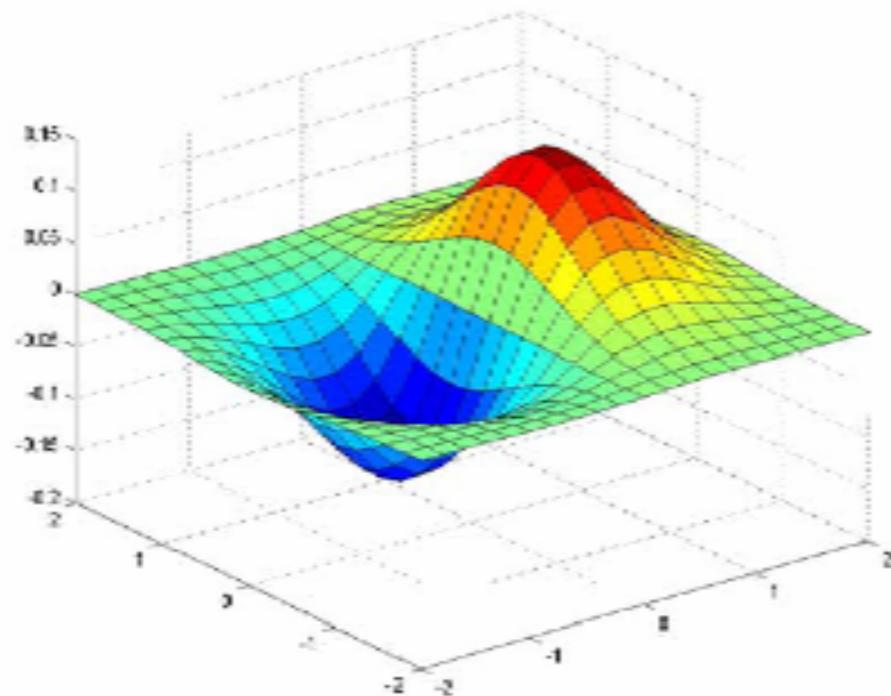
$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$

$$\otimes$$

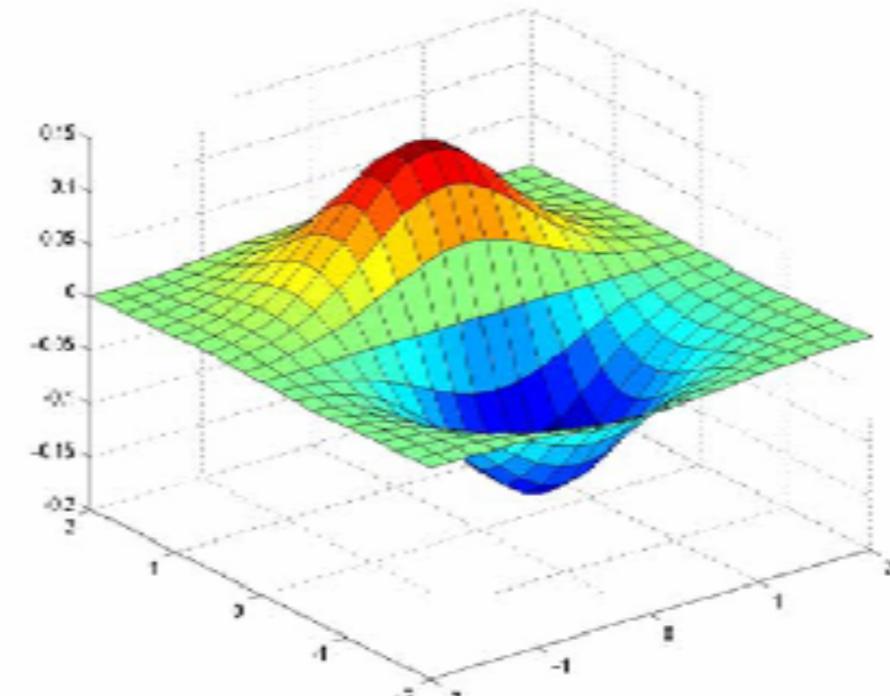
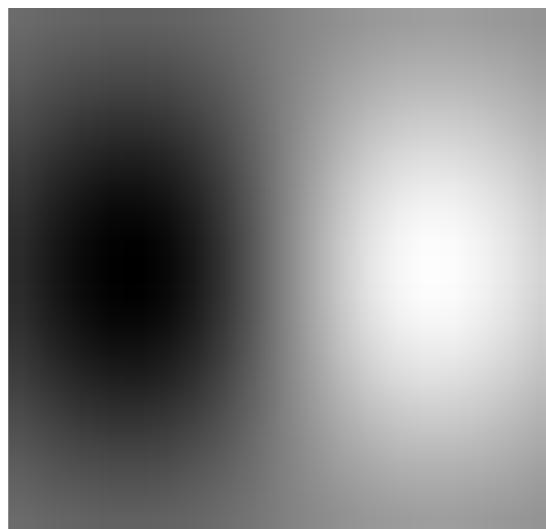
$$\begin{bmatrix} 1 & -1 \end{bmatrix}$$



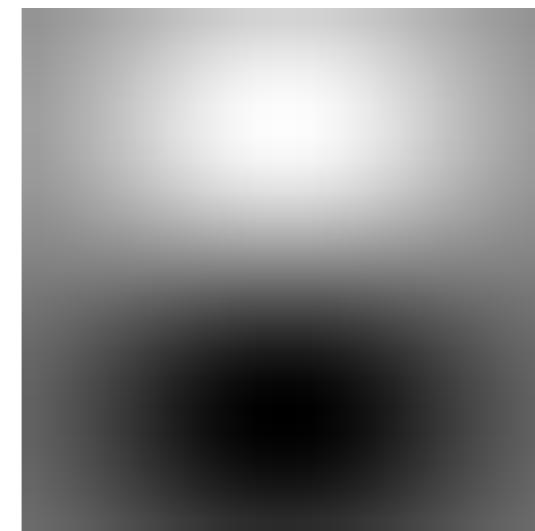
Derivative of Gaussian filters



x-direction



y-direction

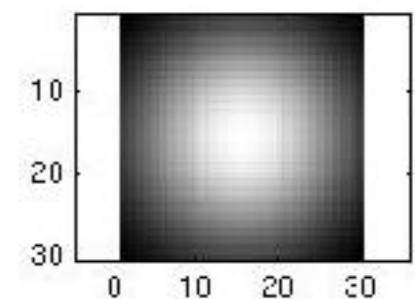
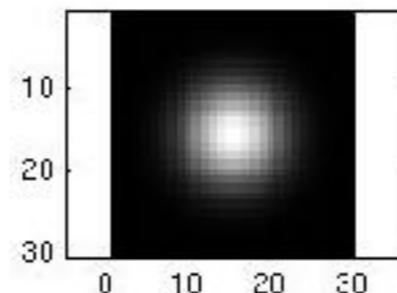
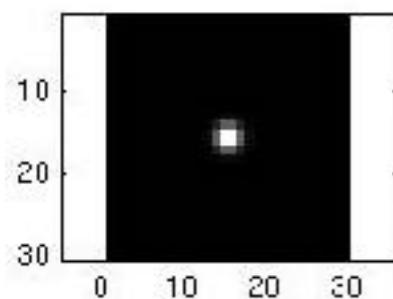


Smoothing with a Gaussian

Recall: parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



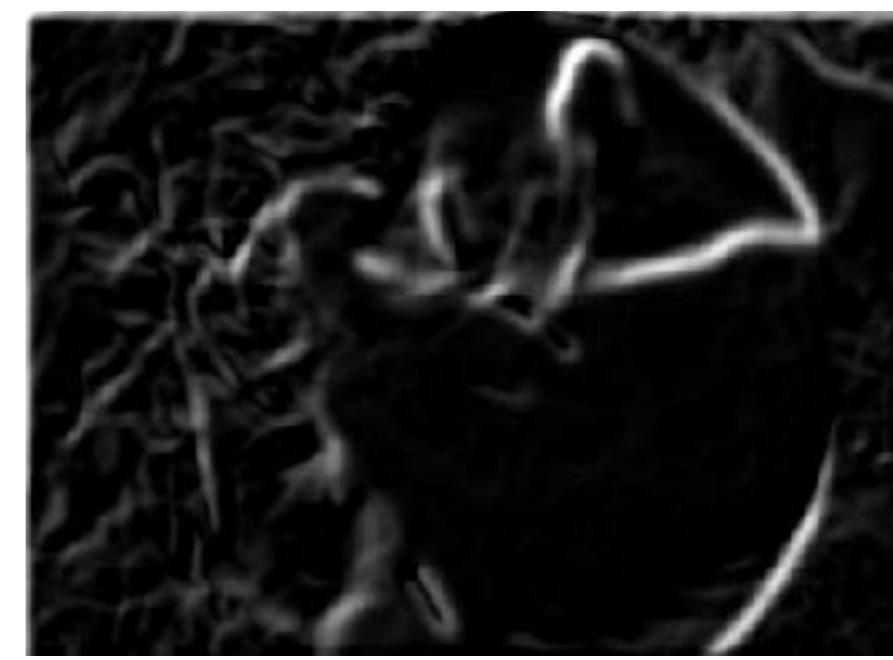
...



Effect of σ on derivatives



$\sigma = 1$ pixel



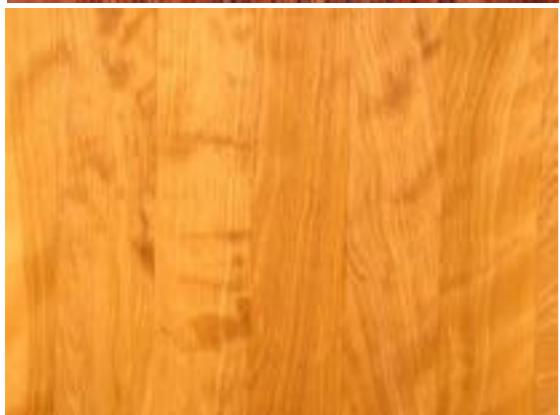
$\sigma = 3$ pixels

The apparent structures differ depending on Gaussian's scale parameter.

Larger values: larger scale edges detected
Smaller values: finer features detected

So, what scale to choose?

It depends on what we're looking for!





Gradients → Edges



Primary edge detection steps:

1. Smoothing: suppress noise
2. Edge enhancement: filter for contrast
3. Edge localization

Determine which local maxima from filter output are actually edges vs. noise

- Threshold, Thin

Thresholding

- Choose a threshold value t
- Set any pixels less than t to zero (off)
- Set any pixels greater than or equal to t to one (on)

Original image



Gradient magnitude image



Thresholding gradient with a lower threshold



Thresholding gradient with a higher threshold



Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- Non-maximum suppression:
 - Thin wide “ridges” down to single pixel width
- Linking and thresholding (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

The Canny edge detector



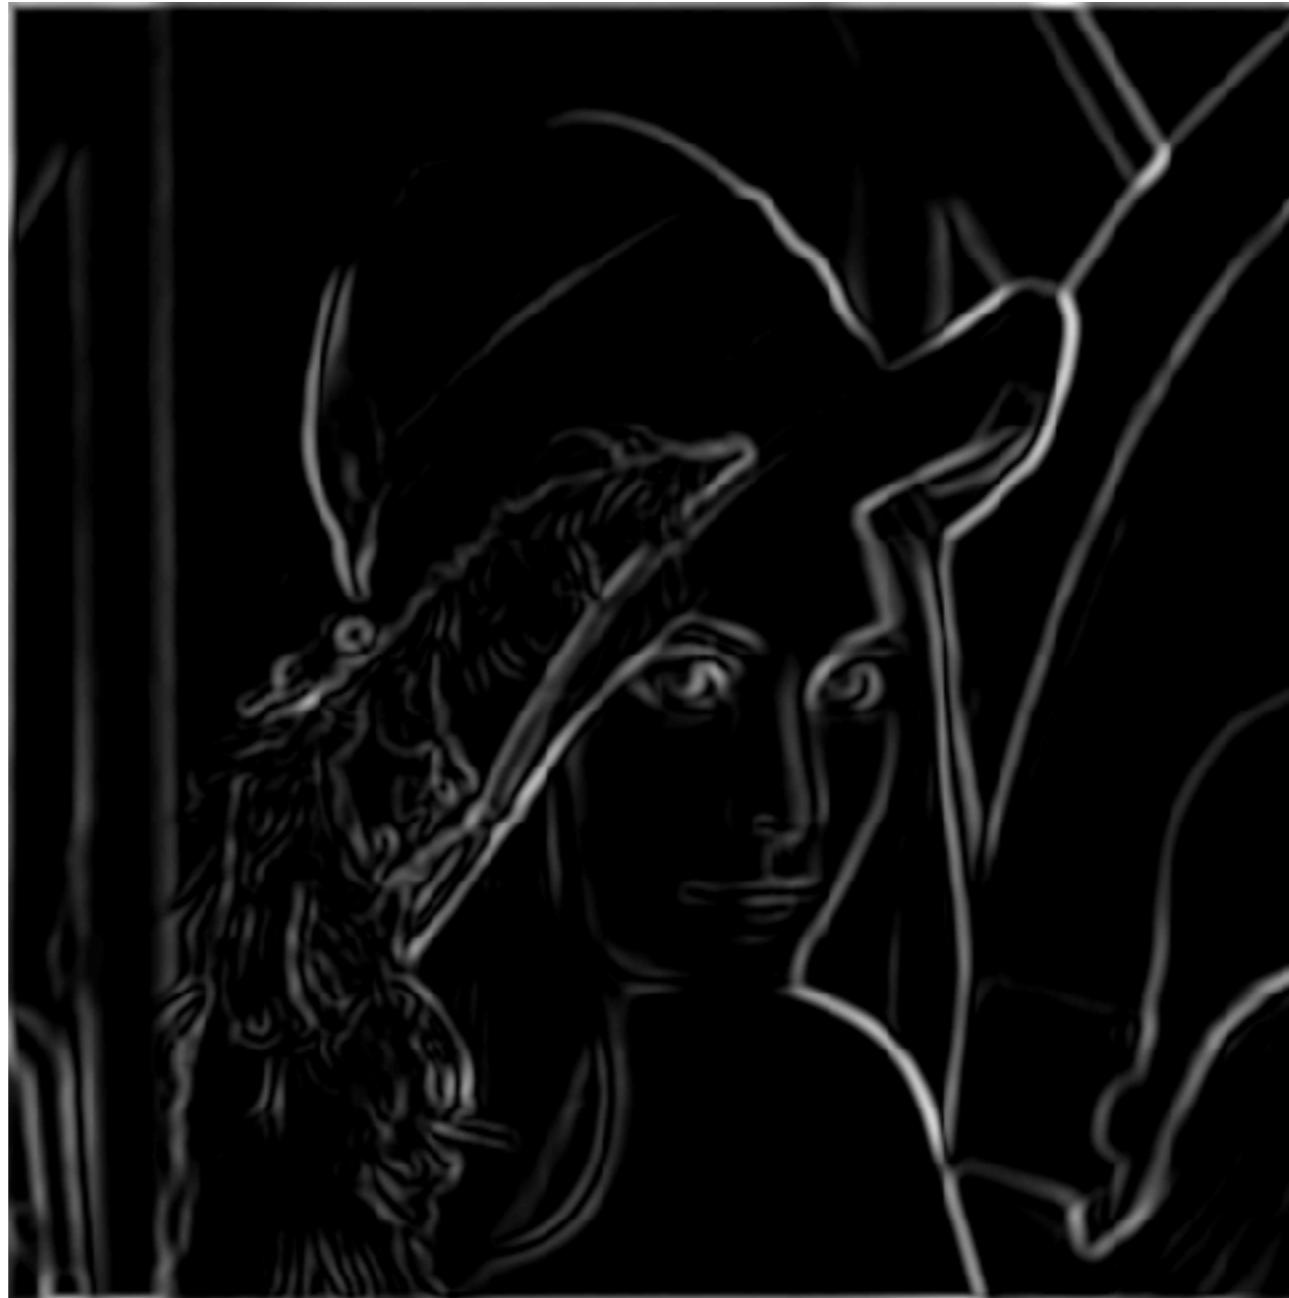
original image (Lena)

The Canny edge detector



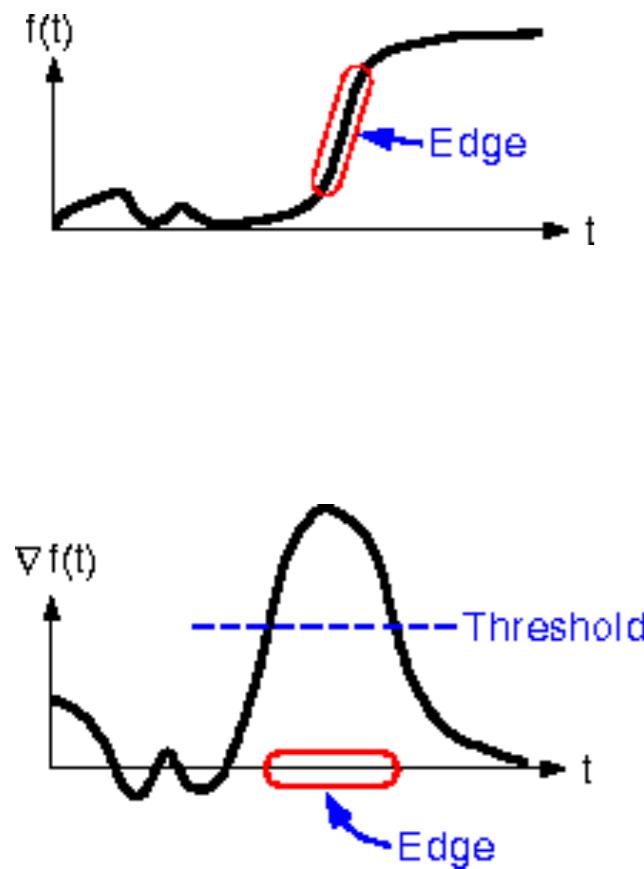
norm of the gradient

The Canny edge detector



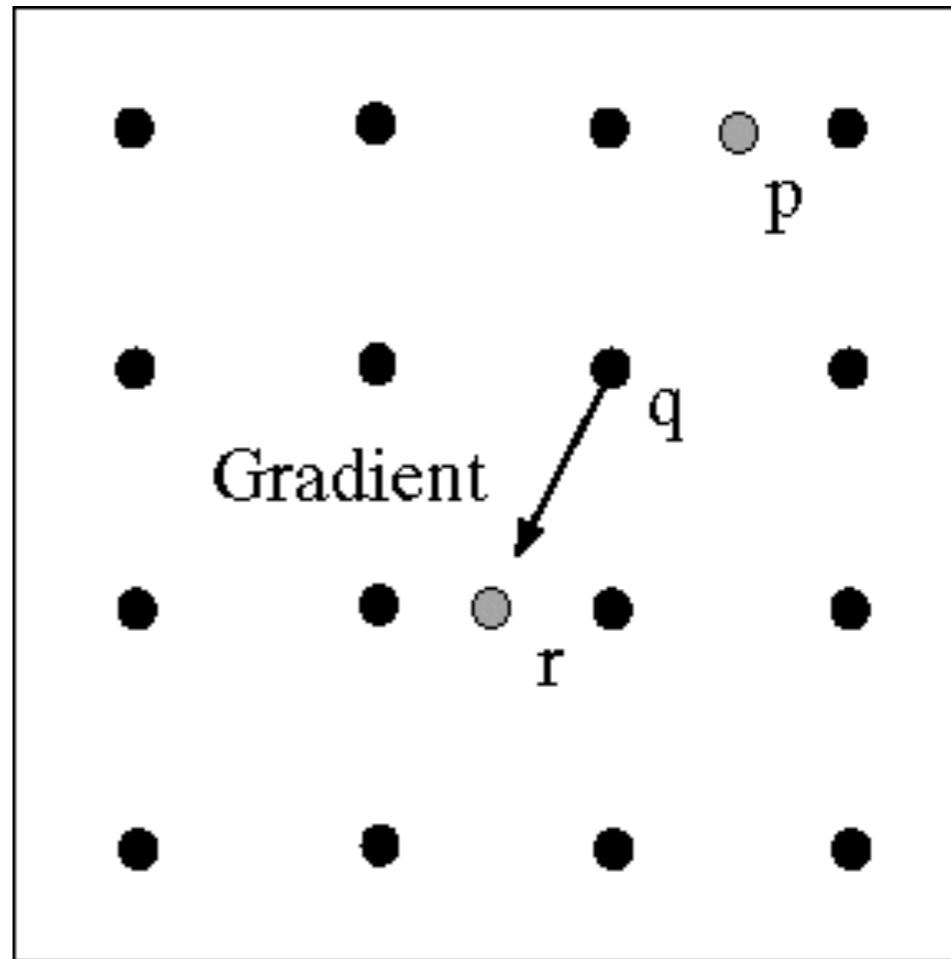
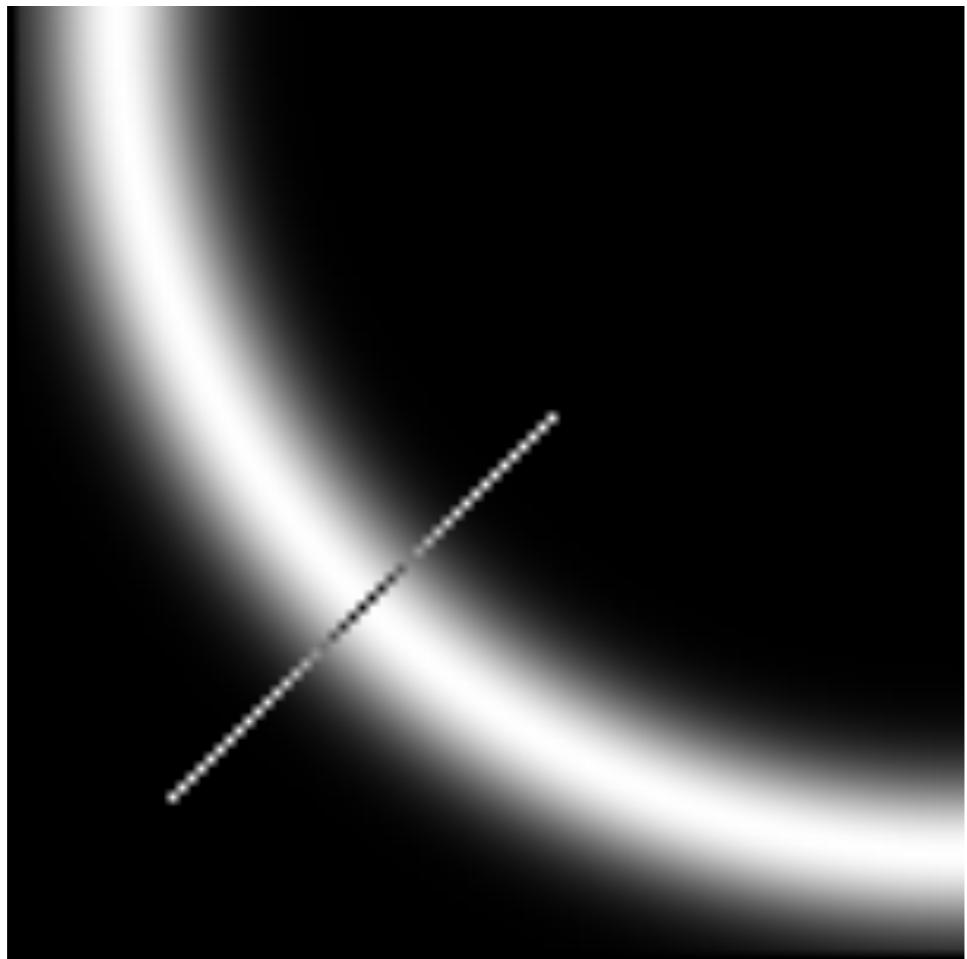
thresholding

The Canny edge detector



How to turn these thick regions of the gradient into curves?

Non-maximum suppression



Check if pixel is local maximum along gradient direction, select single max across width of the edge

- requires checking interpolated pixels p and r

The Canny edge detector



thinning
(non-maximum suppression)

Problem: pixels
along this edge
didn't survive
the
thresholding

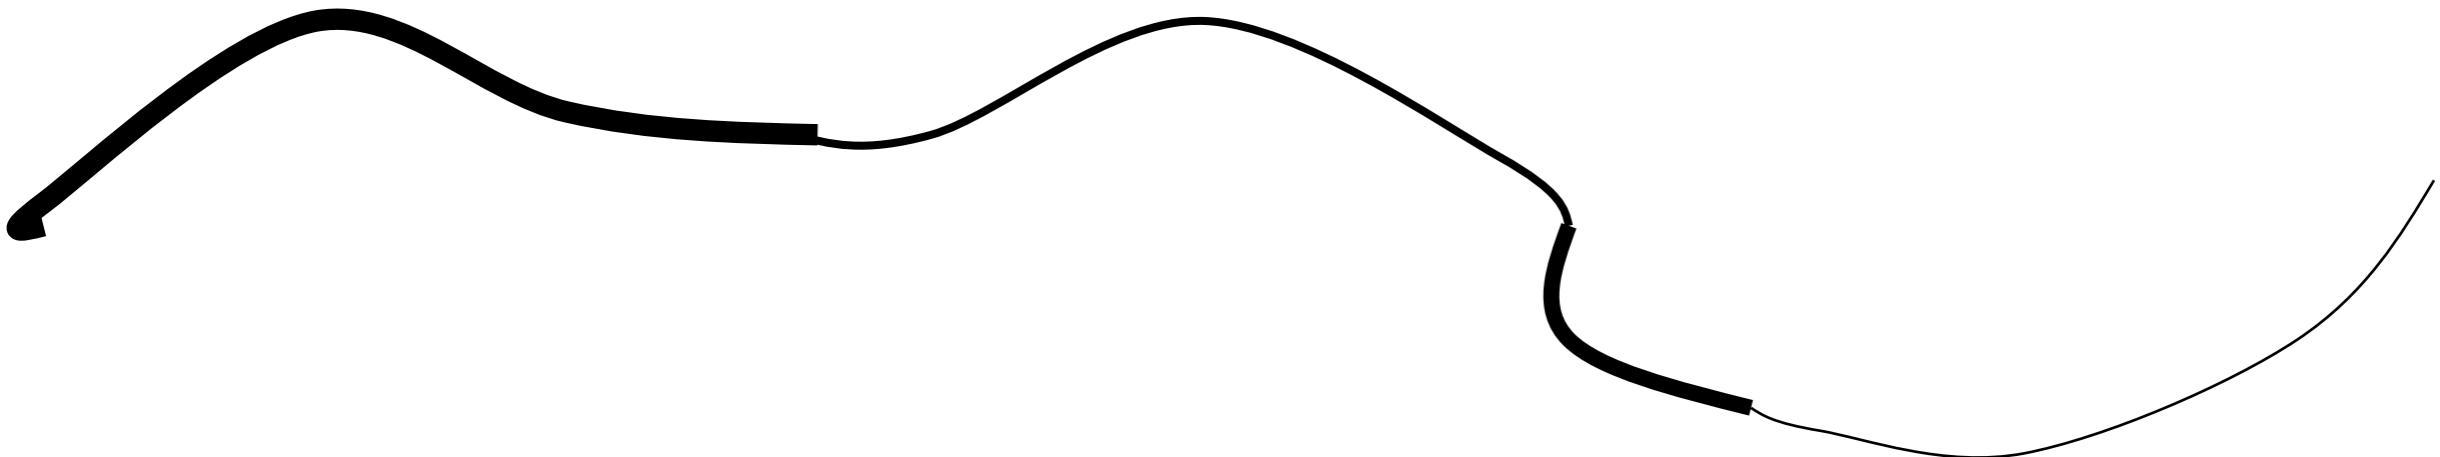
Hysteresis thresholding

- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



Hysteresis thresholding

- Use a high threshold to start edge curves, and a low threshold to continue them.



Canny edge output



Credit: James Hays

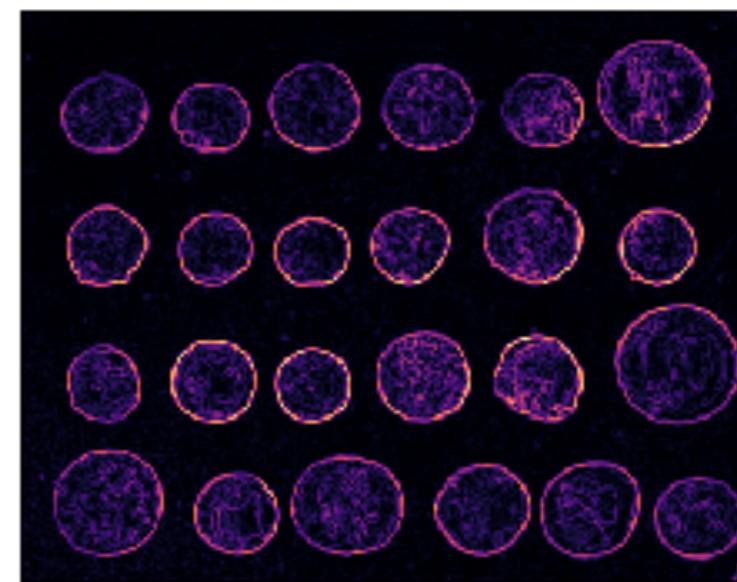
Try it yourself?

Hysteresis detects ‘real’ edges

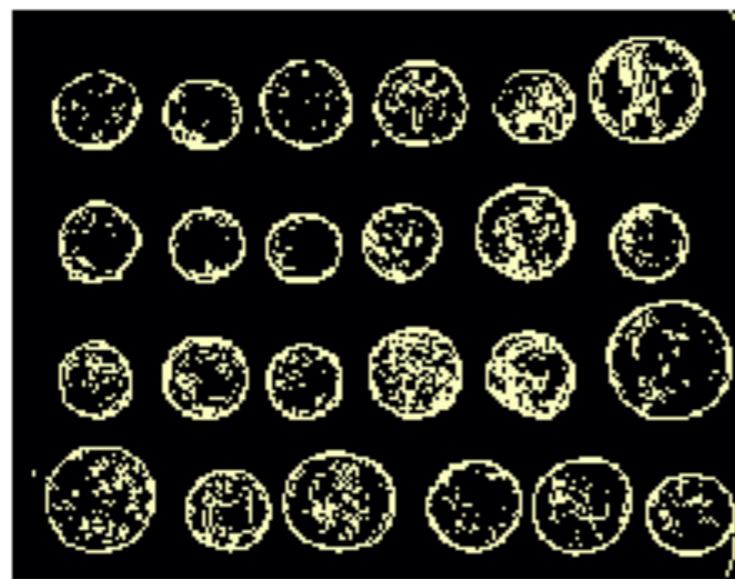
Original image



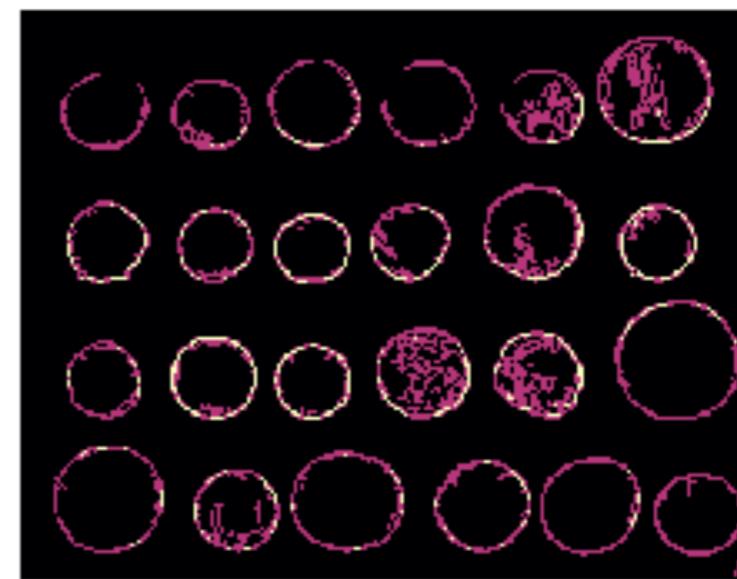
Sobel edges



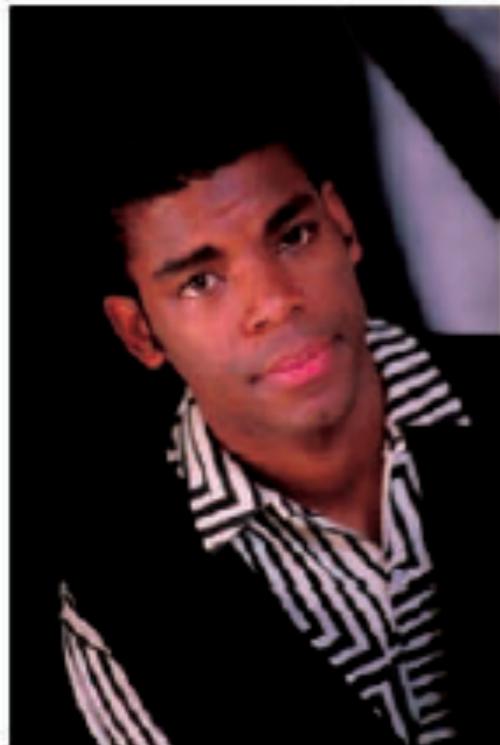
Low threshold



Hysteresis threshold



Low-level edges vs. perceived contours



Background

Texture

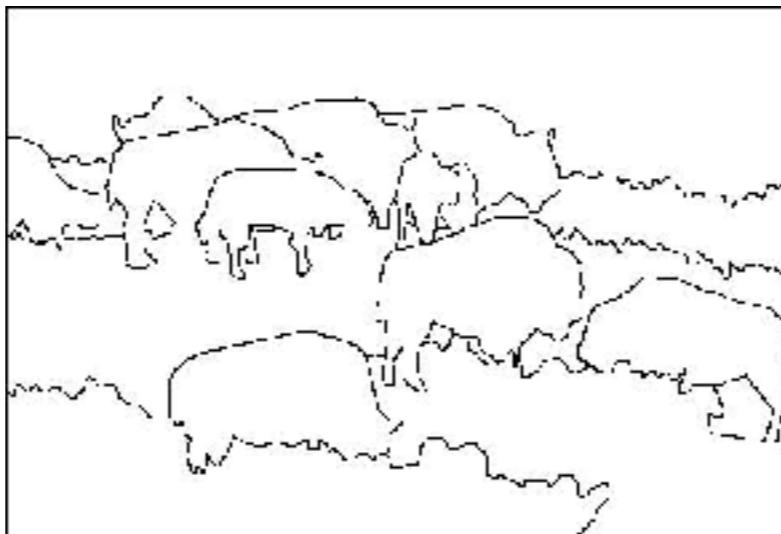
Shadows

Low-level edges vs. perceived contours

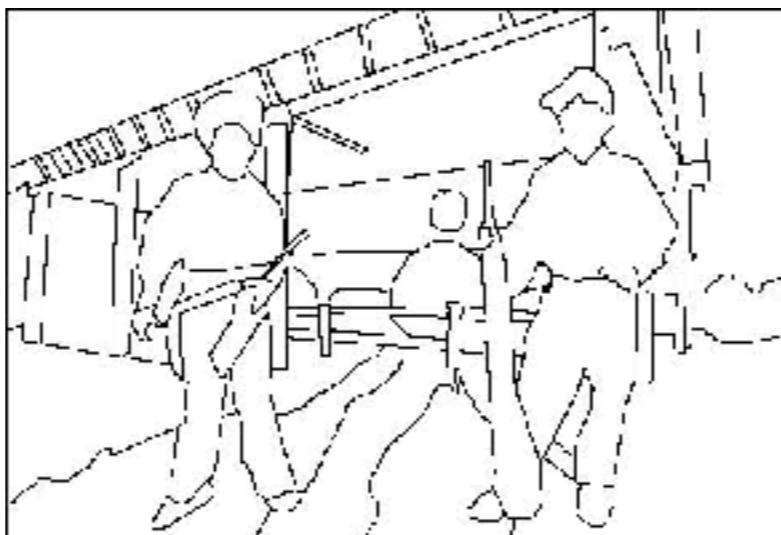
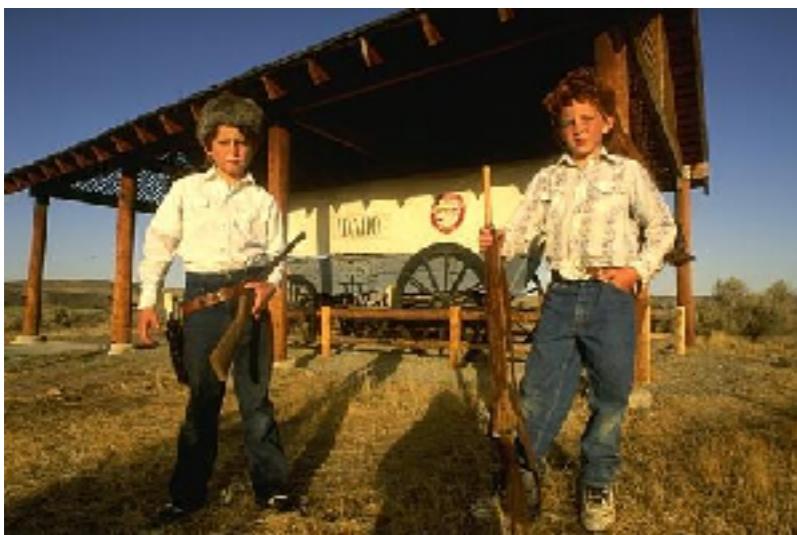
image



human segmentation



gradient magnitude



Seam carving: main idea



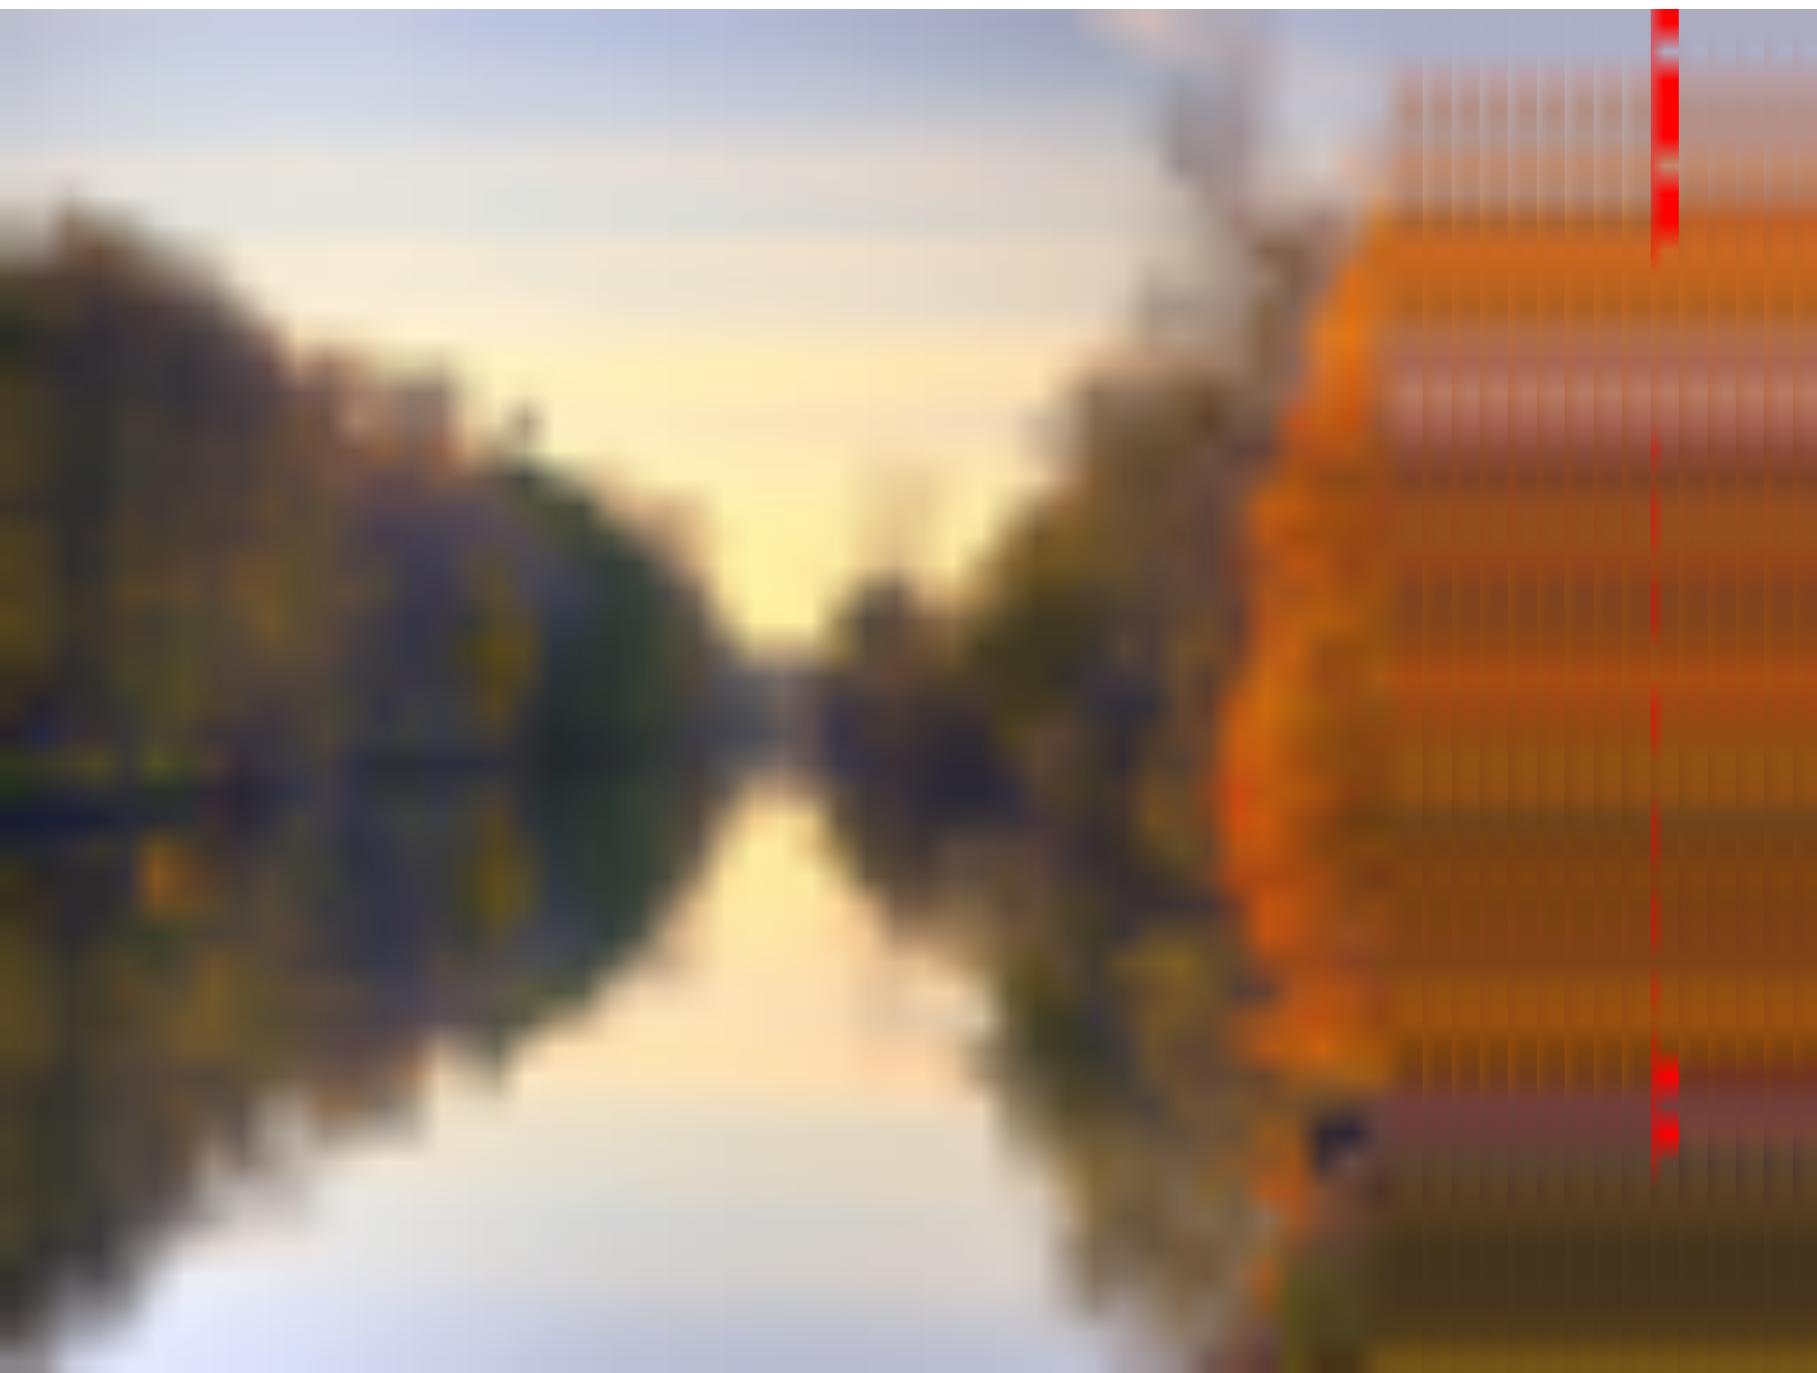
Seam carving: main idea



Content-aware resizing



Traditional resizing



Slide credit: Kristen
Grauman

Seam carving: main idea

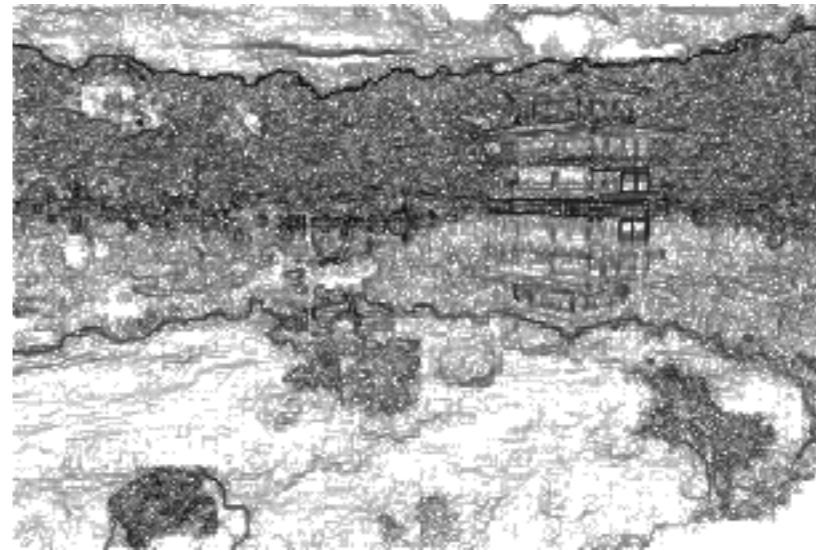


Content-aware resizing

Intuition:

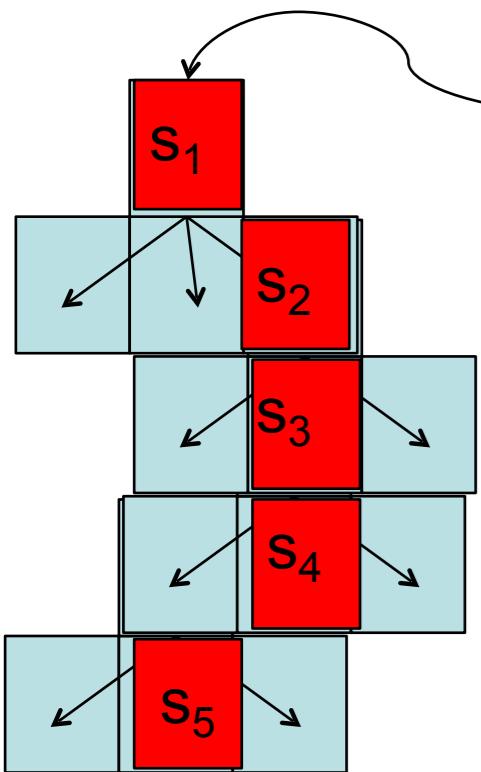
- Preserve the most “interesting” content
 - Prefer to remove pixels with low gradient energy
- To reduce or increase size in one dimension, remove irregularly shaped “seams”
 - Optimal solution via dynamic programming.

Seam carving: main idea



$$Energy(f) = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Seam carving: algorithm



$$Energy(f) = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

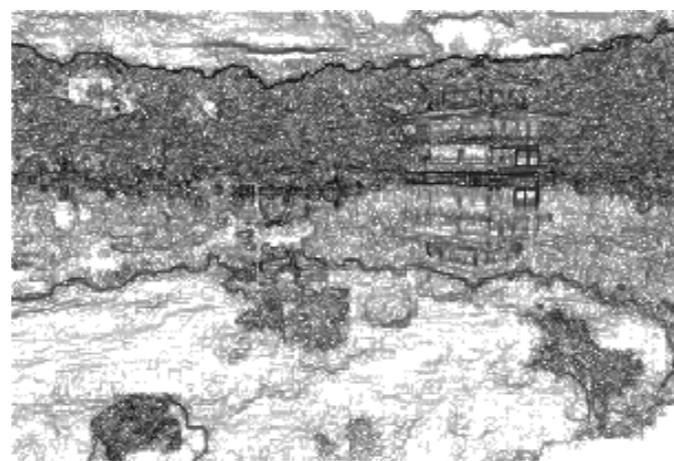
$$Cost(\mathbf{s}) = \sum_{i=1}^h Energy(f(s_i))$$

$$\mathbf{s}^* = \min_{\mathbf{s}} Cost(\mathbf{s})$$

How to identify the minimum cost seam?

- First, consider a greedy approach:

1	3	0
2	8	9
5	2	6

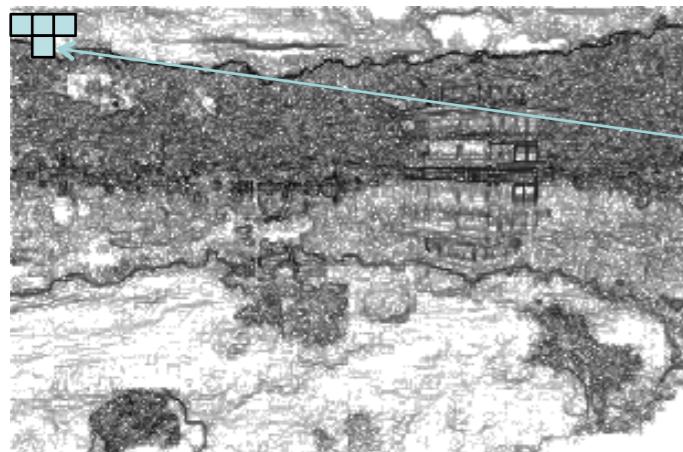


Energy matrix
(gradient magnitude)

Seam carving: algorithm

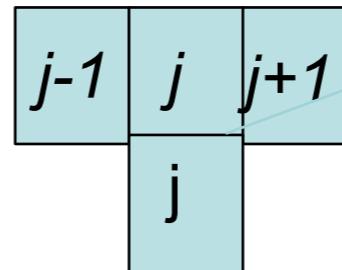
- Compute the *cumulative minimum energy* for all possible connected seams at each entry (i, j) :

$$\mathbf{M}(i, j) = Energy(i, j) + \min(\mathbf{M}(i - 1, j - 1), \mathbf{M}(i - 1, j), \mathbf{M}(i - 1, j + 1))$$



Energy matrix
(gradient magnitude)

row $i-1$
row i



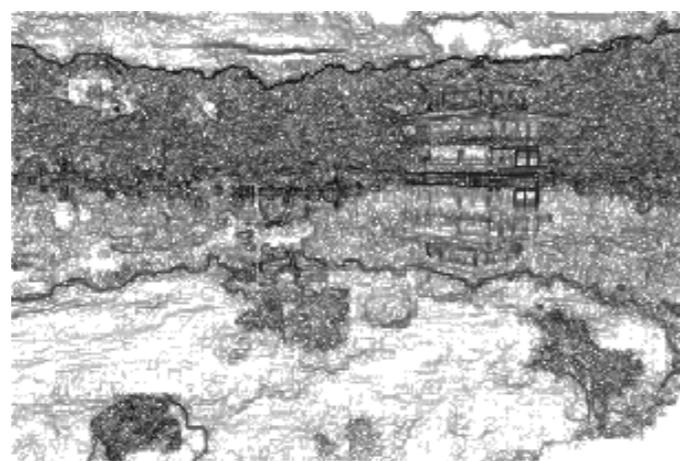
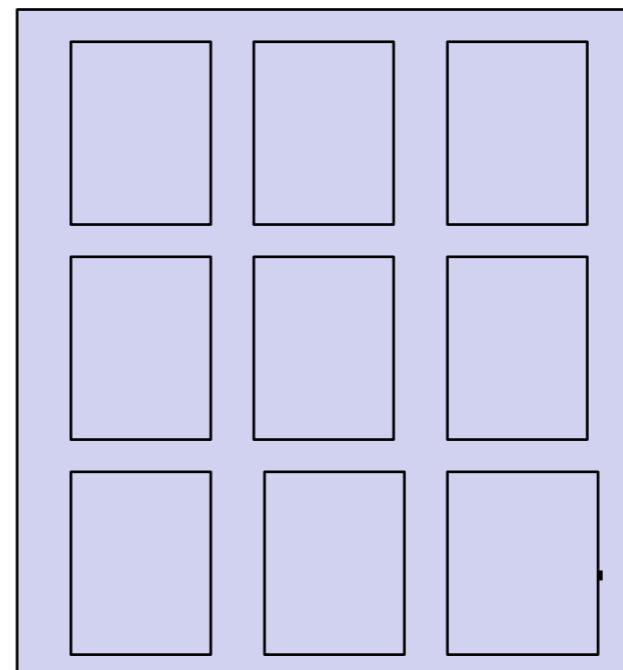
\mathbf{M} matrix:
cumulative min energy
(for vertical seams)

- Then, min value in last row of \mathbf{M} indicates end of the minimal connected vertical seam.
- Backtrack up from there, selecting min of 3 above in \mathbf{M} .

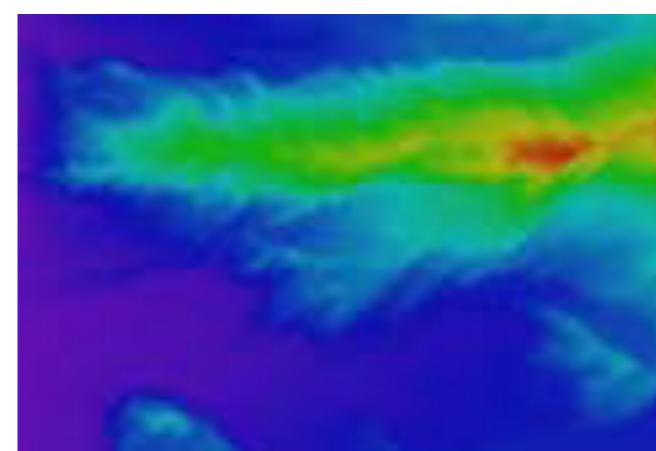
Example

$$\mathbf{M}(i, j) = Energy(i, j) + \min(\mathbf{M}(i - 1, j - 1), \mathbf{M}(i - 1, j), \mathbf{M}(i - 1, j + 1))$$

1	3	0
2	8	9
5	2	6



Energy matrix
(gradient magnitude)



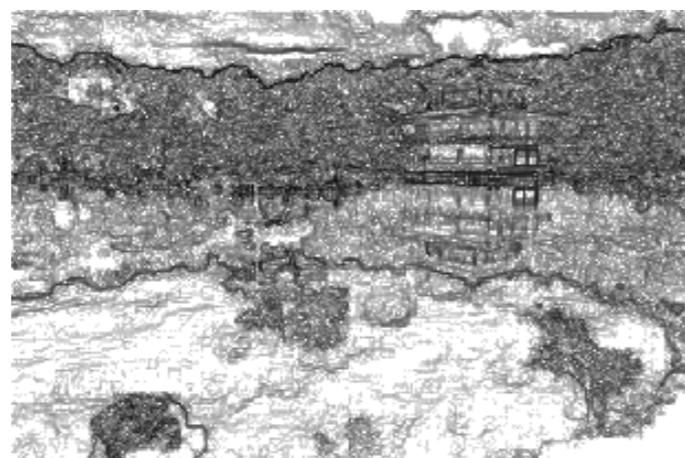
M matrix
(for vertical seams)

Example

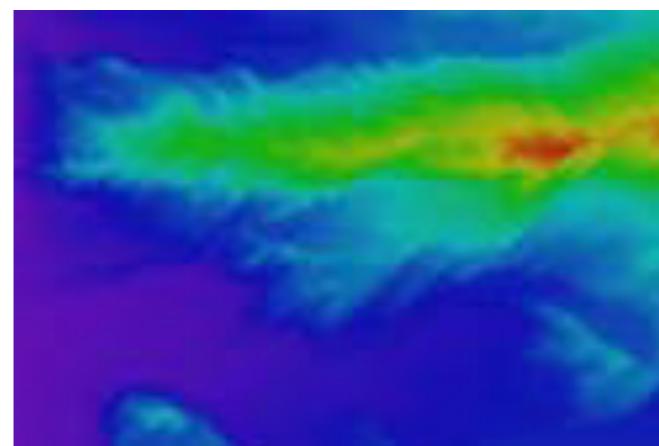
$$\mathbf{M}(i, j) = Energy(i, j) + \min(\mathbf{M}(i - 1, j - 1), \mathbf{M}(i - 1, j), \mathbf{M}(i - 1, j + 1))$$

1	3	0
2	8	9
5	2	6

1	3	0
3	8	9
8	5	14



Energy matrix
(gradient magnitude)



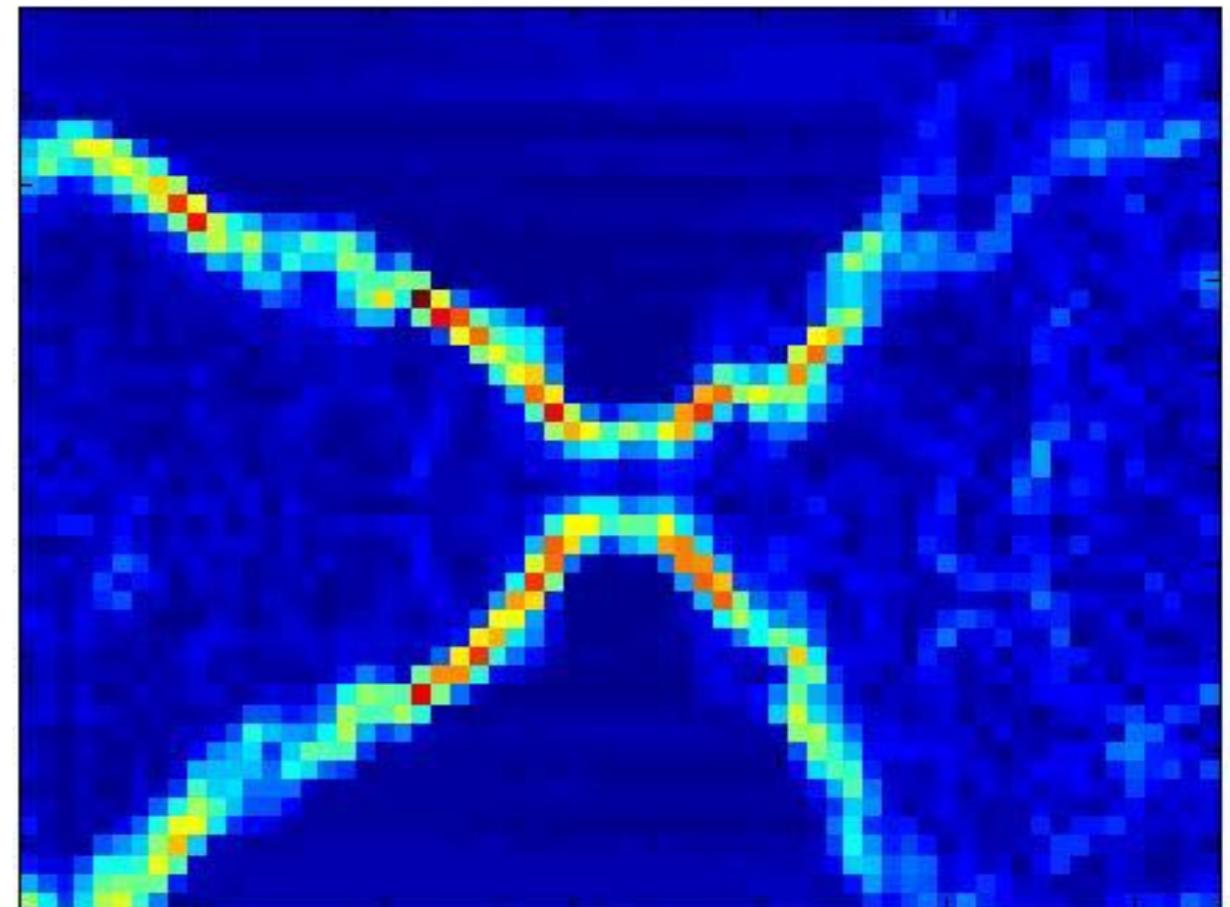
M matrix
(for vertical seams)

Real image example

Original Image



Energy Map



Blue = low energy
Red = high energy

Other notes on seam carving

- Analogous procedure for horizontal seams
- Can also insert seams to increase size of image in either dimension
 - Duplicate optimal seam, averaged with neighbors
- Other energy functions may be plugged in
 - E.g., color-based, interactive,...
- Can use combination of vertical and horizontal seams

One more thing

How to debug
image processing code?



fiona

@fioroco

junior dev: "i found the bug"
senior dev: "i found a bug"

3:24 PM - 4 Dec 2017

Common bugs

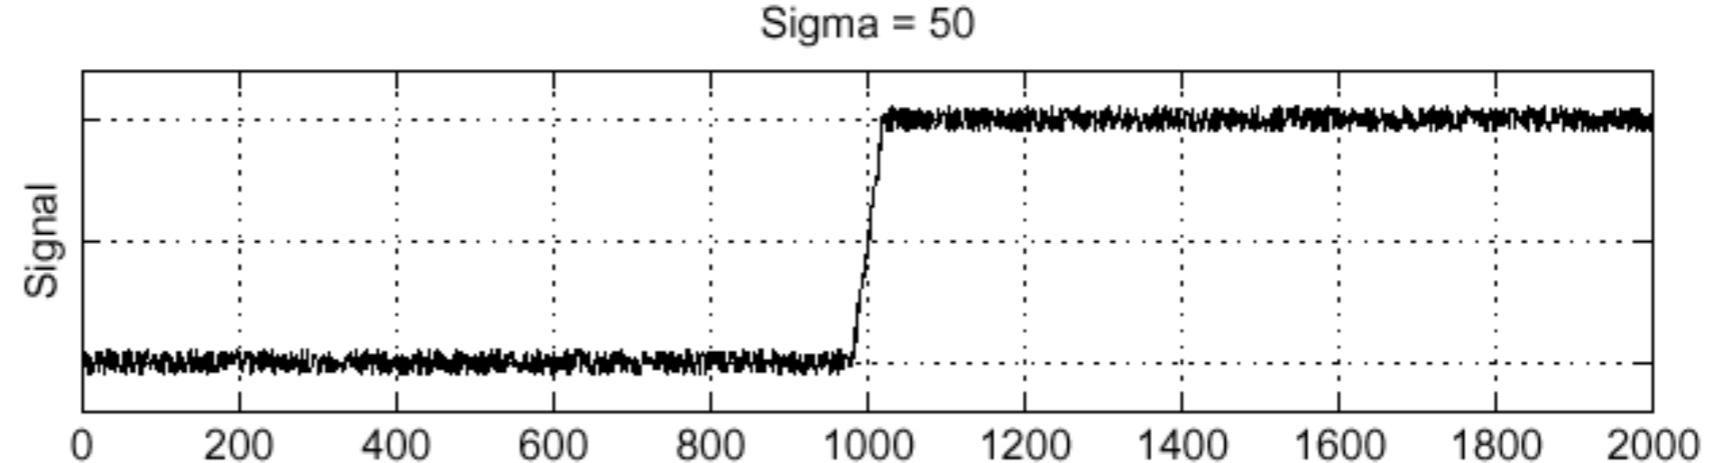
- Color channel order
 - Different libraries handle order differently
 - RGB, BGR, etc.
- Pixel scale
 - uint8: 255 (unsigned integer, default when read in)
 - float: [0, 1]
- Row and column indexing
 - (row, col) or (col, row)
 - Need to be careful when ‘imresize’
- Linear or gamma-corrected color space
 - Gamma curve

Jupyter live demo

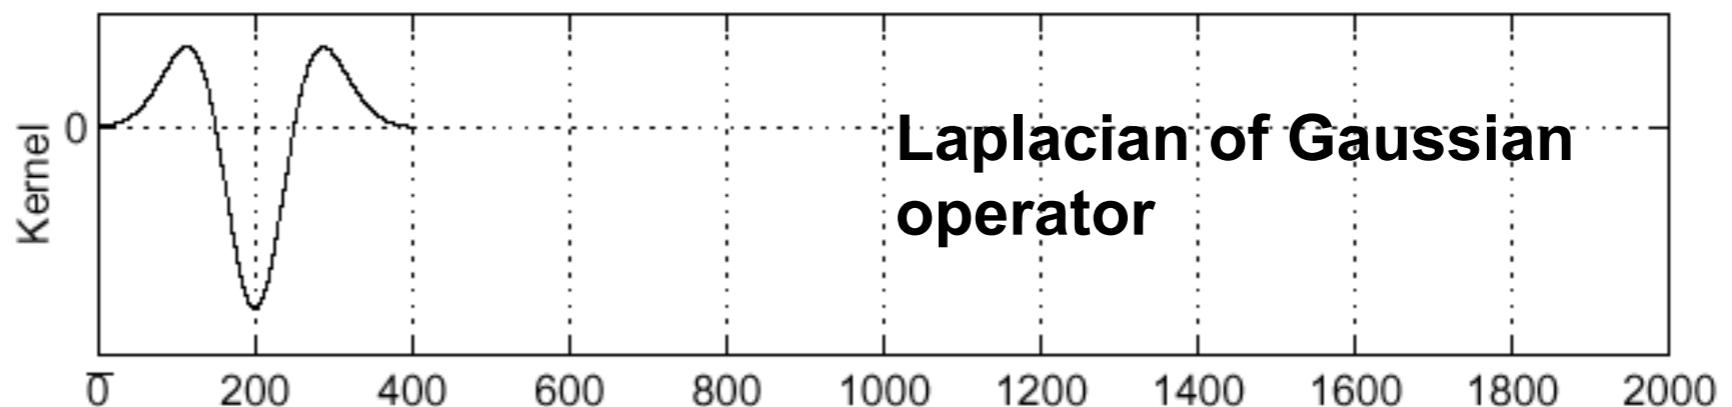
Laplacian of Gaussian

$$\frac{\partial^2}{\partial x^2}(h \star f)$$

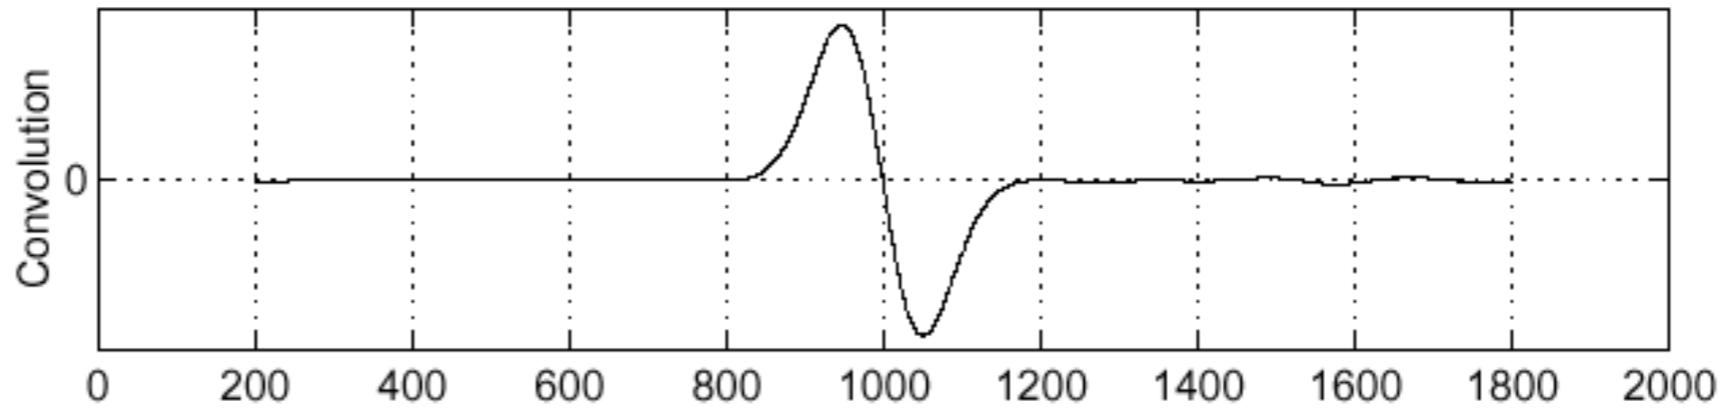
f



$$\frac{\partial^2}{\partial x^2}h$$



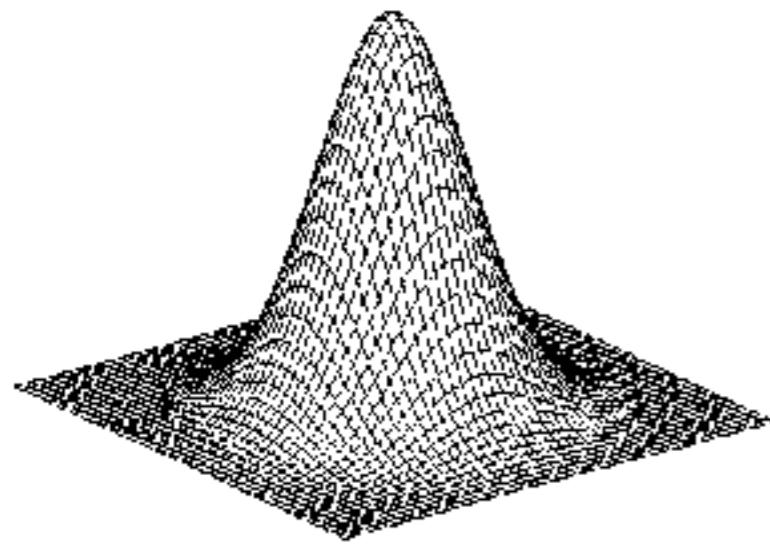
$$(\frac{\partial^2}{\partial x^2}h) \star f$$



Where is the edge?

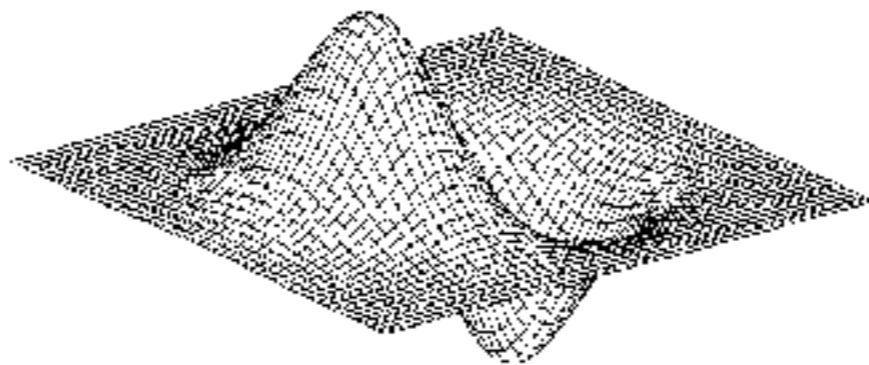
Zero-crossings of bottom graph

2D edge detection filters



Gaussian

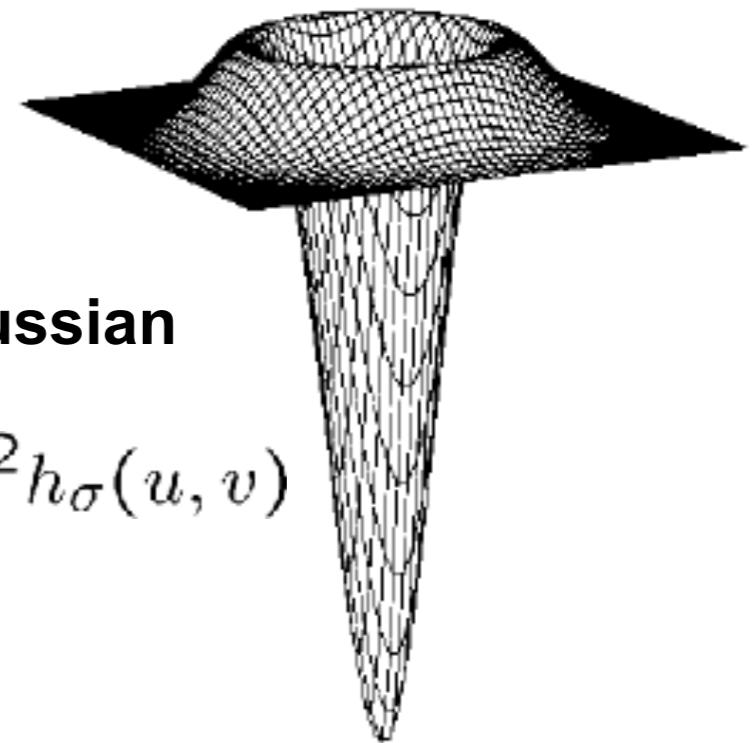
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

$$\nabla^2 h_\sigma(u, v)$$



Laplacian of G

- ∇^2 is the Laplacian operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$